

Data Science 6


Chris Mathys




Master's Degree Programme in Cognitive Science

Spring 2022


Shallow neural networks




Deep neural networks




As we backpropagate errors through the layers, gradients become smaller




Slower learning in earlier layers



Slower learning in earlier layers



Slower learning in earlier layers




Why are gradients unstable?

$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$



Why are gradients unstable?




Why are gradients unstable?

- They are the products of many factors
- The earlier the layer, the more factors the gradients are composed of
- This is an unstable situation
- Mostly, we get vanishing gradients in early layers
- However, by the same multiplication of many factors, we can get exploding gradients


Convolutional neural networks

Local receptive fields




Convolutional neural networks

Local receptive fields




Convolutional neural networks

Local receptive fields



Convolutional neural networks

Local receptive fields



Convolutional neural networks

Shared weights and biases: I've said that each hidden neuron has a bias and 5×5 weights connected to its local receptive field. What I did not yet mention is that we're going to use the *same* weights and bias for each of the 24×24 hidden neurons. In other words, for the j, k th hidden neuron, the output is:

$$\sigma \left(b + \sum_{l=0}^4 \sum_{m=0}^4 w_{l,m} a_{j+l, k+m} \right). \quad (125)$$


Here, σ is the neural activation function - perhaps the [sigmoid function](#) we used in earlier chapters. b is the shared value for the bias. $w_{l,m}$ is a 5×5 array of shared weights. And, finally, we use $a_{x,y}$ to denote the input activation at position x, y .

Convolutional neural networks

This means that all the neurons in the first hidden layer detect exactly the same feature*, just at different locations in the input image. To see why this makes sense, suppose the weights and bias are such that the hidden neuron can pick out, say, a vertical edge in a particular local receptive field. That ability is also likely to be useful at other places in the image. And so it is useful to apply the same feature detector everywhere in the image. To put it in slightly more abstract terms, convolutional networks are well adapted to the translation invariance of images: move a picture of a cat (say) a little ways, and it's still an image of a cat*.

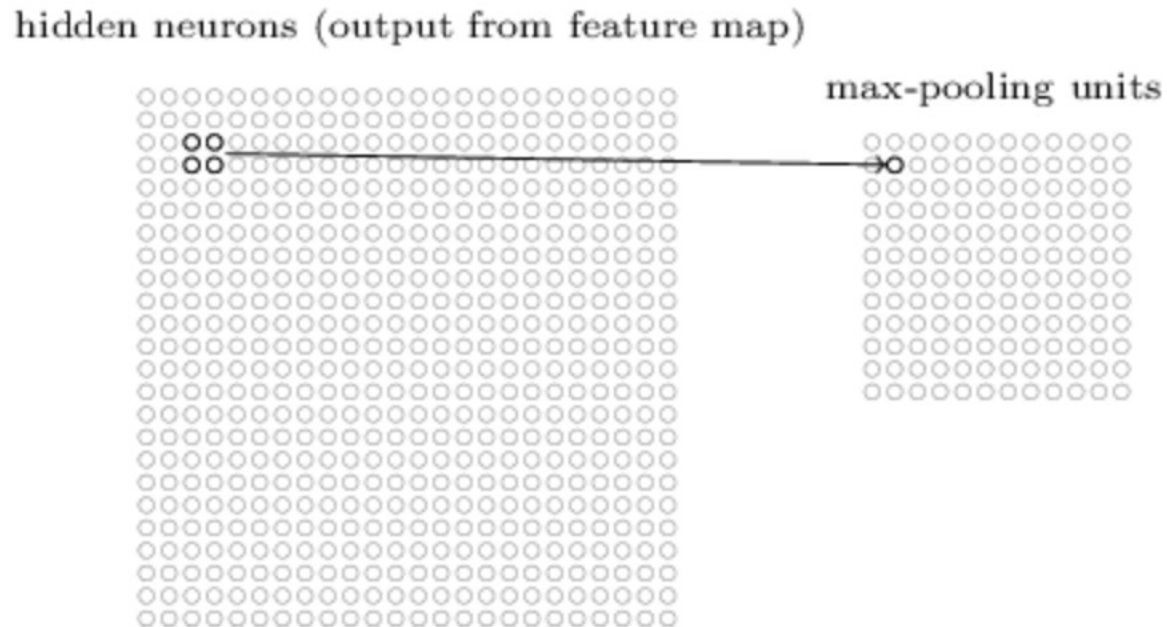
Convolutional neural networks

Feature maps




Convolutional neural networks

Pooling layers are used to condense feature maps




Convolutional neural networks

Pooling layers are used to condense feature maps



Convolutional neural networks

Putting it all together



Convolutional neural networks

Further resources:

ICTP Winter School on Learning and Artificial Intelligence

<https://indico.ictp.it/event/8339/>

Materials and videos:

<https://indico.ictp.it/event/8339/other-view?view=ictp timetable>