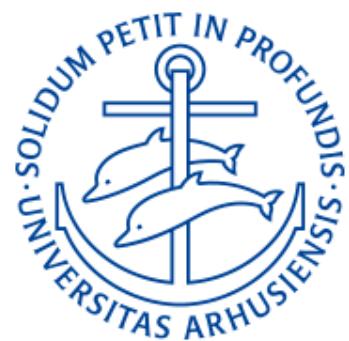


# Data Science 8

**Chris Mathys**



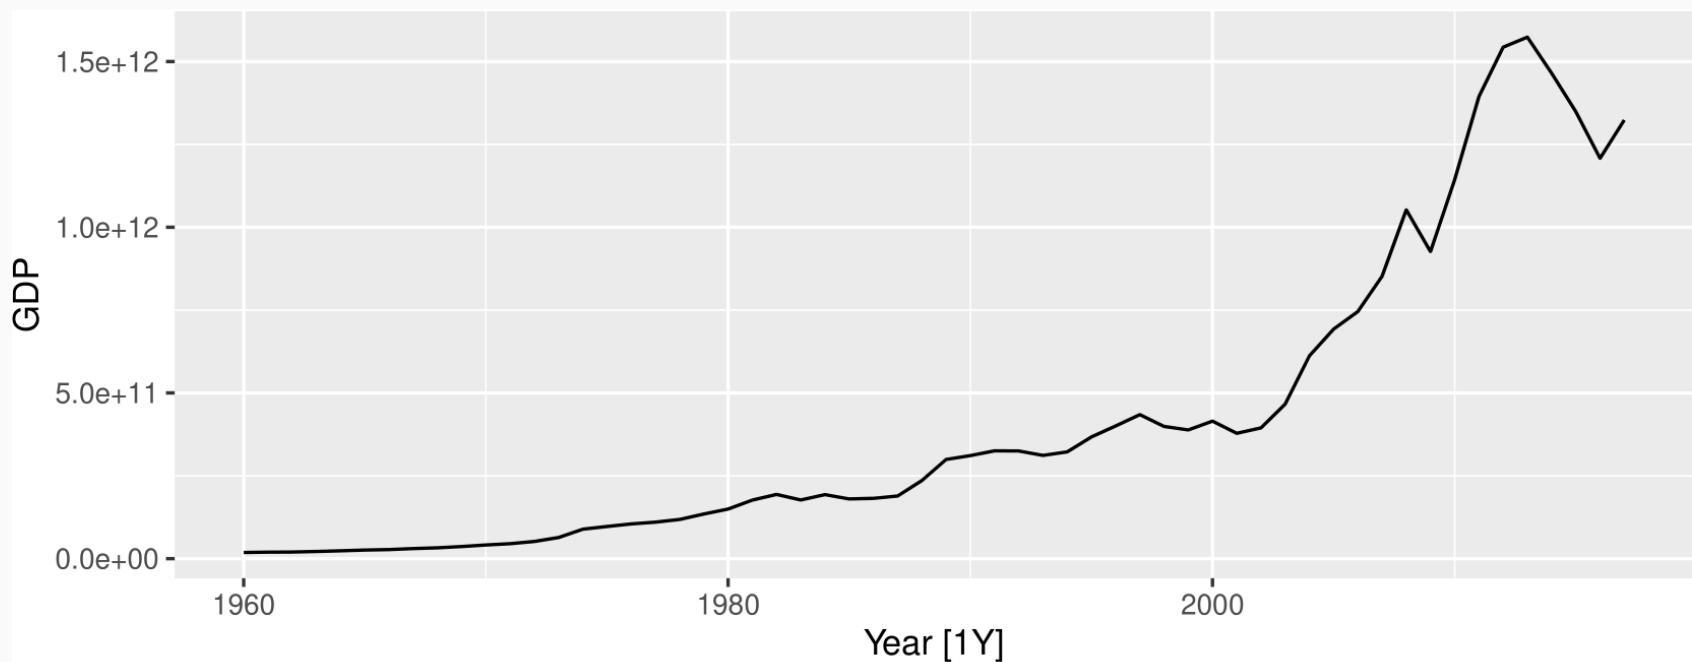
Master's Degree Programme in Cognitive Science  
Spring 2022

# Outline

- 1 Transformations and adjustments
- 2 Time series components
- 3 History of time series decomposition
- 4 STL decomposition
- 5 When things go wrong

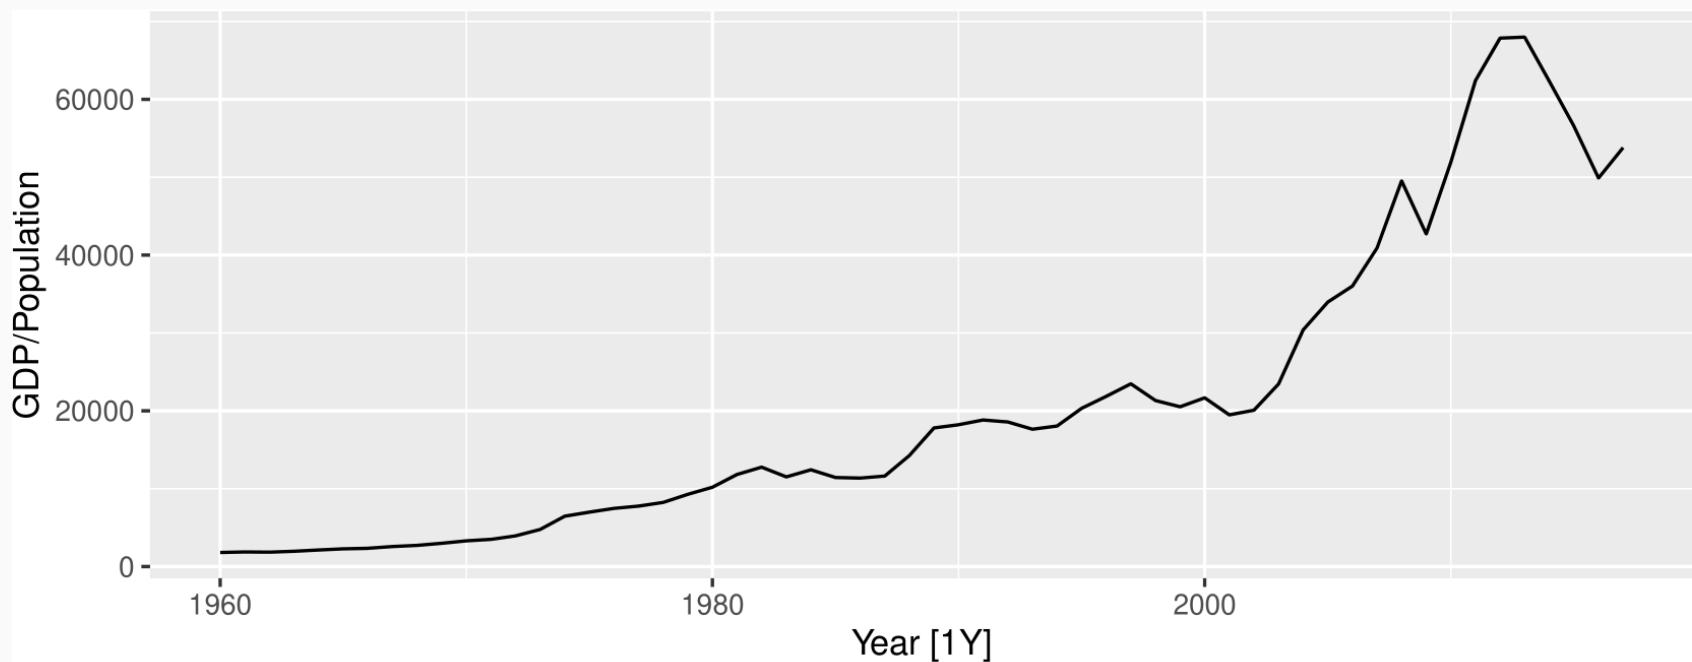
# Per capita adjustments

```
global_economy %>%
  filter(Country == "Australia") %>%
  autoplot(GDP)
```



# Per capita adjustments

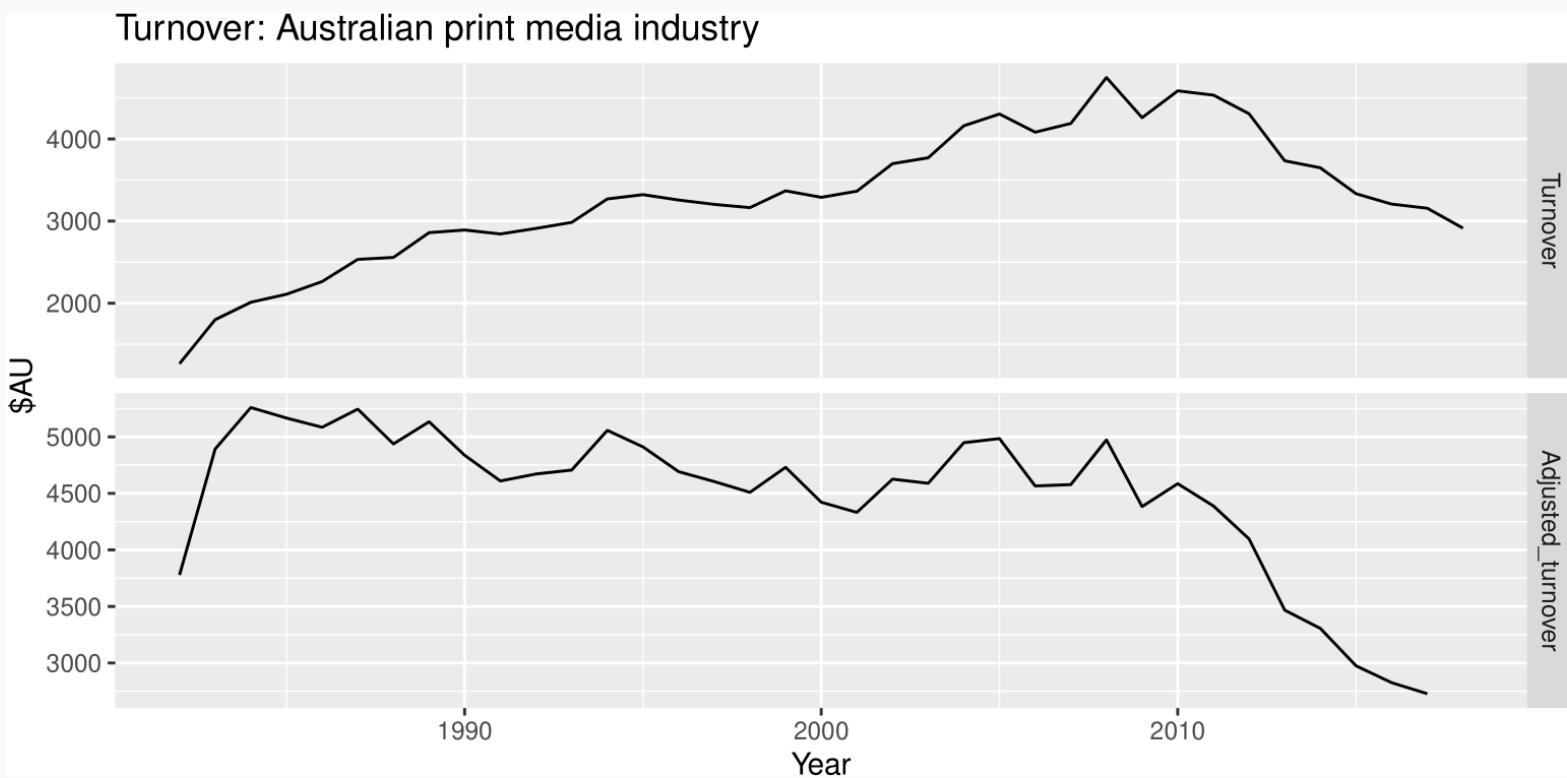
```
global_economy %>%
  filter(Country == "Australia") %>%
  autoplot(GDP / Population)
```



# Inflation adjustments

```
print_retail <- aus_retail %>%
  filter(Industry == "Newspaper and book retailing") %>%
  group_by(Industry) %>%
  index_by(Year = year(Month)) %>%
  summarise(Turnover = sum(Turnover))
aus_economy <- global_economy %>%
  filter(Code == "AUS")
print_retail %>%
  left_join(aus_economy, by = "Year") %>%
  mutate(Adjusted_turnover = Turnover / CPI * 100) %>%
  pivot_longer(c(Turnover, Adjusted_turnover), values_to = "Turnover") %>%
  mutate(name = factor(name, levels=c("Turnover","Adjusted_turnover"))) %>%
  ggplot(aes(x = Year, y = Turnover)) +
  geom_line() +
  facet_grid(name ~ ., scales = "free_y") +
  labs(title = "Turnover: Australian print media industry", y = "$AU")
```

# Inflation adjustments



# Mathematical transformations

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

## Mathematical transformations for stabilizing variation

Square root     $w_t = \sqrt{y_t}$               ↓

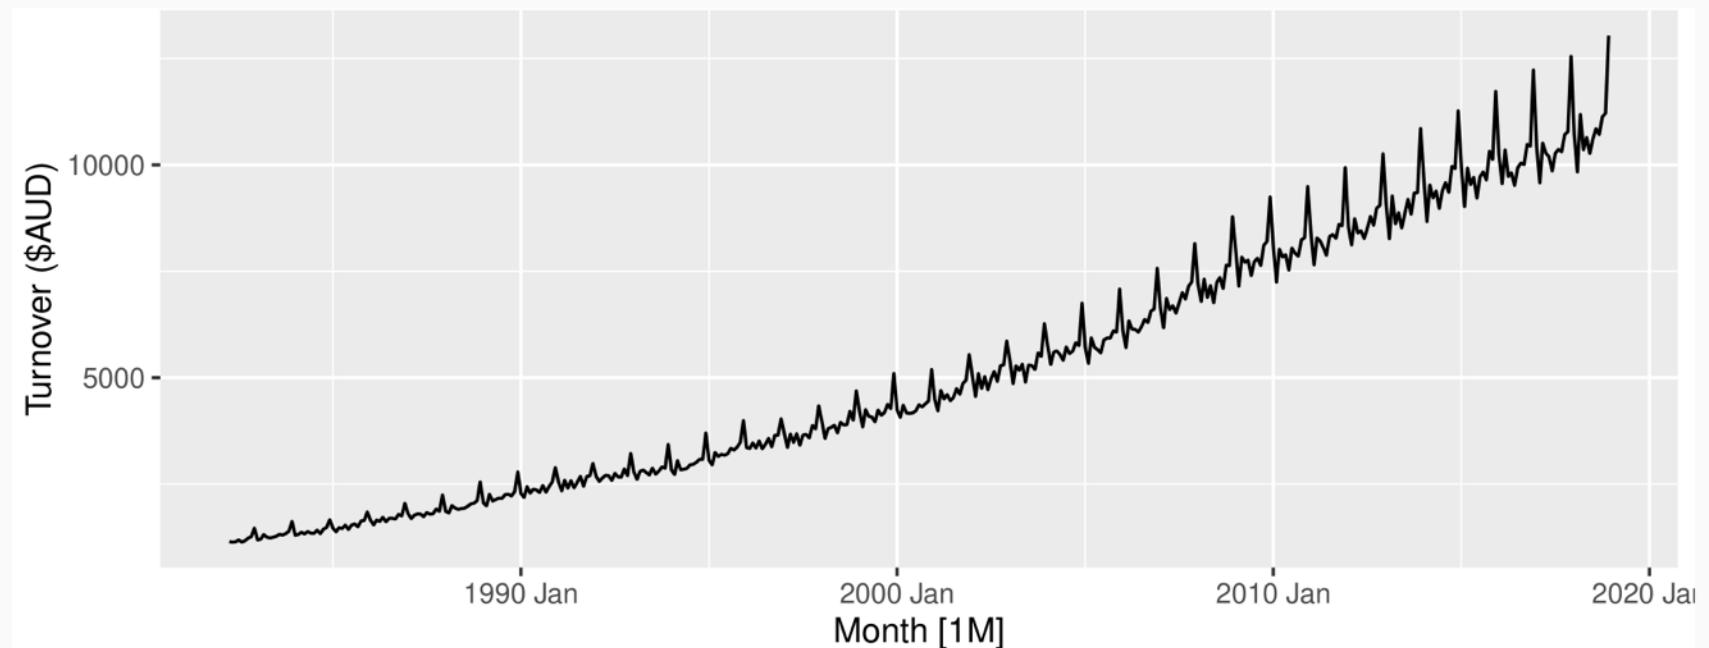
Cube root     $w_t = \sqrt[3]{y_t}$       Increasing

Logarithm     $w_t = \log(y_t)$       strength

Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes on the original scale**.

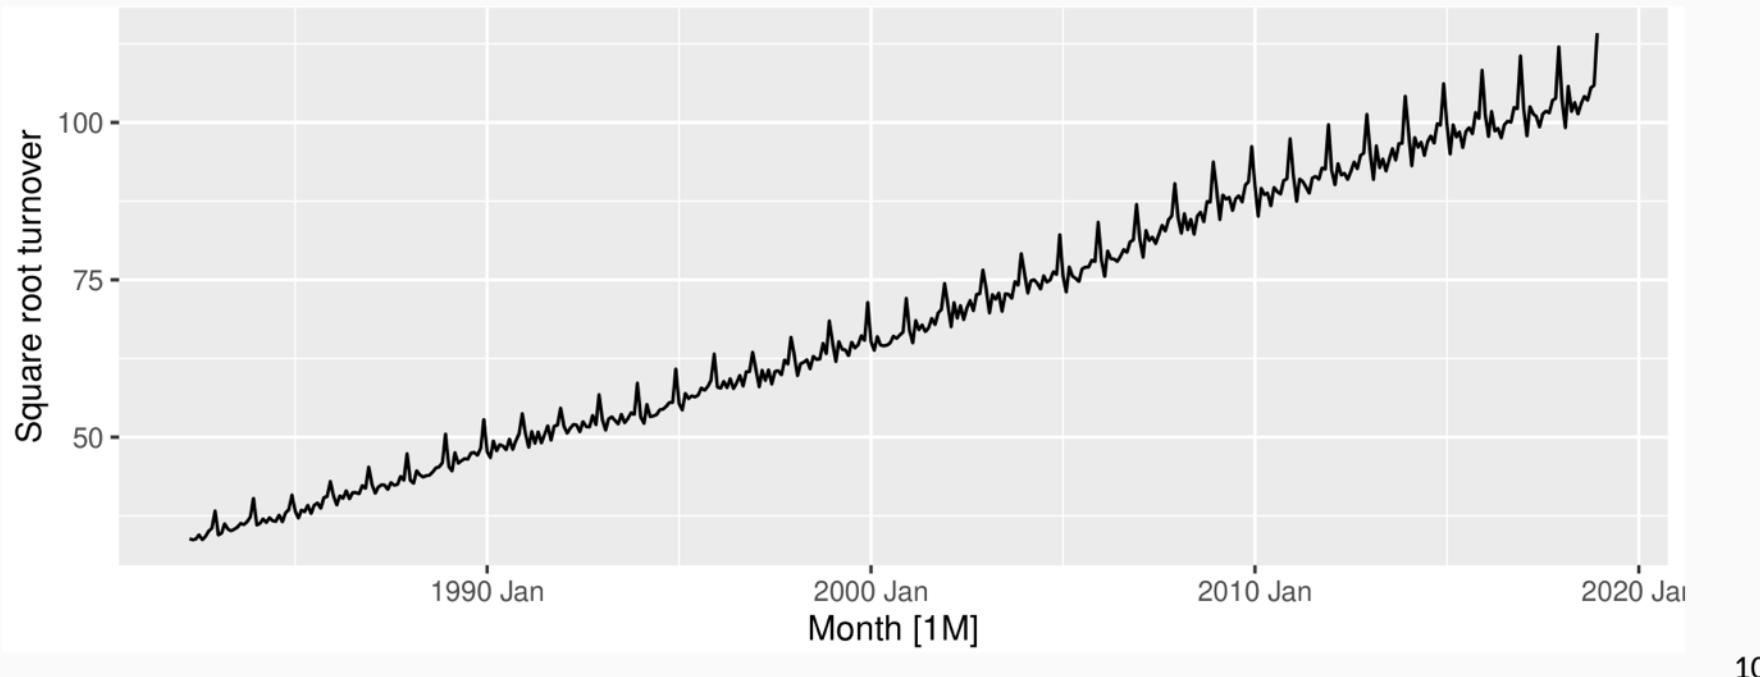
# Mathematical transformations

```
food <- aus_retail %>%
  filter(Industry == "Food retailing") %>%
  summarise(Turnover = sum(Turnover))
```



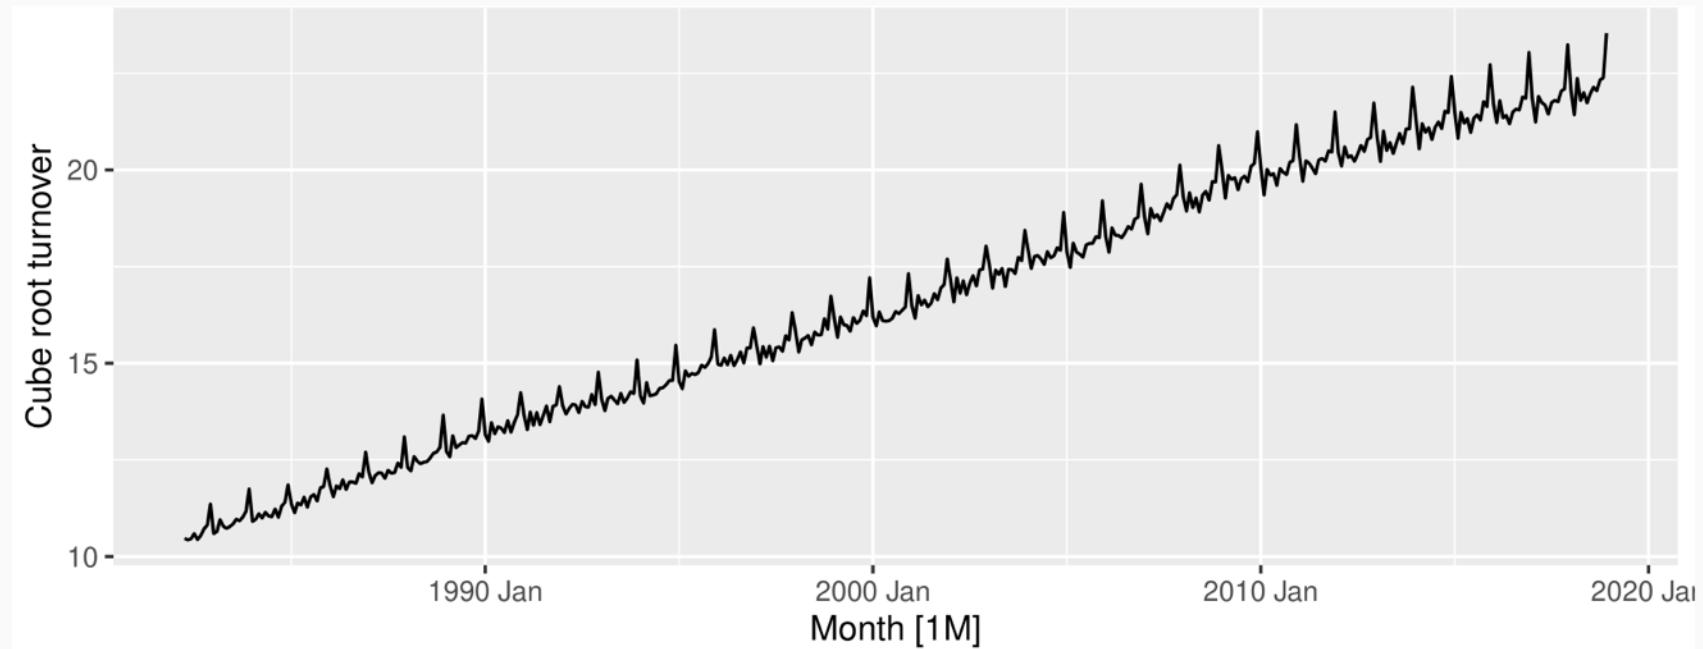
# Mathematical transformations

```
food %>% autoplot(sqrt(Turnover)) +  
  labs(y = "Square root turnover")
```



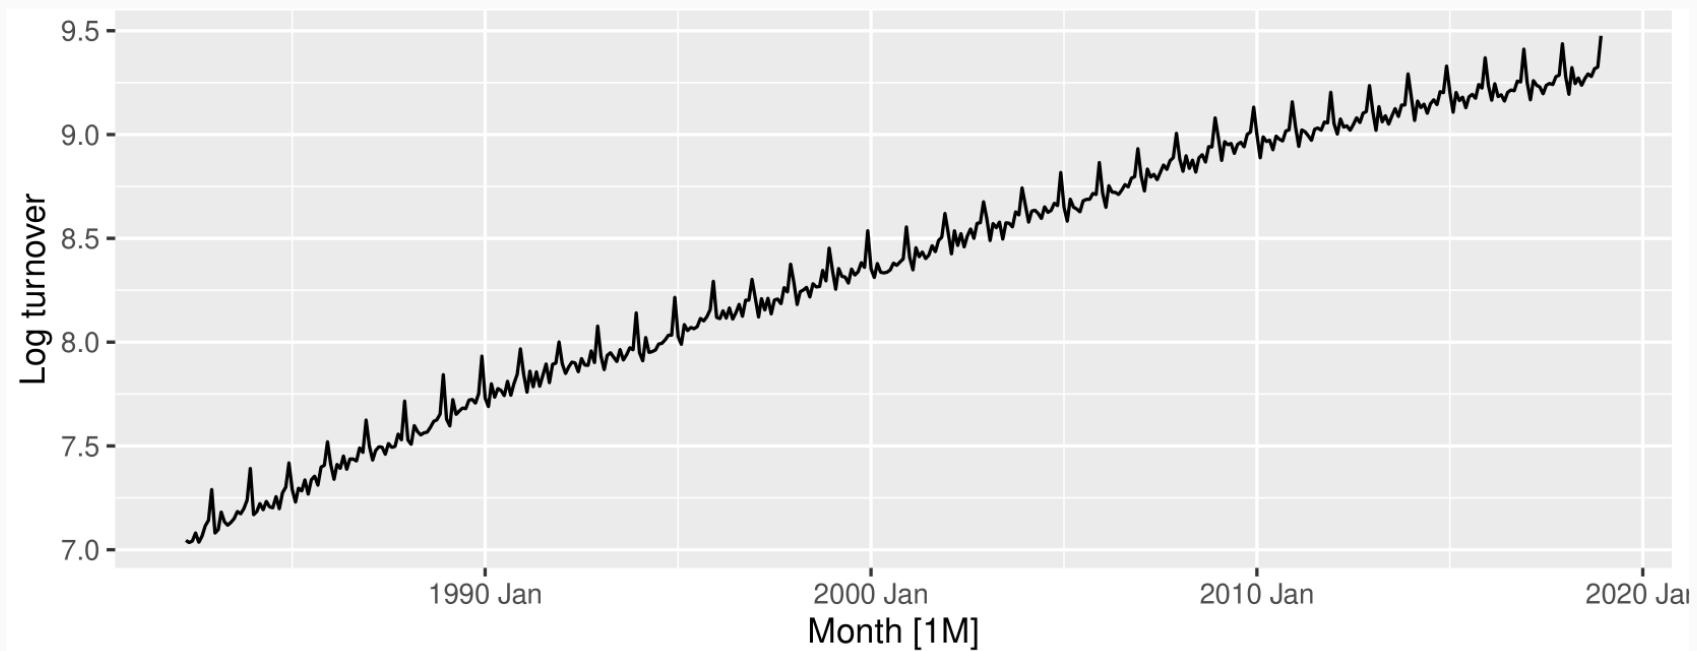
# Mathematical transformations

```
food %>% autoplot(Turnover^(1/3)) +  
  labs(y = "Cube root turnover")
```



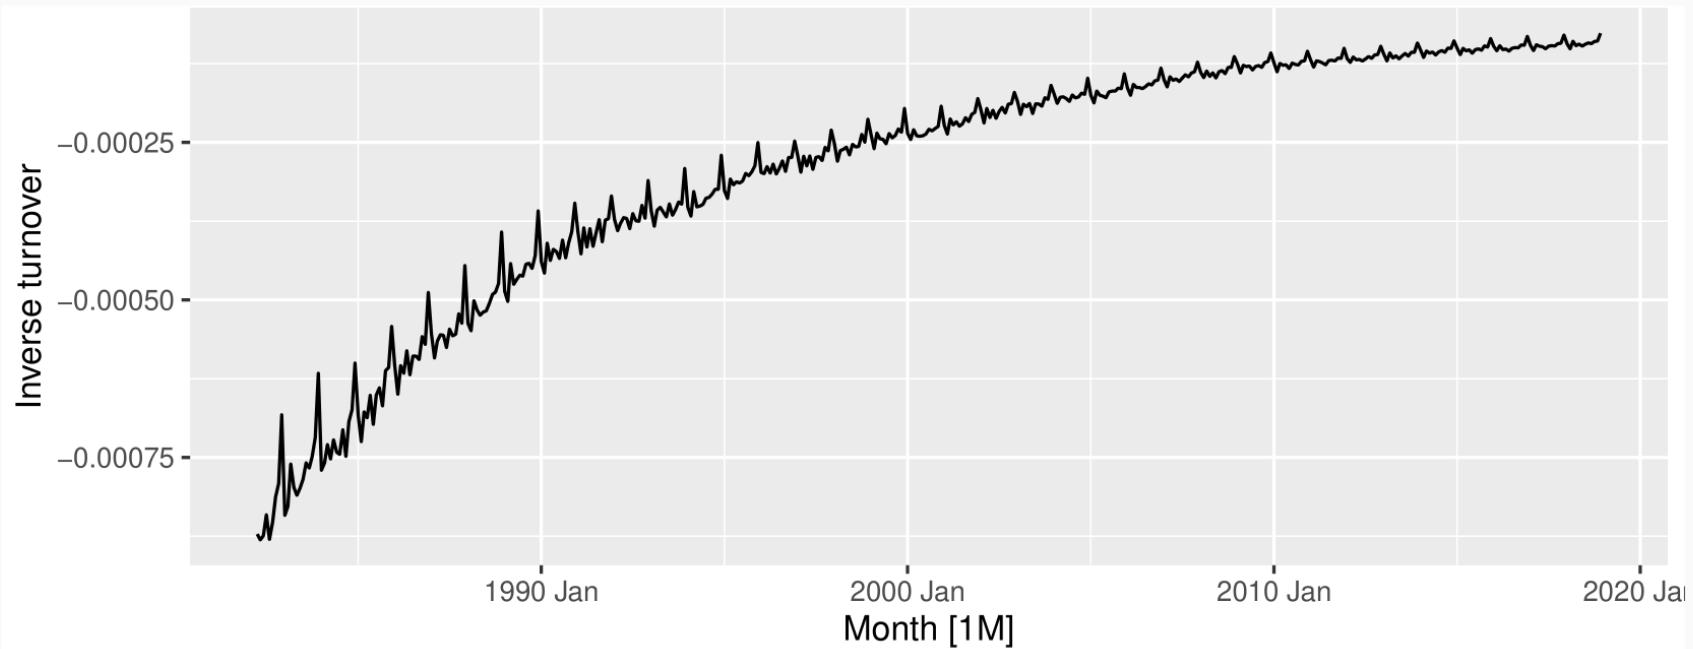
# Mathematical transformations

```
food %>% autoplot(log(Turnover)) +  
  labs(y = "Log turnover")
```



# Mathematical transformations

```
food %>% autoplot(-1/Turnover) +  
  labs(y = "Inverse turnover")
```



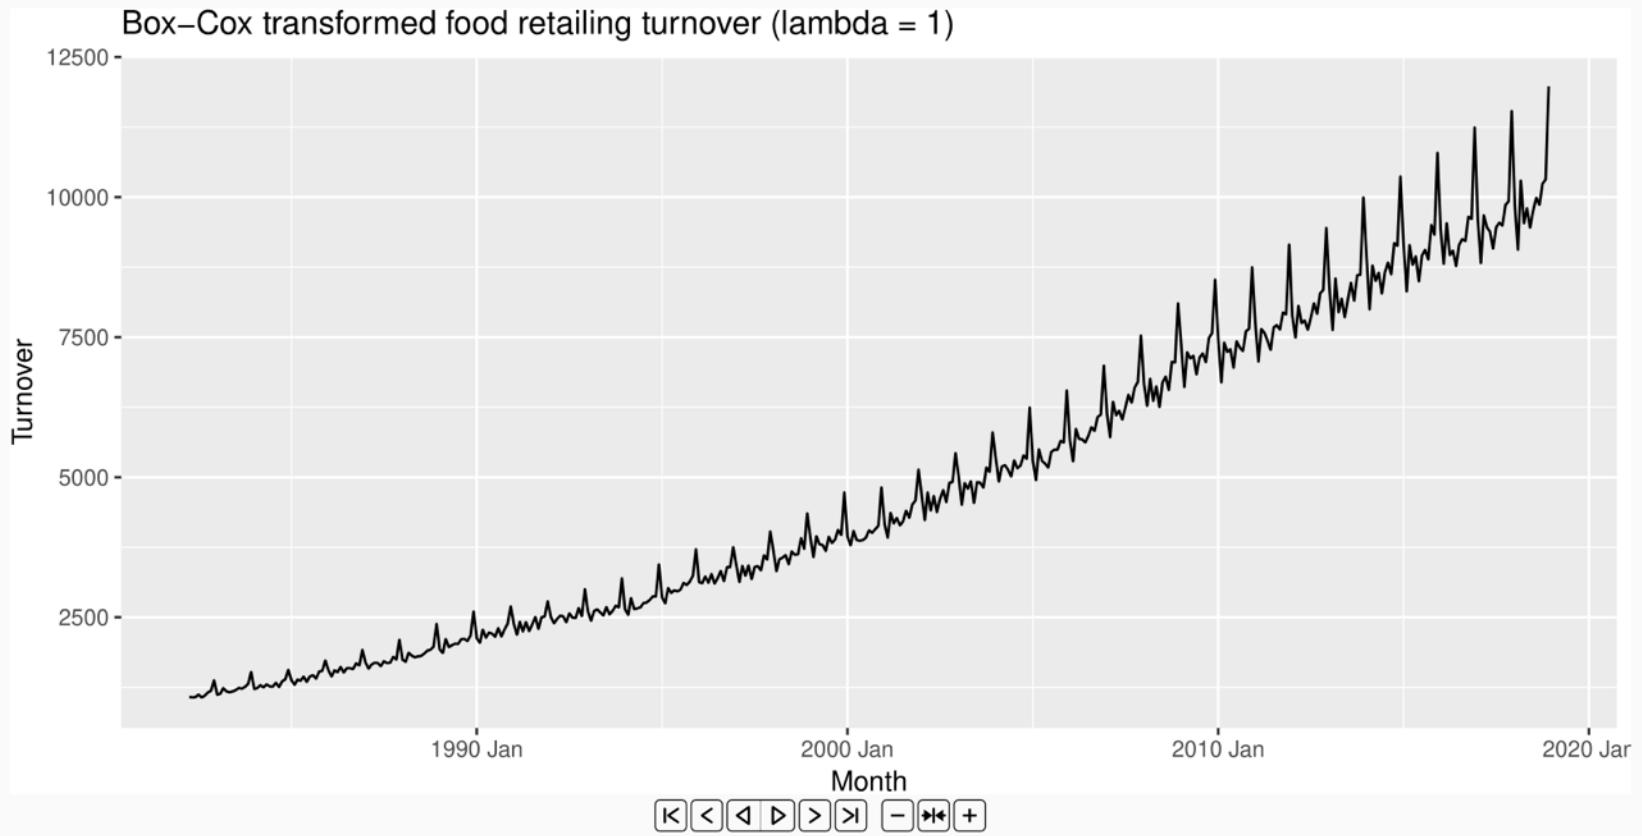
# Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (\text{sign}(y_t)|y_t|^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- Actually the Bickel-Doksum transformation (allowing for  $y_t < 0$ )
- $\lambda = 1$ : (No substantive transformation)
- $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)
- $\lambda = 0$ : (Natural logarithm)
- $\lambda = -1$ : (Inverse plus 1)

# Box-Cox transformations



# Box-Cox transformations

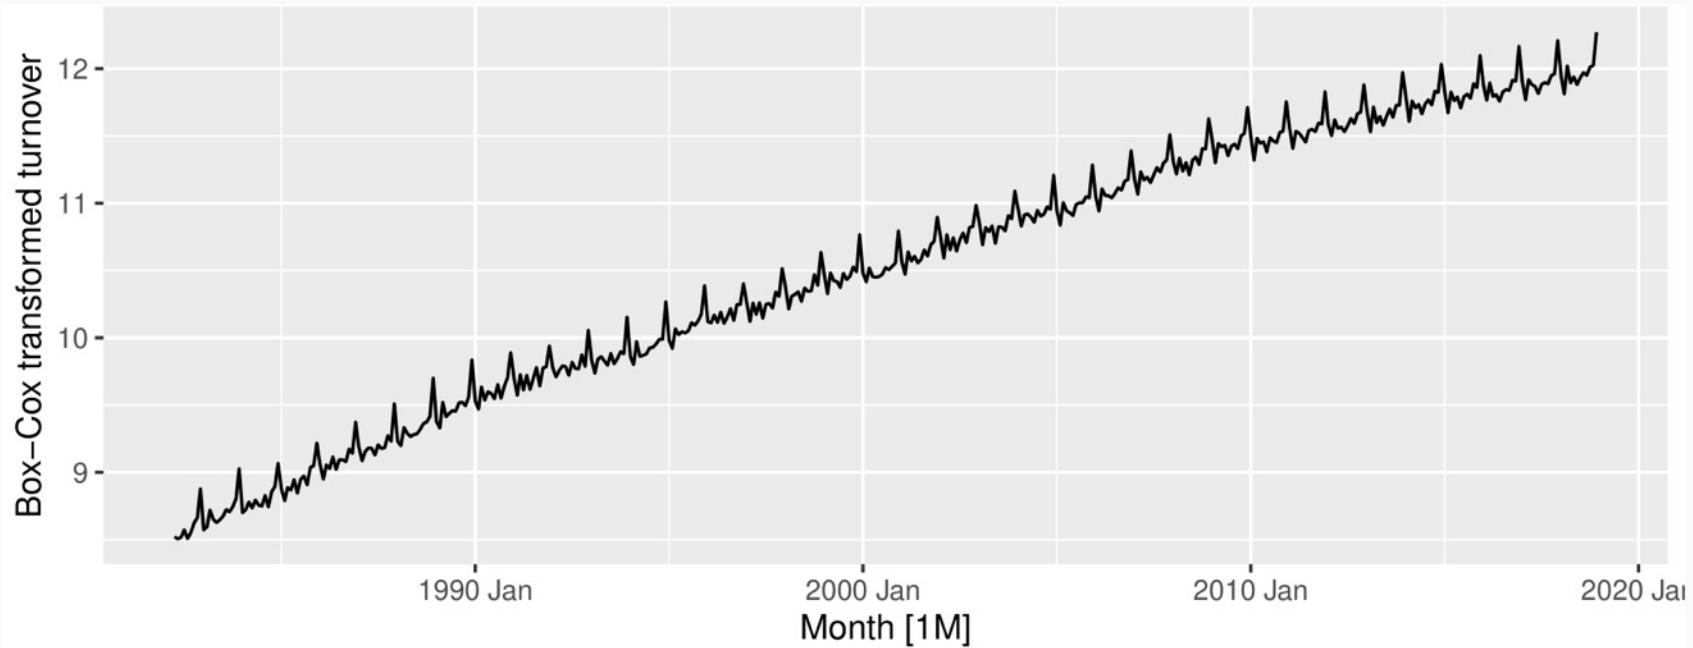
```
food %>%
  features(Turnover, features = guerero)
```

```
## # A tibble: 1 x 1
##   lambda_guerero
##       <dbl>
## 1      0.0524
```

- This attempts to balance the seasonal fluctuations and random variation across the series.
- Always check the results.
- A low value of  $\lambda$  can give extremely large prediction intervals.

# Box-Cox transformations

```
food %>% autoplot(box_cox(Turnover, 0.0524)) +  
  labs(y = "Box-Cox transformed turnover")
```



# Transformations

- Often no transformation needed.
- Simple transformations are easier to explain and work well enough.
- Transformations can have very large effect on PI.
- If some data are zero or negative, then use  $\lambda > 0$ .
- `log1p()` can also be useful for data with zeros.
- Choosing logs is a simple way to force forecasts to be positive
- Transformations must be reversed to obtain forecasts on the original scale. (Handled automatically by `fable`.)

# Outline

- 1 Transformations and adjustments
- 2 Time series components
- 3 History of time series decomposition
- 4 STL decomposition
- 5 When things go wrong

# Time series patterns

## Recall

**Trend** pattern exists when there is a long-term increase or decrease in the data.

**Cyclic** pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).

**Seasonal** pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

# Time series decomposition

$$y_t = f(S_t, T_t, R_t)$$

where  $y_t$  = data at period  $t$   
 $T_t$  = trend-cycle component at period  $t$   
 $S_t$  = seasonal component at period  $t$   
 $R_t$  = remainder component at period  $t$

# Time series decomposition

$$y_t = f(S_t, T_t, R_t)$$

where  $y_t$  = data at period  $t$

$T_t$  = trend-cycle component at period  $t$

$S_t$  = seasonal component at period  $t$

$R_t$  = remainder component at period  $t$

**Additive decomposition:**  $y_t = S_t + T_t + R_t.$

**Multiplicative decomposition:**  $y_t = S_t \times T_t \times R_t.$

# Time series decomposition

- Additive model appropriate if magnitude of seasonal fluctuations does not vary with level.
- If seasonal are proportional to level of series, then multiplicative model appropriate.
- Multiplicative decomposition more prevalent with economic series
- Alternative: use a Box-Cox transformation, and then use additive decomposition.
- Logs turn multiplicative relationship into an additive relationship:

$$y_t = S_t \times T_t \times R_t \quad \Rightarrow \quad \log y_t = \log S_t + \log T_t + \log R_t.$$

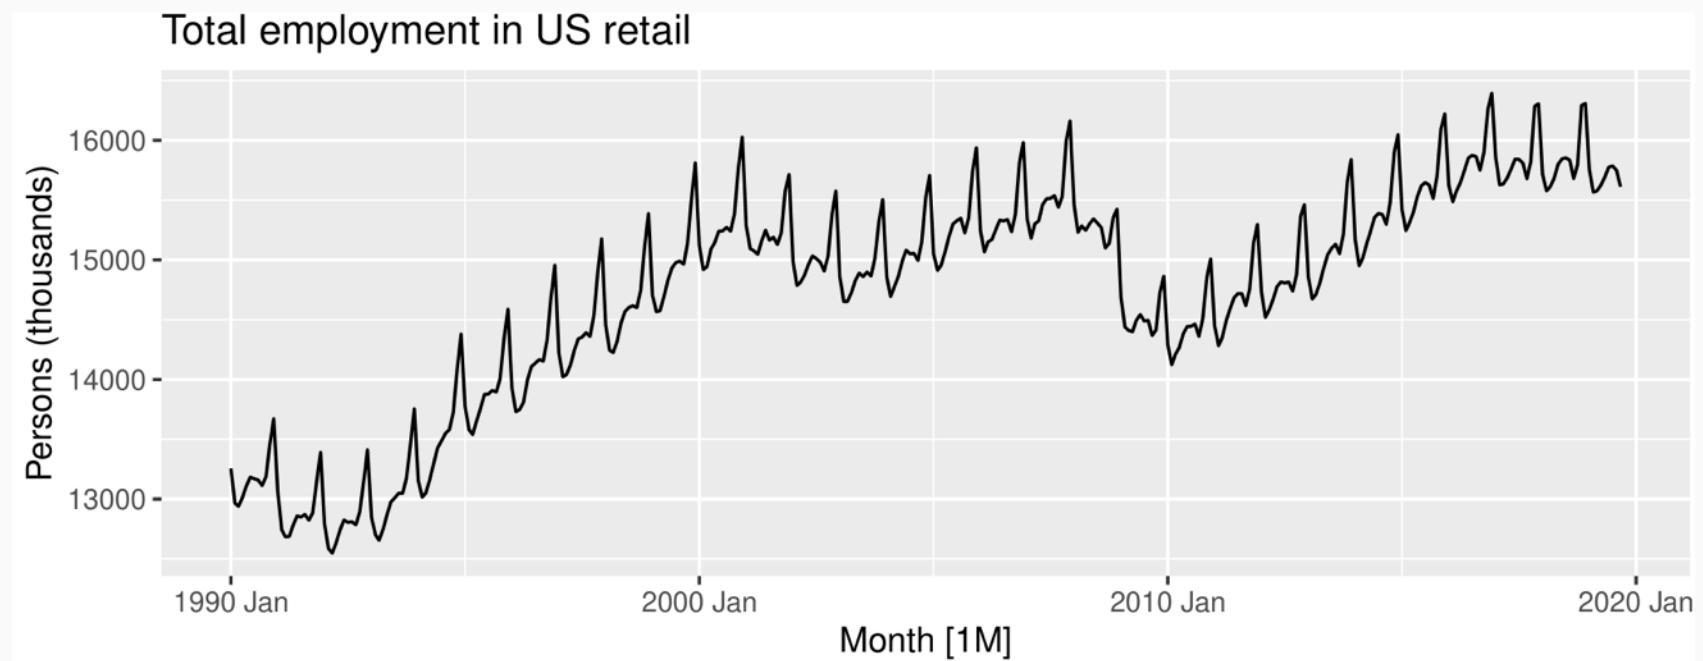
# US Retail Employment

```
us_retail_employment <- us_employment %>%
  filter(year(Month) >= 1990, Title == "Retail Trade") %>%
  select(-Series_ID)
us_retail_employment
```

```
## # A tsibble: 357 x 3 [1M]
##       Month Title     Employed
##       <mth> <chr>     <dbl>
## 1 1990 Jan Retail Trade 13256.
## 2 1990 Feb Retail Trade 12966.
## 3 1990 Mar Retail Trade 12938.
## 4 1990 Apr Retail Trade 13012.
## 5 1990 May Retail Trade 13108.
## 6 1990 Jun Retail Trade 13183.
## 7 1990 Jul Retail Trade 13170.
## 8 1990 Aug Retail Trade 13160.
## 9 1990 Sep Retail Trade 12112
```

# US Retail Employment

```
us_retail_employment %>%
  autoplot(Employed) +
  labs(y="Persons (thousands)", title="Total employment in US retail")
```



# US Retail Employment

```
us_retail_employment %>%
  model(stl = STL(Employed))
```

```
## # A mable: 1 x 1
##       stl
##     <model>
## 1   <STL>
```

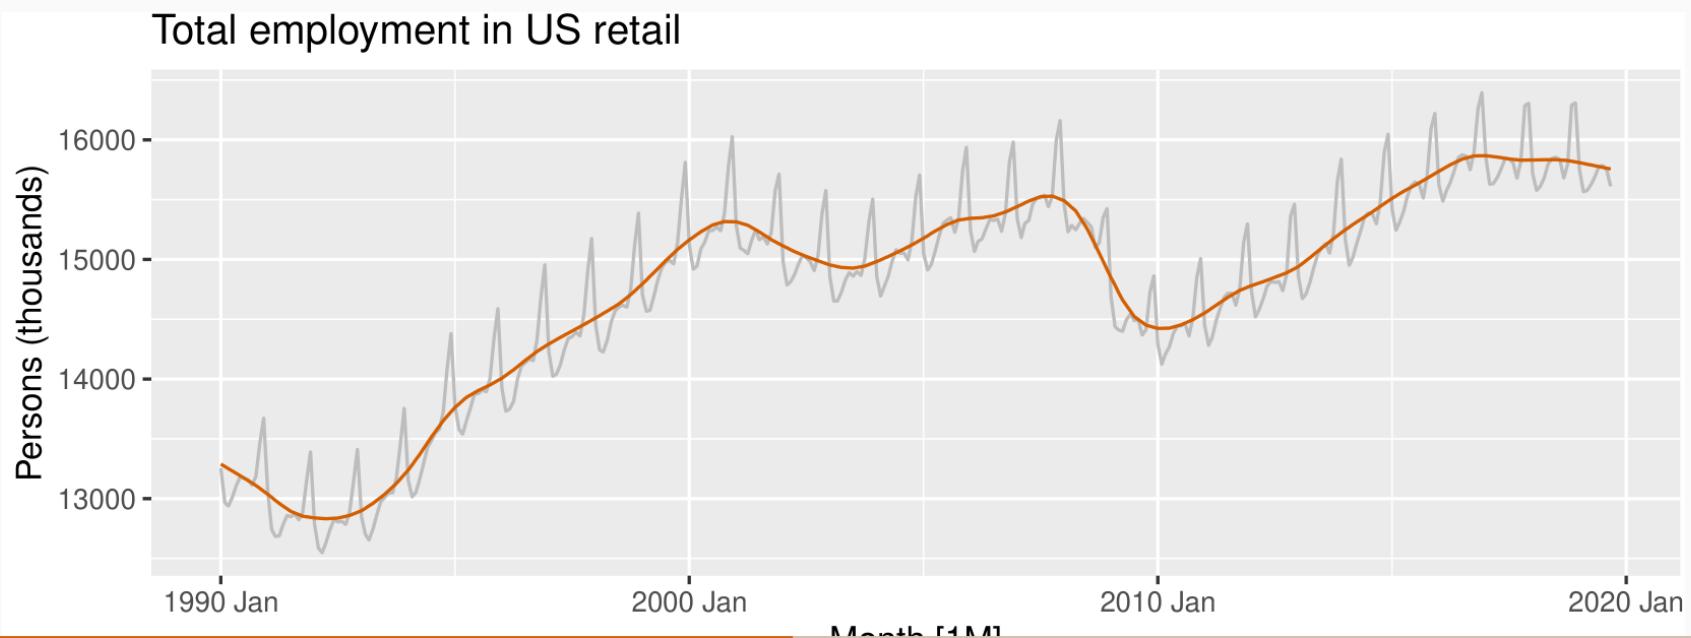
# US Retail Employment

```
dcmp <- us_retail_employment %>%
  model(stl = STL(Employed))
components(dcmp)
```

```
## # A dable: 357 x 7 [1M]
## # Key:   .model [1]
## # :     Employed = trend + season_year + remainder
##   .model   Month Employed trend season_year remainder season_adjust
##   <chr>    <mth>   <dbl>  <dbl>      <dbl>      <dbl>      <dbl>
## 1 stl     1990 Jan    13256. 13288.     -33.0      0.836    13289.
## 2 stl     1990 Feb    12966. 13269.     -258.     -44.6     13224.
## 3 stl     1990 Mar    12938. 13250.     -290.     -22.1     13228.
## 4 stl     1990 Apr    13012. 13231.     -220.      1.05     13232.
## 5 stl     1990 May    13108. 13211.     -114.      11.3     13223.
## 6 stl     1990 Jun    13183. 13192.     -24.3      15.5     13207.
## 7 stl     1990 Jul    13170. 13172.     -23.2      21.6     13193.26
## # ... with 350 more rows, and 1 more variable:
## #   season_adjust <dbl>
```

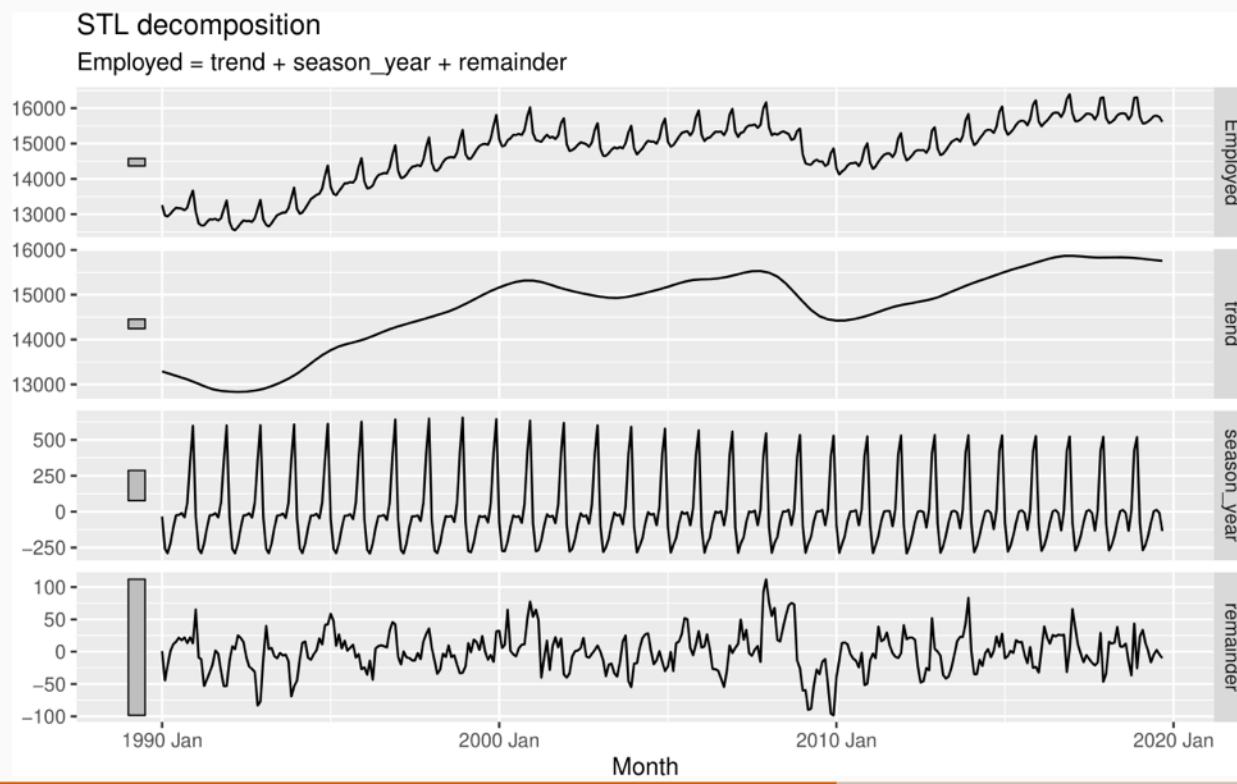
# US Retail Employment

```
us_retail_employment %>%
  autoplot(Employed, color='gray') +
  autolayer(components(dcmp), trend, color="#D55E00") +
  labs(y="Persons (thousands)", title="Total employment in US retail")
```



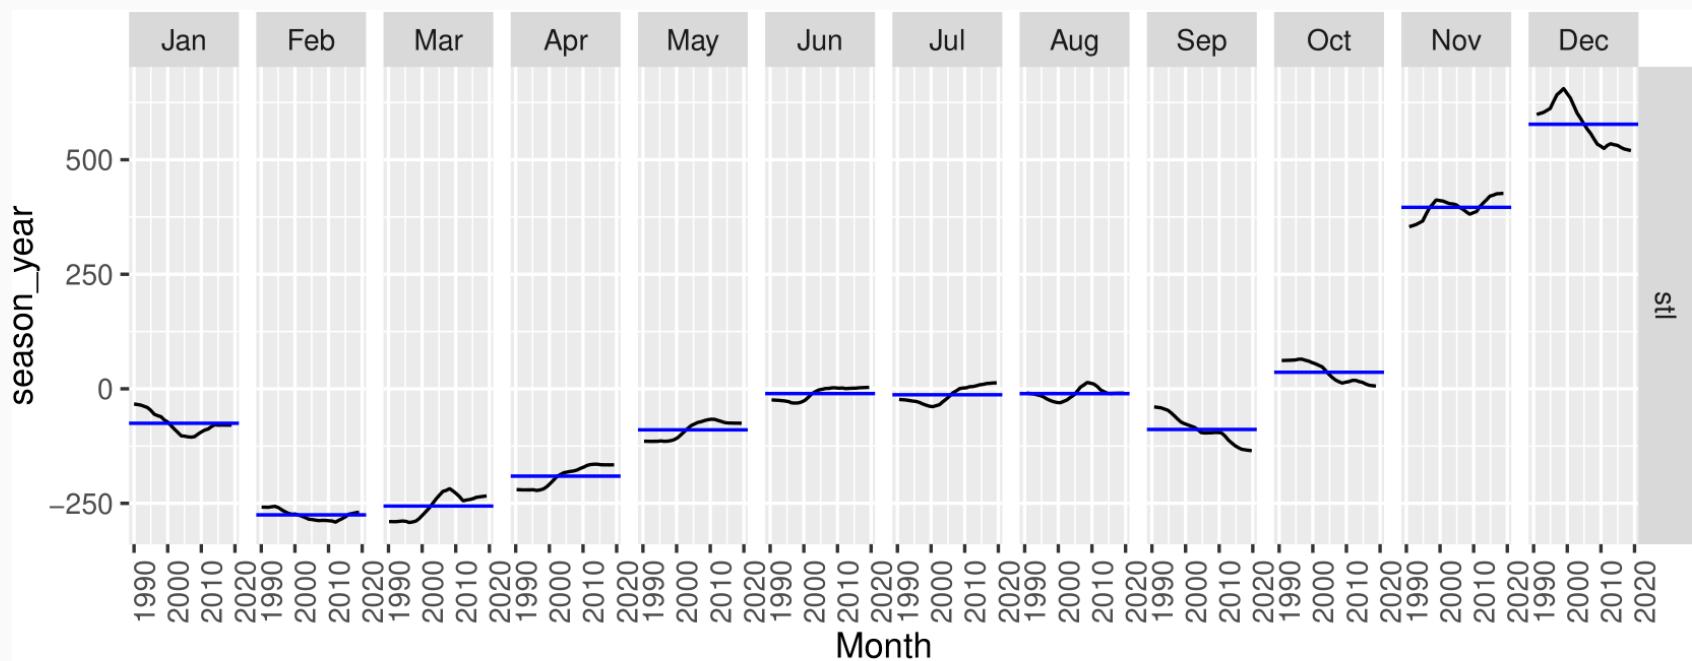
# US Retail Employment

```
components(dcmp) %>% autoplot()
```



# US Retail Employment

```
components(dcmp) %>% gg_subseries(season_year)
```



# Seasonal adjustment

- Useful by-product of decomposition: an easy way to calculate seasonally adjusted data.
- Additive decomposition: seasonally adjusted data given by

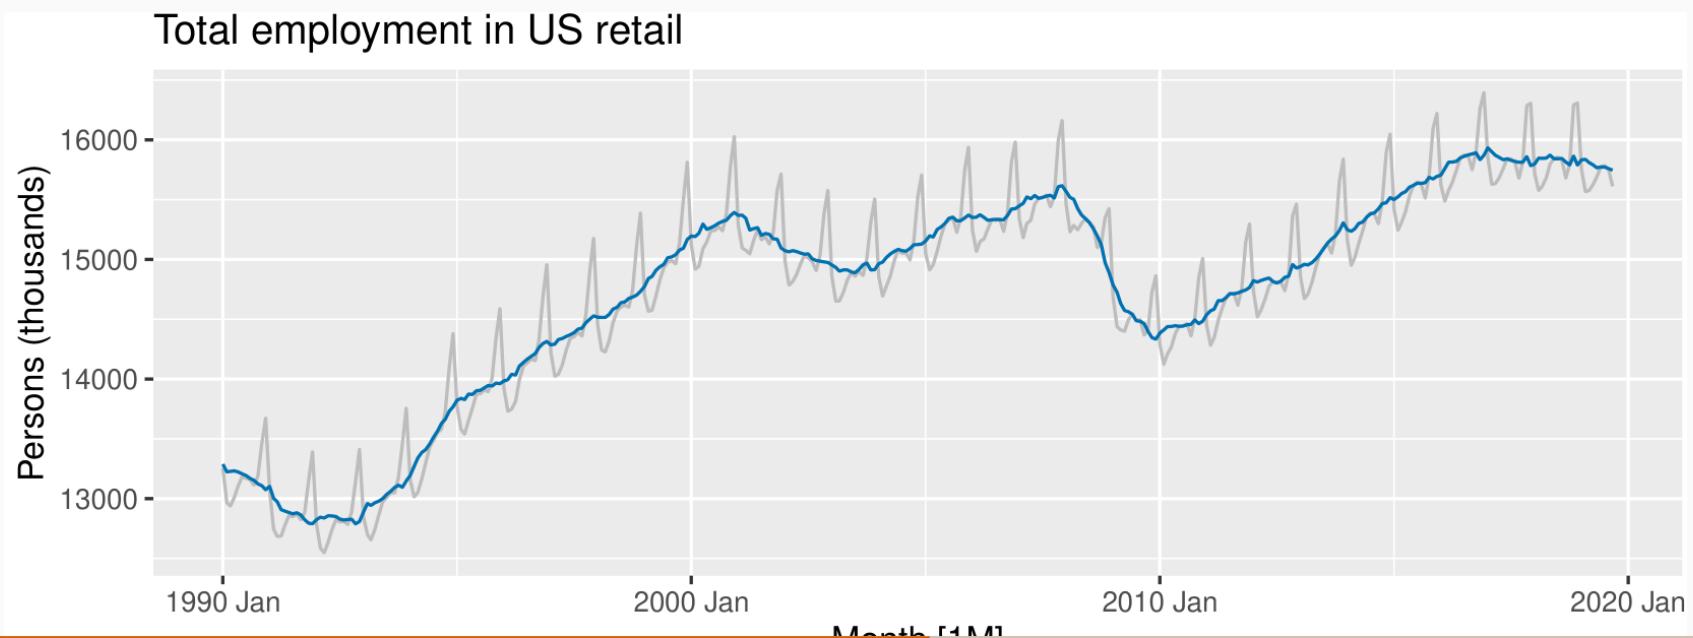
$$y_t - S_t = T_t + R_t$$

- Multiplicative decomposition: seasonally adjusted data given by

$$y_t / S_t = T_t \times R_t$$

# US Retail Employment

```
us_retail_employment %>%
  autoplot(Employed, color='gray') +
  autolayer(components(dcmp), season_adjust, color='#0072B2') +
  labs(y="Persons (thousands)", title="Total employment in US retail")
```



# Seasonal adjustment

- We use estimates of  $S$  based on past values to seasonally adjust a current value.
- Seasonally adjusted series reflect **remainders** as well as **trend**. Therefore they are not “smooth” and “downturns” or “upturns” can be misleading.
- It is better to use the trend-cycle component to look for turning points.

# Outline

- 1 Transformations and adjustments
- 2 Time series components
- 3 History of time series decomposition
- 4 STL decomposition
- 5 When things go wrong

# History of time series decomposition

- Classical method originated in 1920s.
- Census II method introduced in 1957. Basis for X-11 method and variants (including X-12-ARIMA, X-13-ARIMA)
- STL method introduced in 1983
- TRAMO/SEATS introduced in 1990s.

## National Statistics Offices

- ABS uses X-12-ARIMA
- US Census Bureau uses X-13ARIMA-SEATS
- Statistics Canada uses X-12-ARIMA
- ONS (UK) uses X-12-ARIMA
- EuroStat use X-13ARIMA-SEATS

# X-11 decomposition

## Advantages

- Relatively robust to outliers
- Completely automated choices for trend and seasonal changes
- Very widely tested on economic data over a long period of time.

## Disadvantages

- No prediction/confidence intervals
- Ad hoc method with no underlying model
- Only developed for quarterly and monthly data

## Extensions: X-12-ARIMA and X-13-ARIMA

- The X-11, X-12-ARIMA and X-13-ARIMA methods are based on Census II decomposition.
- These allow adjustments for trading days and other explanatory variables.
- Known outliers can be omitted.
- Level shifts and ramp effects can be modelled.
- Missing values estimated and replaced.
- Holiday factors (e.g., Easter, Labour Day) can be estimated.

# X-13ARIMA-SEATS

## Advantages

- Model-based
- Smooth trend estimate
- Allows estimates at end points
- Allows changing seasonality
- Developed for economic data

## Disadvantages

- Only developed for quarterly and monthly data

# Outline

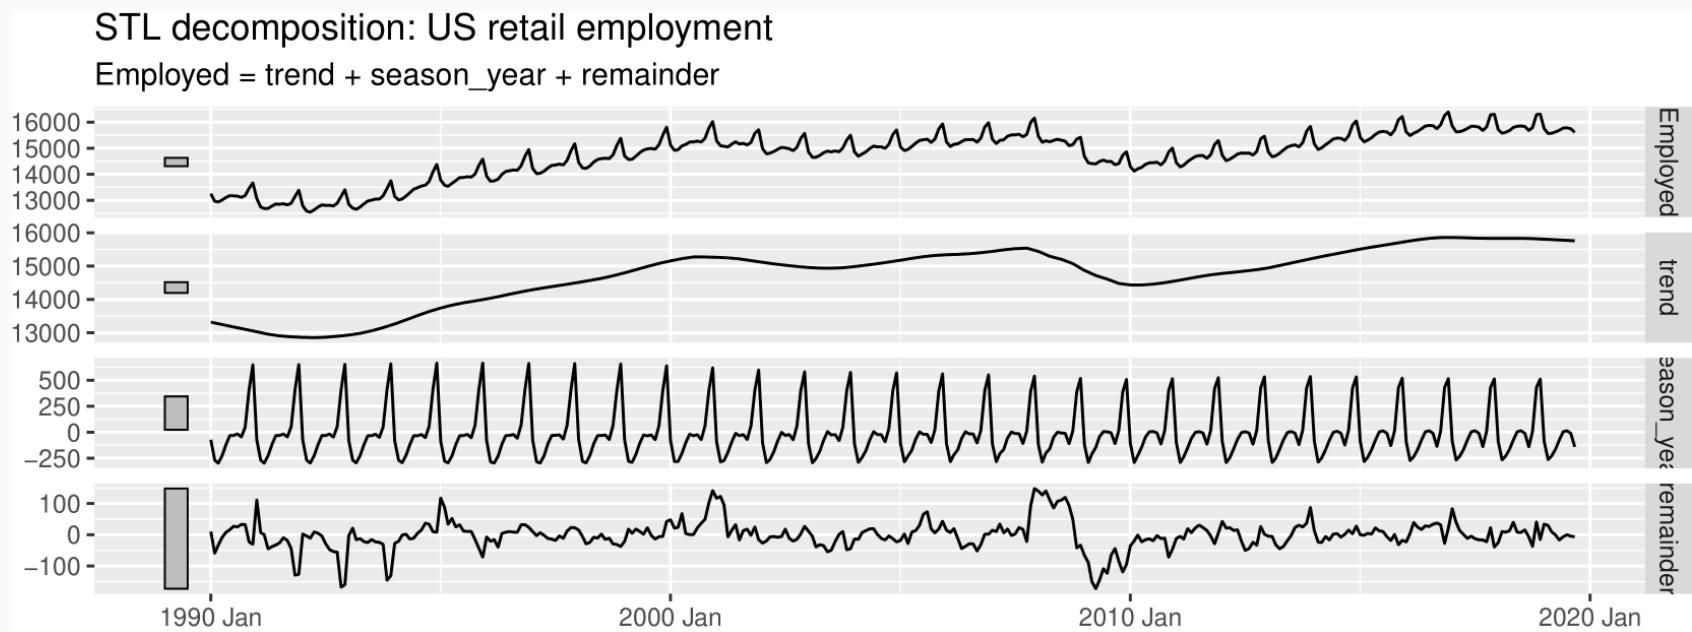
- 1 Transformations and adjustments
- 2 Time series components
- 3 History of time series decomposition
- 4 STL decomposition
- 5 When things go wrong

# STL decomposition

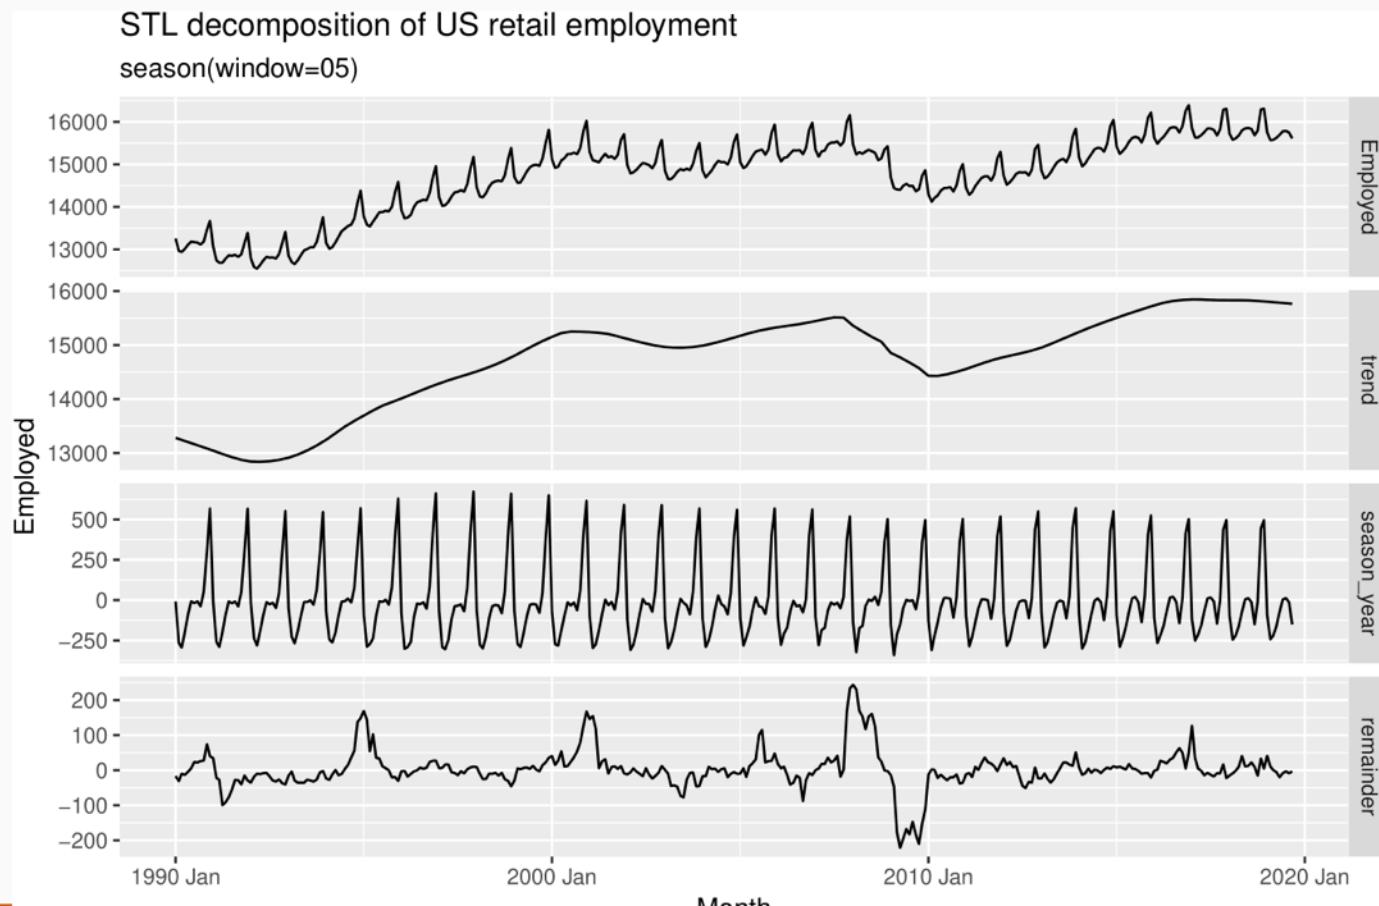
- STL: “Seasonal and Trend decomposition using Loess”
- Very versatile and robust.
- Unlike X-12-ARIMA, STL will handle any type of seasonality.
- Seasonal component allowed to change over time, and rate of change controlled by user.
- Smoothness of trend-cycle also controlled by user.
- Robust to outliers
- Not trading day or calendar adjustments.
- Only additive.
- Take logs to get multiplicative decomposition.
- Use Box-Cox transformations to get other decompositions.

# STL decomposition

```
us_retail_employment %>%
  model(STL(Employed ~ season(window=9), robust=TRUE)) %>%
  components() %>% autoplot() +
  labs(title = "STL decomposition: US retail employment")
```



# STL decomposition



# STL decomposition

```
us_retail_employment %>%
  model(STL(Employed ~ season(window=5))) %>%
  components()

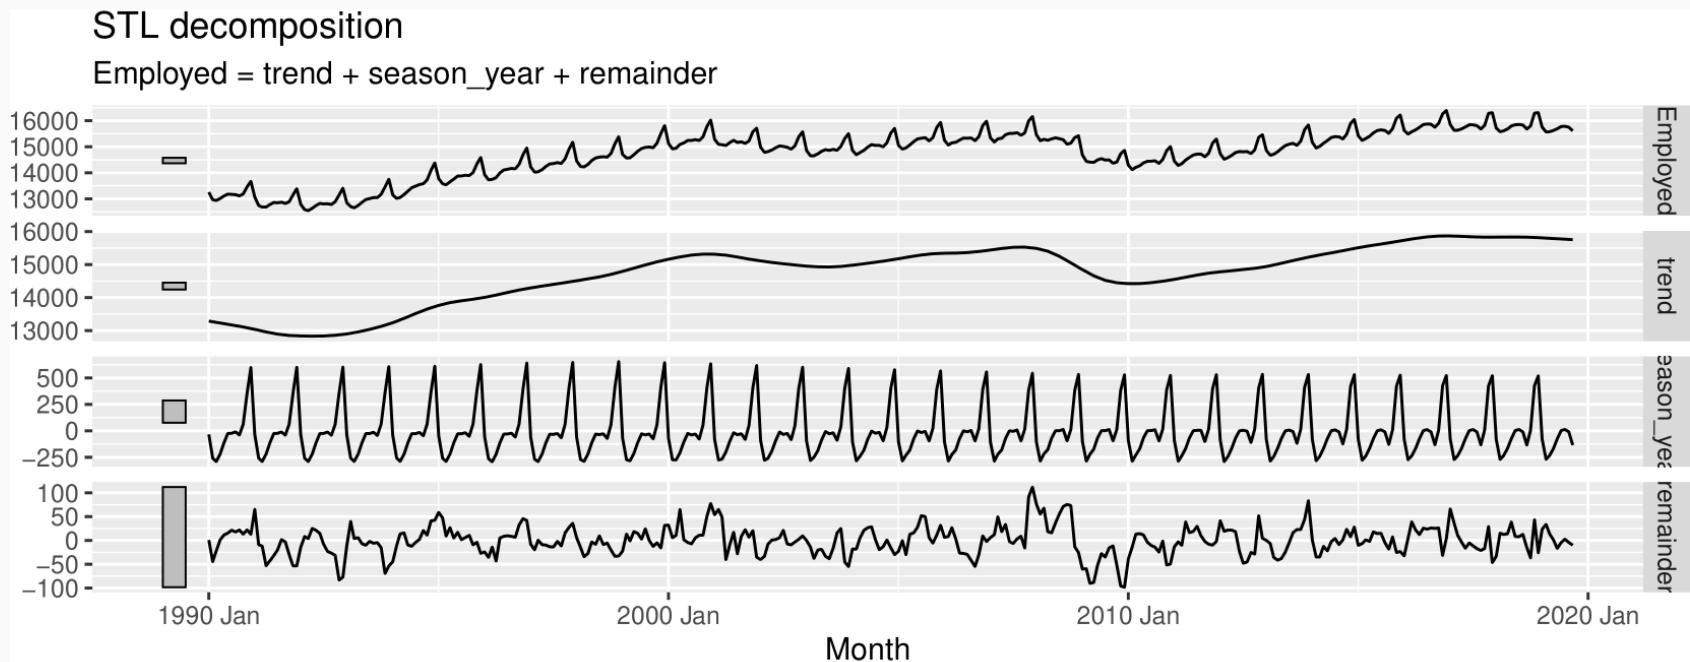
us_retail_employment %>%
  model(STL(Employed ~ trend(window=15) +
              season(window="periodic"),
              robust = TRUE)
  ) %>% components()
```

- `trend(window = ?)` controls wiggliness of trend component.
- `season(window = ?)` controls variation on seasonal component.
- `season(window = 'periodic')` is equivalent to an infinite window.

# STL decomposition

- `STL()` chooses season(`window=13`) by default
- Can include transformations.

```
us_retail_employment %>%  
  model(STL(Employed)) %>%  
  components() %>% autoplot()
```



# STL decomposition

- Algorithm that updates trend and seasonal components iteratively.
- Starts with  $\hat{T}_t = 0$
- Uses a mixture of loess and moving averages to successively refine the trend and seasonal estimates.
- The trend window controls loess bandwidth applied to deasonalised values.
- The season window controls loess bandwidth applied to detrended subseries.
- Robustness weights based on remainder.
- Default season window = 13
- Default trend window = `nextodd(ceiling((1.5*period)/(1-(1.5/s.window))))`

# Outline

- 1 Transformations and adjustments
- 2 Time series components
- 3 History of time series decomposition
- 4 STL decomposition
- 5 When things go wrong

# The ABS stuff-up

NEWS 

LOCATION: Clayton, Vic [Change](#)

[Home](#) Just In Australia World Business Sport Analysis & Opinion Fact Check Programs

**BREAKING NEWS** Police arrest man in connection with stabbing death of 17-year-old Masa Vukotic in M

[Print](#) [Email](#) [Facebook](#) [Twitter](#) [More](#)

## Treasurer Joe Hockey calls for answers over Australian Bureau of Statistics jobs data

By [Michael Vincent](#) and [Simon Frazer](#)  
Updated 9 Oct 2014, 12:17pm

**Federal Treasurer Joe Hockey says he wants answers to the problems the Australian Bureau of Statistics (ABS) has had with unemployment figures.**

Mr Hockey, who is in the US to discuss Australia's G20 agenda, said last month's unemployment figures were "extraordinary".

The rate was 6.1 per cent after jumping to a 12-year high of 6.4 per cent the previous month.

The ABS has now taken the rare step of abandoning seasonal adjustment for its latest employment data.



**PHOTO:** Joe Hockey says he is unhappy with the volatility of ABS unemployment figures. (AAP: Alan Porritt)

**RELATED STORY:** ABS abandons seasonal adjustment for latest jobs data

# The ABS stuff-up

The screenshot shows the ABC News homepage with a blue header. In the top left is a 'NEWS' button with a television icon. To its right is a 'LOCATION' dropdown set to 'Clayton, Vic' with a 'Change' link. On the far right is a search bar with a magnifying glass icon. Below the header is a navigation menu with links: 'Just In', 'Australia', 'World', 'Business', 'Sport', 'Analysis & Opinion', 'Fact Check', 'Programs', and 'More'. A yellow 'BREAKING NEWS' box is visible, followed by a headline: 'Police arrest man in connection with stabbing death of 17-year-old Masa Vukotic in Mel'. Below the headline are sharing options: 'Print', 'Email', 'Facebook', 'Twitter', and 'More'. The main article title is 'ABS abandons seasonal adjustment for latest jobs data' by Michael Janda, updated on 8 Oct 2014, 4:19pm. The text discusses the ABS's decision to abandon seasonal adjustment for employment data. It includes quotes from the ABS and Westpac chief economist Bill Evans. There are also related stories and a map of Australia.

## ABC abandons seasonal adjustment for latest jobs data

By business reporter Michael Janda  
Updated 8 Oct 2014, 4:19pm

**The Australian Bureau of Statistics is taking the rare step of abandoning seasonal adjustment for its latest employment data.**

The ABS uses seasonal adjustment, based on historical experience, to account for the normal variation between hiring and firing patterns between different months.

However, after a winter where the seasonally adjusted unemployment rate swung wildly from 6.1 to 6.4 and back to 6.1 per cent, [the bureau released a statement](#) saying it will not adjust the original figure for September for seasonal factors.

It will also reset the seasonal adjustment for July and August to one, meaning that these months will also reflect the original figures.

**VIDEO:** Westpac chief economist Bill Evans discusses the ABS jobs data changes (ABC News)

**RELATED STORY:** Doubt the record breaking jobs figures? So does the ABS

**RELATED STORY:** Jobs increase record sees unemployment slashed

**RELATED STORY:** Unemployment surges to 12-year high at 6.4 pc

**MAP:** Australia

# The ABS stuff-up

## ABS jobs and unemployment figures - key questions answered by an expert

A professor of statistics at Monash University explains exactly what is seasonal adjustment, why it matters and what went wrong in the July and August figures



School leavers come on to the jobs market at the same time, causing a seasonal fluctuation. Photograph: Brian Snyder/Reuters

The Australian Bureau of Statistics has retracted its seasonally adjusted employment data for July and August, which recorded huge swings in the jobless rate. The ABS is also planning to review the methods it uses for seasonal adjustment to ensure its figures are as accurate as possible. Rob Hyndman, a professor of statistics at Monash University and member of the bureau's methodology advisory board, answers our questions:

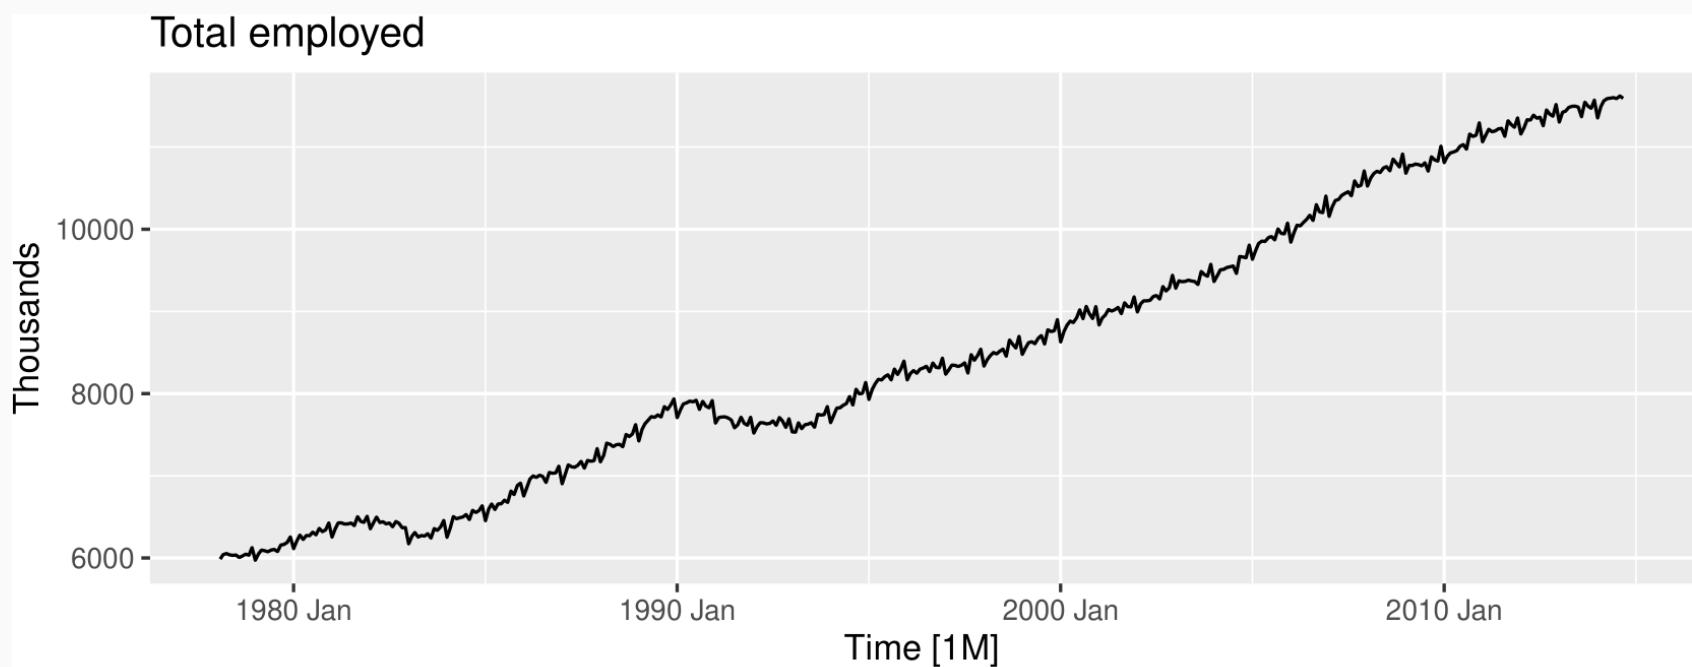
# The ABS stuff-up

employed

```
## # A tsibble: 440 x 4 [1M]
##       Time Month Year Employed
##       <mth> <ord> <dbl>    <dbl>
## 1 1978 Feb   Feb    1978    5986.
## 2 1978 Mar   Mar    1978    6041.
## 3 1978 Apr   Apr    1978    6054.
## 4 1978 May   May    1978    6038.
## 5 1978 Jun   Jun    1978    6031.
## 6 1978 Jul   Jul    1978    6036.
## 7 1978 Aug   Aug    1978    6005.
## 8 1978 Sep   Sep    1978    6024.
## 9 1978 Oct   Oct    1978    6046.
## 10 1978 Nov   Nov   1978    6034.
## # ... with 430 more rows
```

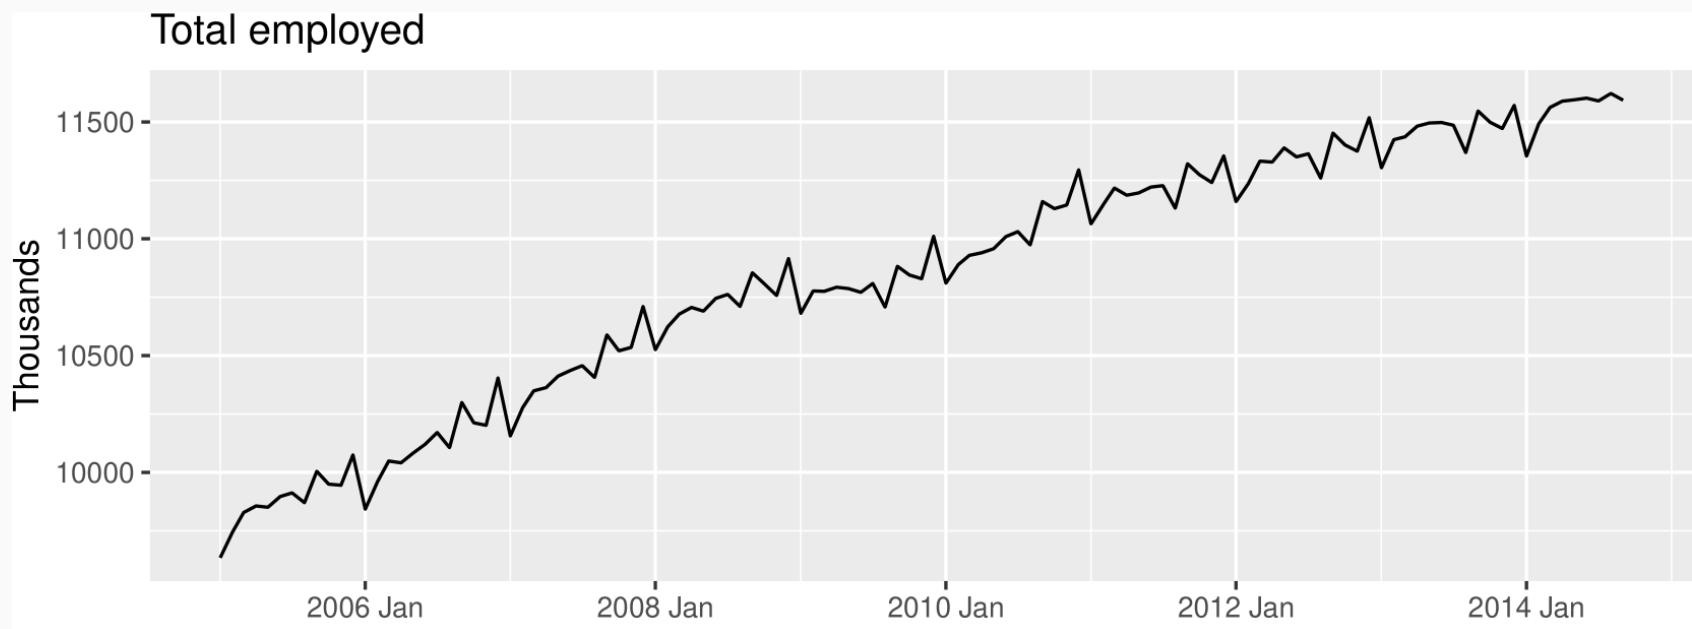
# The ABS stuff-up

```
employed %>%
  autoplot(Employed) +
  labs(title = "Total employed", y = "Thousands")
```



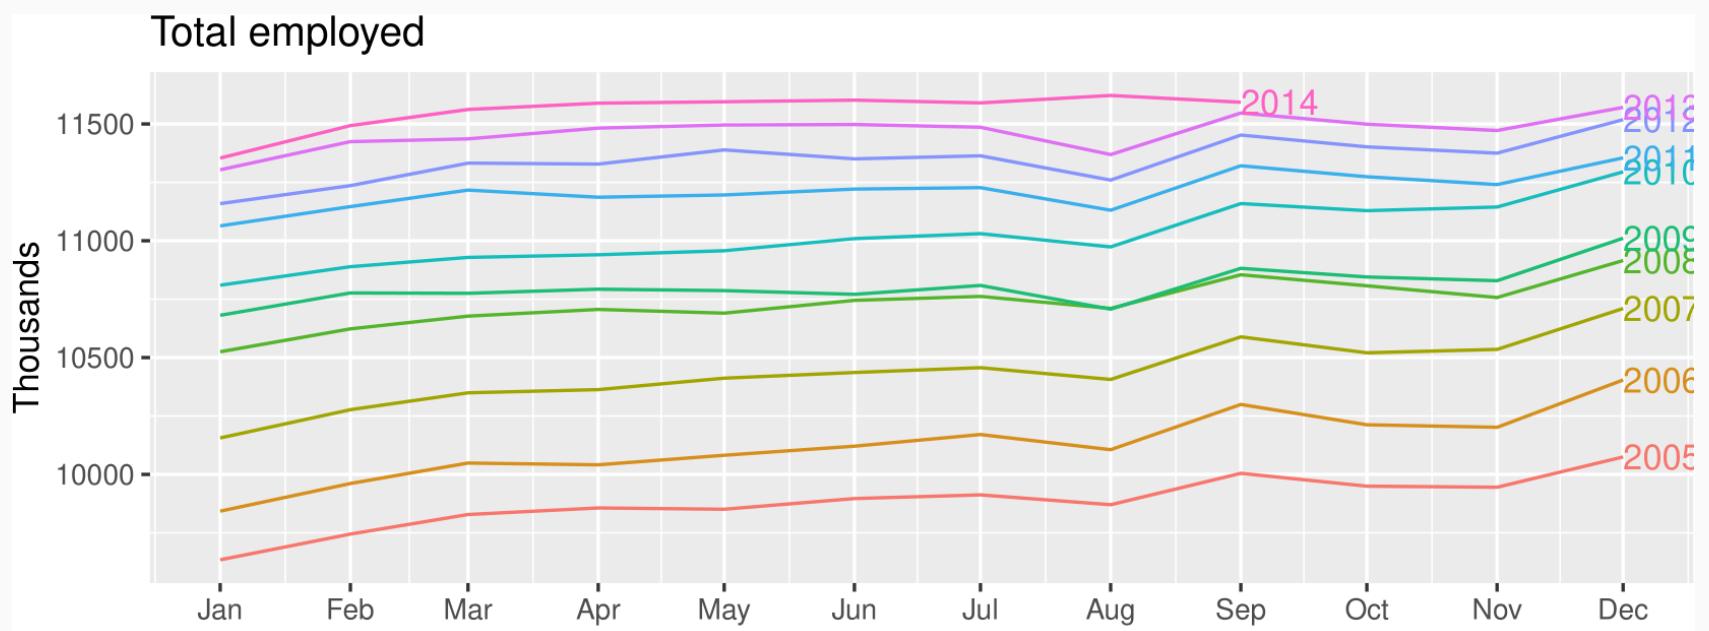
# The ABS stuff-up

```
employed %>%
  filter(Year >= 2005) %>%
  autoplot(Employed) +
  labs(title = "Total employed", y = "Thousands")
```



# The ABS stuff-up

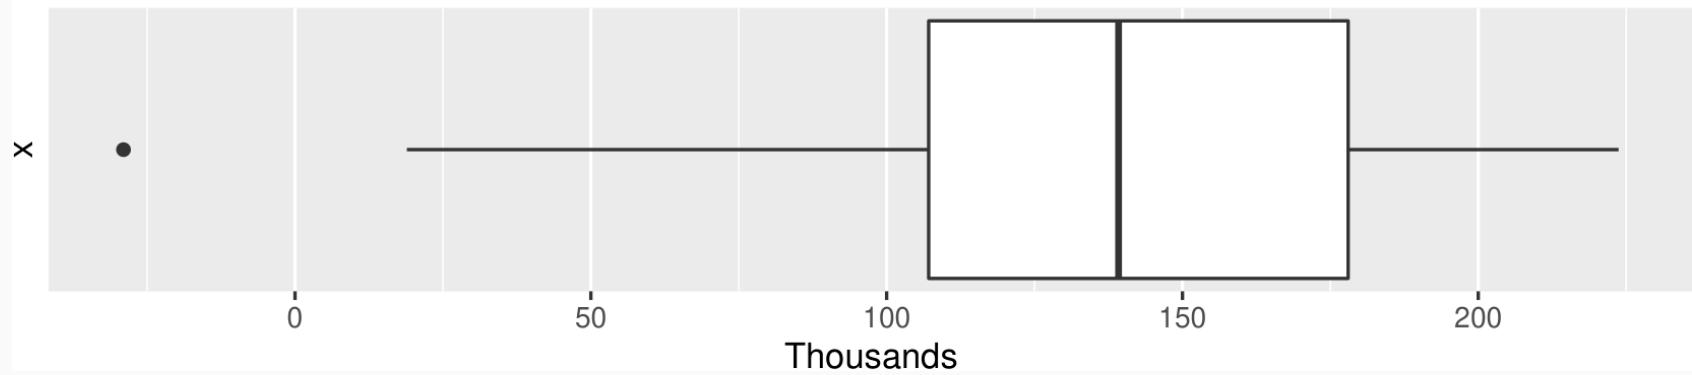
```
employed %>%
  filter(Year >= 2005) %>%
  gg_season(Employed, label = "right") +
  labs(title = "Total employed", y = "Thousands")
```



# The ABS stuff-up

```
employed %>%
  mutate(diff = difference(Employed)) %>%
  filter(Month == "Sep") %>%
  ggplot(aes(y = diff, x = 1)) +
  geom_boxplot() + coord_flip() +
  labs(title = "Sep - Aug: total employed", y = "Thousands") +
  scale_x_continuous(breaks = NULL, labels = NULL)
```

Sep – Aug: total employed

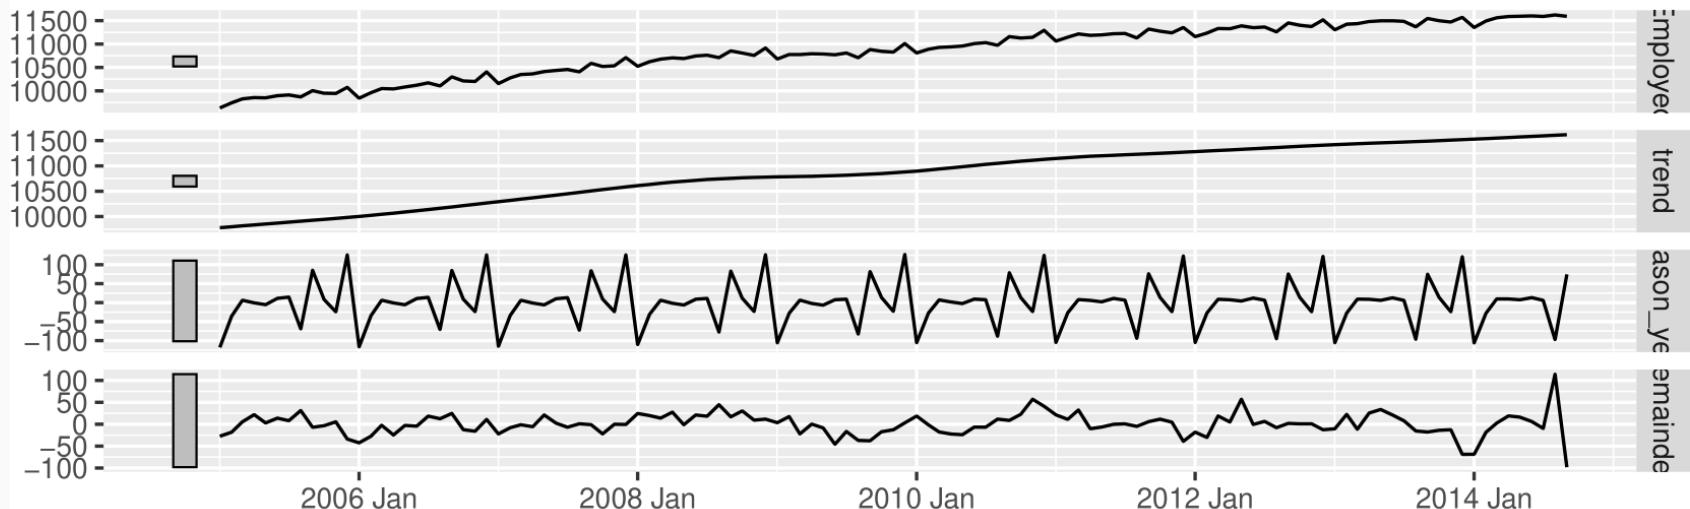


# The ABS stuff-up

```
dcmp <- employed %>%
  filter(Year >= 2005) %>%
  model(stl = STL(Employed ~ season(window = 11), robust = TRUE))
components(dcmp) %>% autoplot()
```

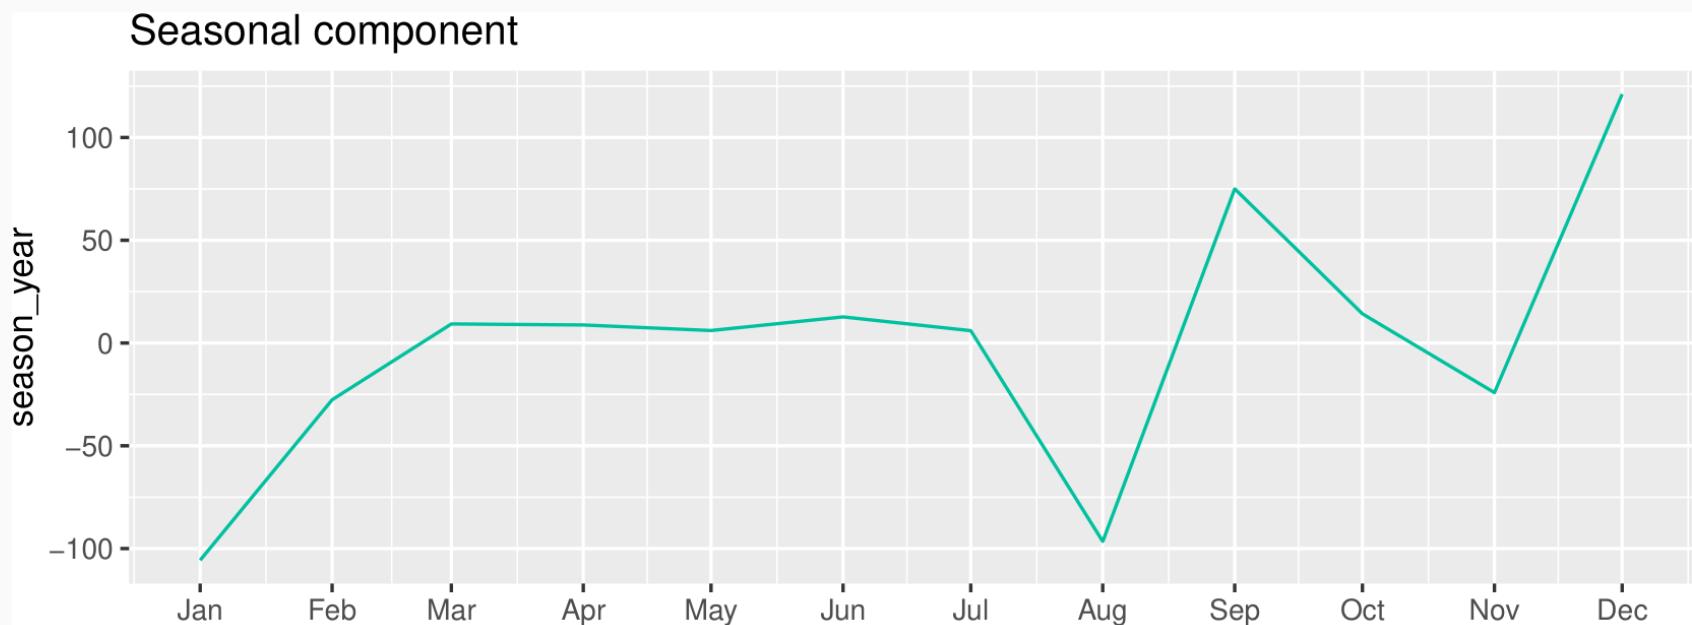
## STL decomposition

Employed = trend + season\_year + remainder



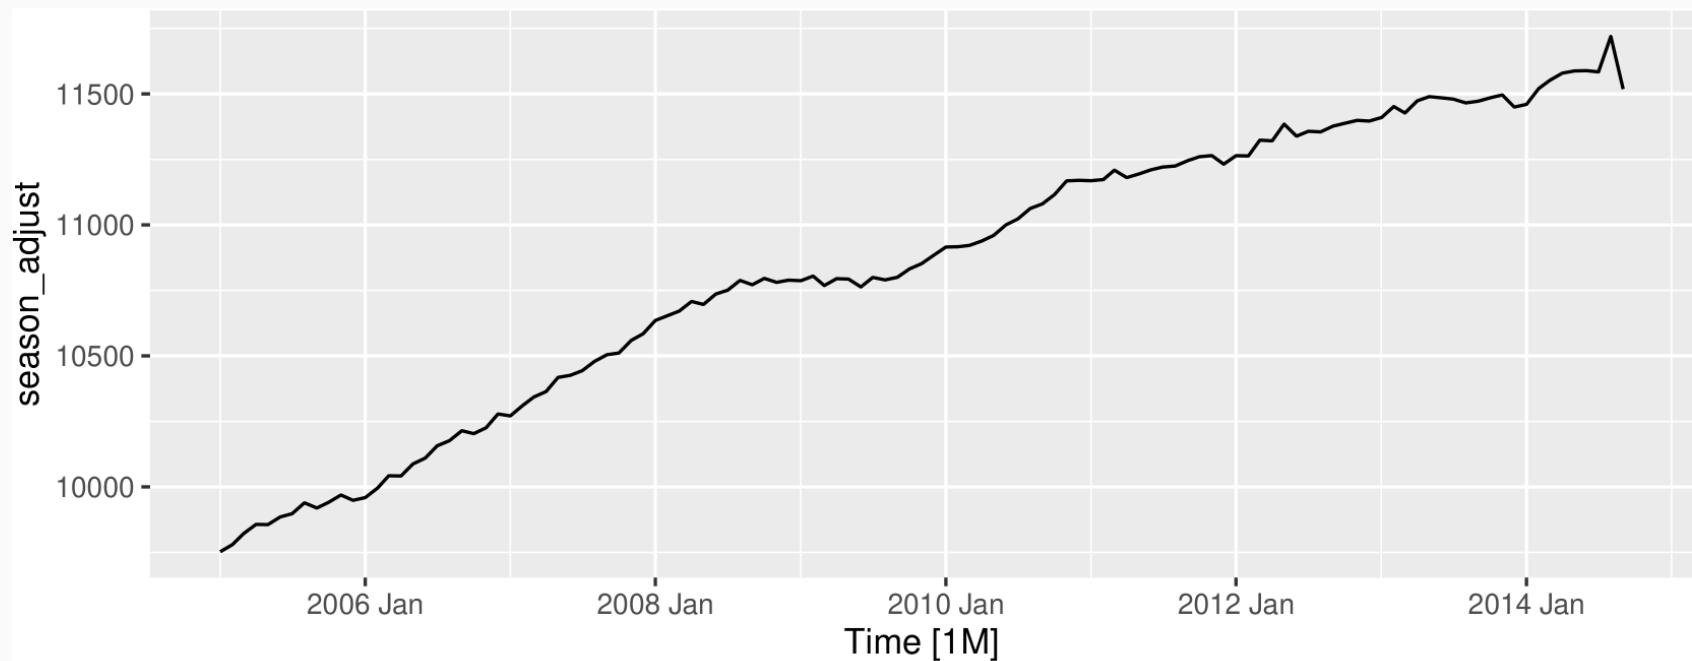
# The ABS stuff-up

```
components(dcmp) %>%
  filter(year(Time) == 2013) %>%
  gg_season(season_year) +
  labs(title = "Seasonal component") + guides(colour = "none")
```



# The ABS stuff-up

```
components(dcmp) %>%
  as_tsibble() %>%
  autoplot(season_adjust)
```



# The ABS stuff-up

- August 2014 employment numbers higher than expected.
- Supplementary survey usually conducted in August for employed people.
- Most likely, some employed people were claiming to be unemployed in August to avoid supplementary questions.
- Supplementary survey not run in 2014, so no motivation to lie about employment.
- In previous years, seasonal adjustment fixed the problem.
- The ABS has now adopted a new method to avoid the bias.

# Outline

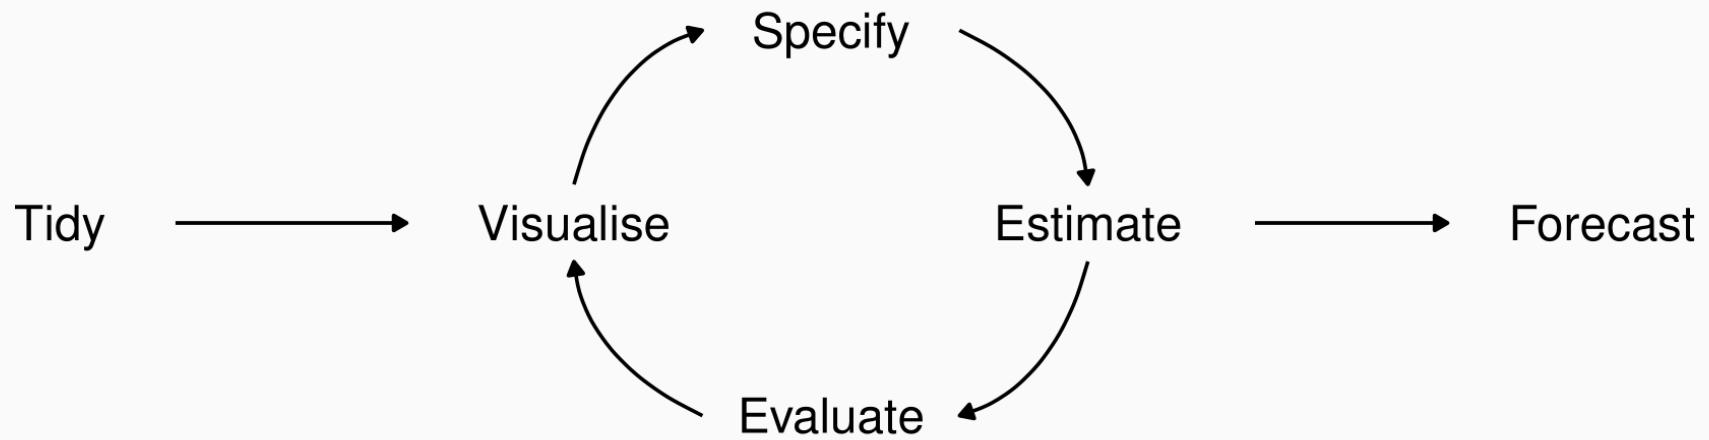
- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# A tidy forecasting workflow

The process of producing forecasts can be split up into a few fundamental steps.

- 1 Preparing data
- 2 Data visualisation
- 3 Specifying a model
- 4 Model estimation
- 5 Accuracy & performance evaluation
- 6 Producing forecasts

# A tidy forecasting workflow



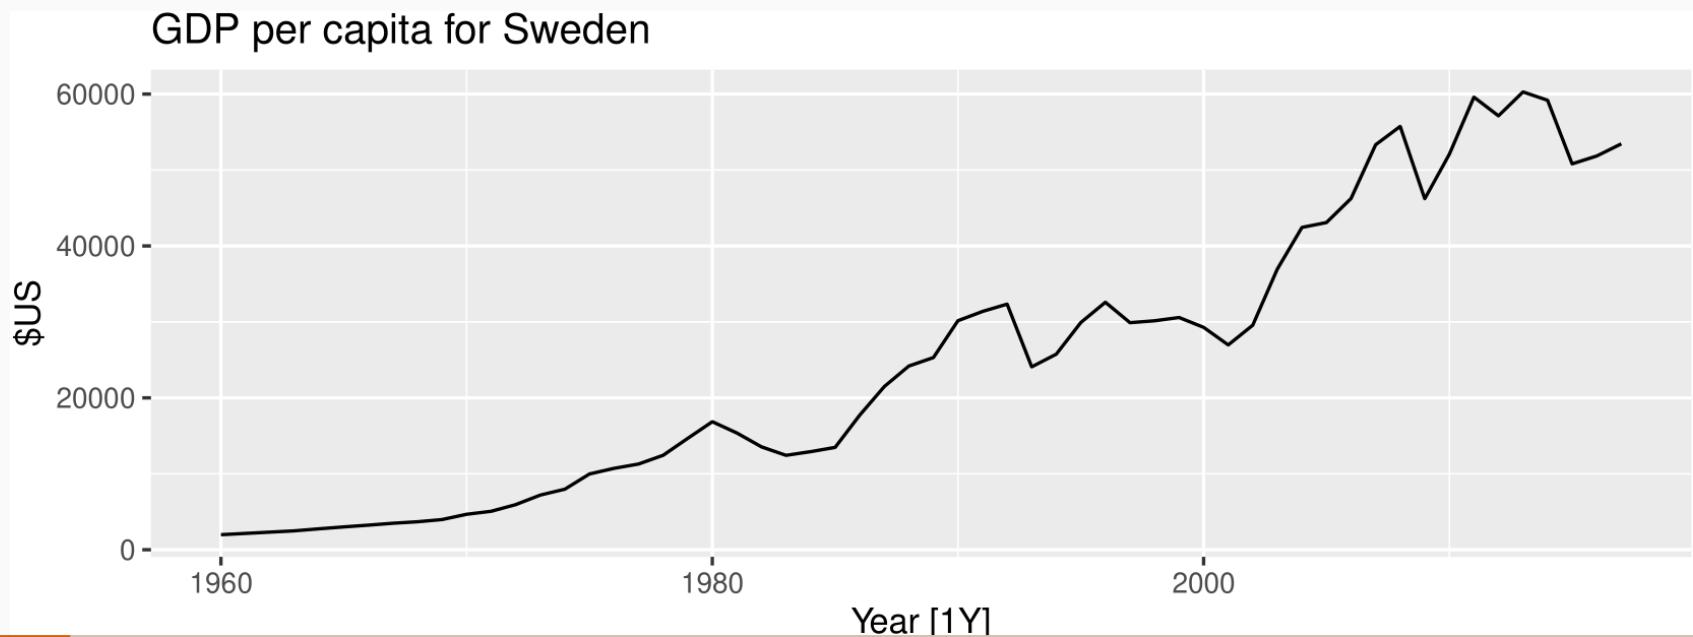
# Data preparation (tidy)

```
gdppc <- global_economy %>%
  mutate(GDP_per_capita = GDP/Population) %>%
  select(Year, Country, GDP, Population, GDP_per_capita)
gdppc
```

```
## # A tsibble: 15,150 x 5 [1Y]
## # Key:     Country [263]
## #       Year Country          GDP Population GDP_per_capita
## #   <dbl> <fct>    <dbl>      <dbl>        <dbl>
## 1 1960 Afghanistan 537777811.    8996351      59.8
## 2 1961 Afghanistan 548888896.    9166764      59.9
## 3 1962 Afghanistan 546666678.    9345868      58.5
## 4 1963 Afghanistan 751111191.    9533954      78.8
## 5 1964 Afghanistan 800000044.    9731361      82.2
## 6 1965 Afghanistan 1006666638.   9938414      101.
```

# Data visualisation

```
gdppc %>%
  filter(Country=="Sweden") %>%
  autoplot(GDP_per_capita) +
  labs(title = "GDP per capita for Sweden", y = "$US")
```



# Model estimation

The `model()` function trains models to data.

```
fit <- gdppc %>%
  model(trend_model = TSLM(GDP_per_capita ~ trend()))
fit
```

```
## # A mable: 263 x 2
## # Key:     Country [263]
##   Country            trend_model
##   <fct>              <model>
## 1 Afghanistan        <TSLM>
## 2 Albania            <TSLM>
## 3 Algeria            <TSLM>
## 4 American Samoa    <TSLM>
## 5 Andorra            <TSLM>
```

A `mable` is a model table, each cell corresponds to a fitted model.

# Producing forecasts

```
fit %>% forecast(h = "3 years")  
  
## # A fable: 789 x 5 [1Y]  
## # Key:     Country, .model [263]  
##   Country      .model     Year GDP_per_capita .mean  
##   <fct>        <chr>      <dbl>    <dist>     <dbl>  
## 1 Afghanistan trend_model 2018 N(526, 9653) 526.  
## 2 Afghanistan trend_model 2019 N(534, 9689) 534.  
## 3 Afghanistan trend_model 2020 N(542, 9727) 542.  
## 4 Albania      trend_model 2018 N(4716, 476419) 4716.  
## 5 Albania      trend_model 2019 N(4867, 481086) 4867.  
## 6 Albania      trend_model 2020 N(5018, 486012) 5018.  
## 7 Algeria      trend_model 2018 N(4410, 643094) 4410.  
## 8 Algeria      trend_model 2019 N(4489, 645311) 4489.
```

# Producing forecasts

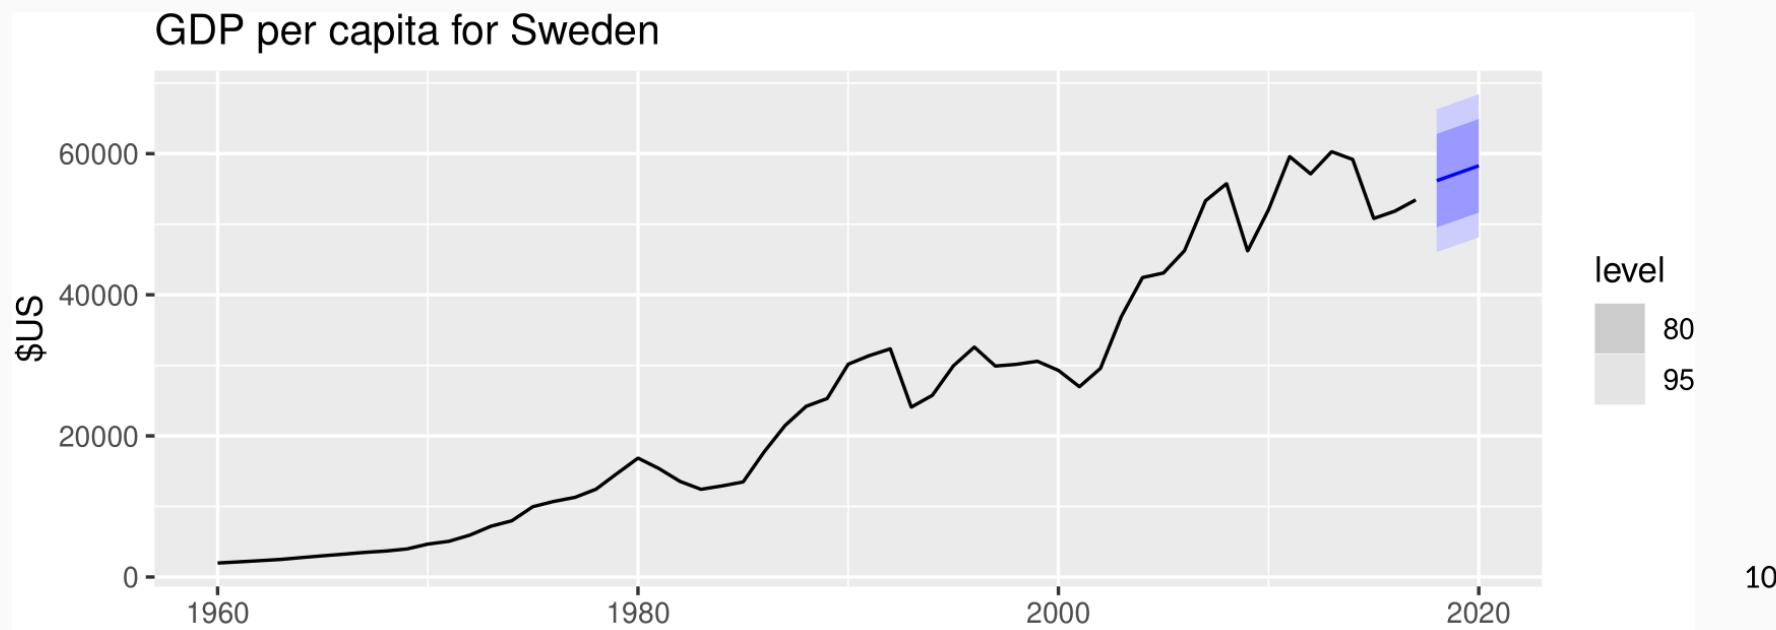
```
fit %>% forecast(h = "3 years")
```

```
## # A fable: 789 x 5 [1Y]
## # Key:     Country, .model [263]
##   Country      .model     Year GDP_per_capita .mean
##   <fct>        <chr>      <dbl>    <dist>      <dbl>
## 1 Afghanistan trend_model 2018   N(526, 9653)  526.
## 2 Afghanistan trend_model 2019   N(534, 9689)  534.
## 3 Afghanistan trend_model 2020   N(542, 9727)  542.
## 4 Albania      trend_model 2018   N(4716, 476419) 4716.
## 5 Albania      trend_model 2019   N(4867, 481086) 4867.
## 6 Albania      trend_model 2020   N(5018, 486012) 5018.
## 7 Algeria      trend_model 2018   N(4410, 643094) 4410.
## 8 Algeria      trend_model 2019   N(4489, 645311) 4489.
```

A fable is a forecast table with point forecasts and distributions.

# Visualising forecasts

```
fit %>% forecast(h = "3 years") %>%
  filter(Country=="Sweden") %>%
  autoplot(gdppc) +
  labs(title = "GDP per capita for Sweden", y = "$US")
```



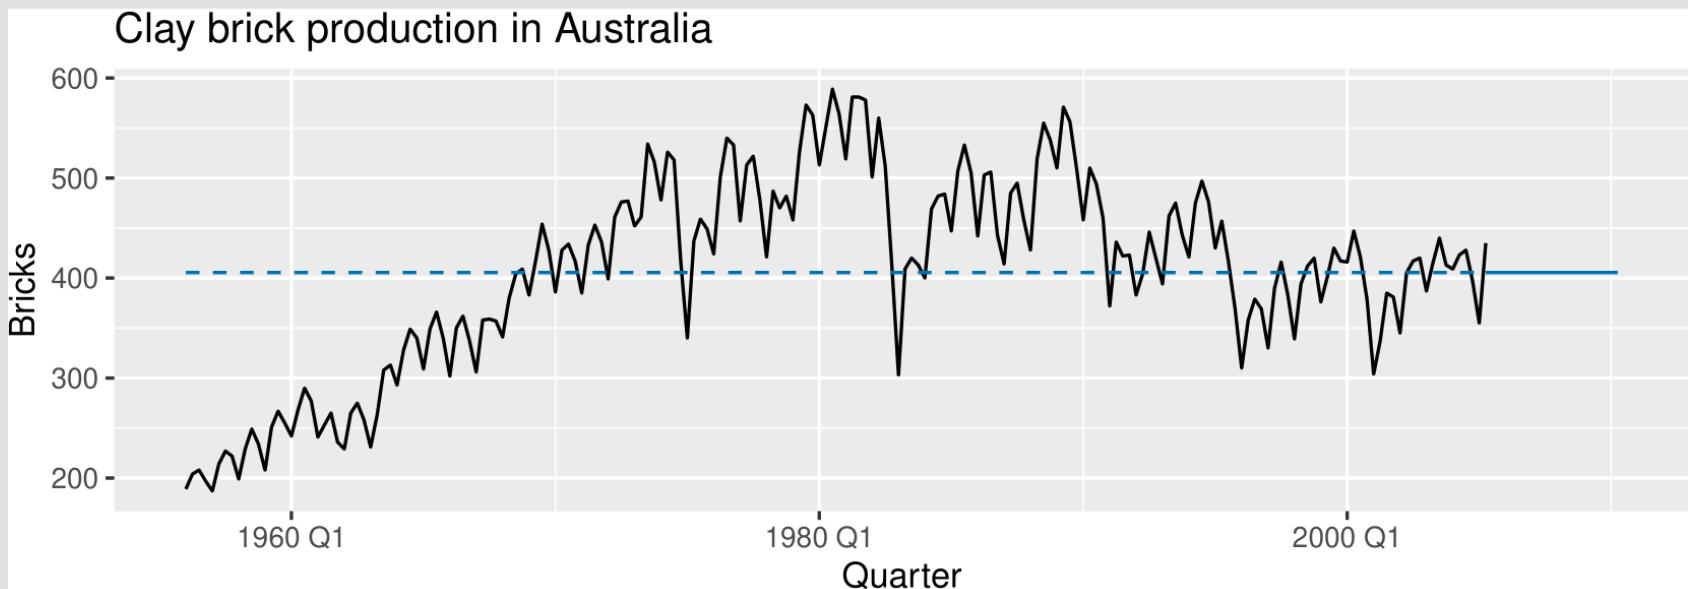
# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Some simple forecasting methods

## MEAN( $y$ ): Average method

- Forecast of all future values is equal to mean of historical data  $\{y_1, \dots, y_T\}$ .
- Forecasts:  $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$

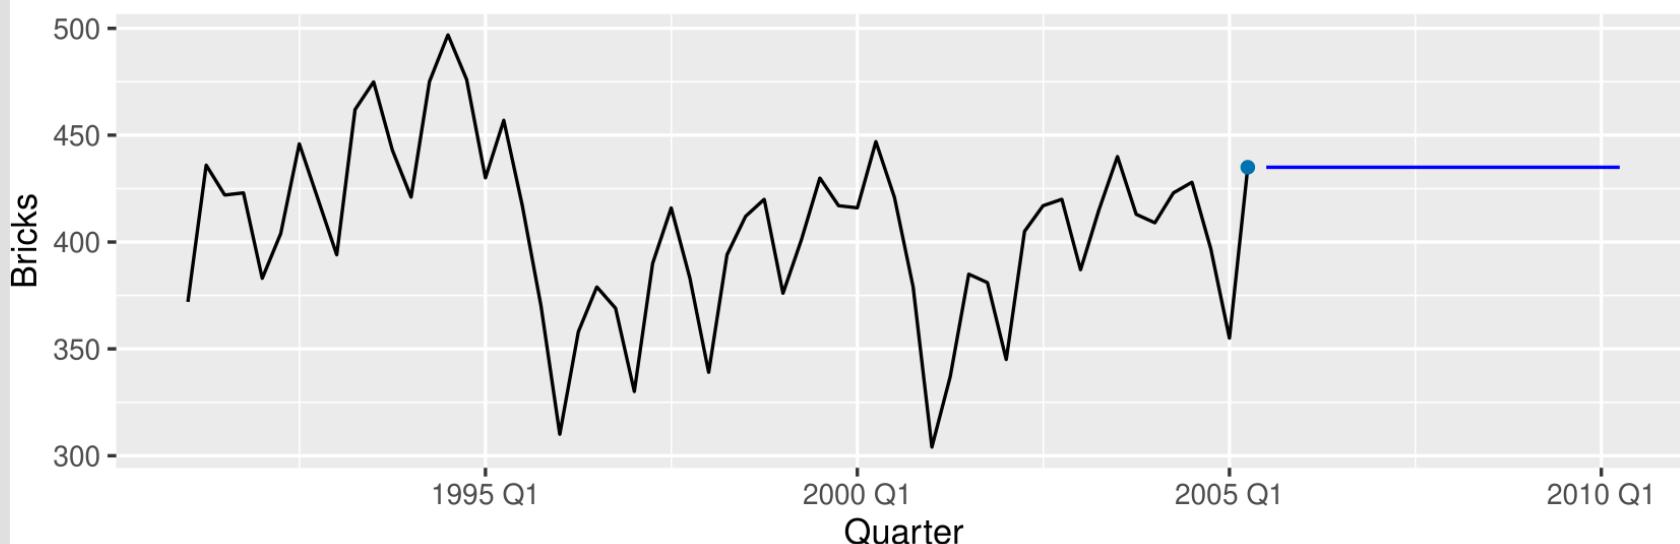


# Some simple forecasting methods

## NAIVE( $y$ ): Naïve method

- Forecasts equal to last observed value.
- Forecasts:  $\hat{y}_{T+h|T} = y_T$ .
- Consequence of efficient market hypothesis.

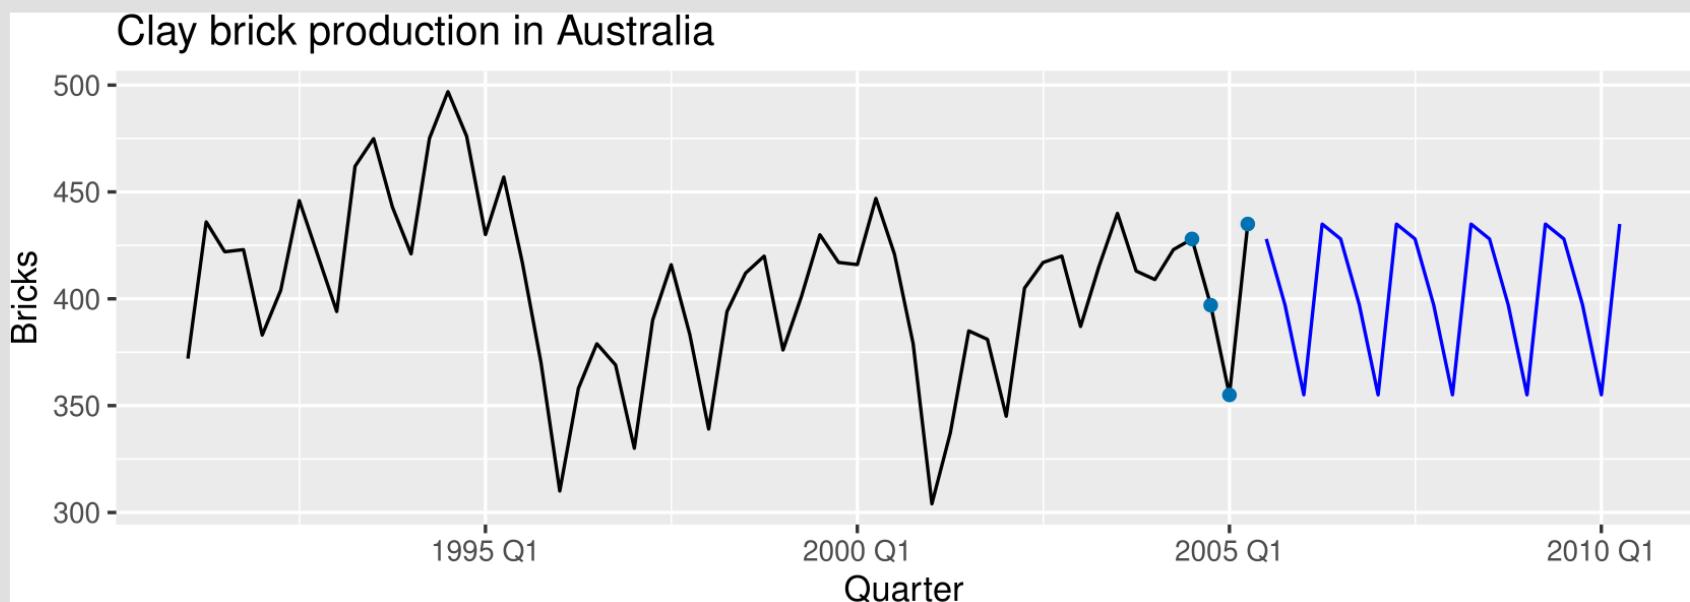
Clay brick production in Australia



# Some simple forecasting methods

## SNAIVE( $y \sim \text{lag}(m)$ ): Seasonal naïve method

- Forecasts equal to last value from same season.
- Forecasts:  $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$ , where  $m$  = seasonal period and  $k$  is the integer part of  $(h - 1)/m$ .



# Some simple forecasting methods

## Rw(y ~ drift()): Drift method

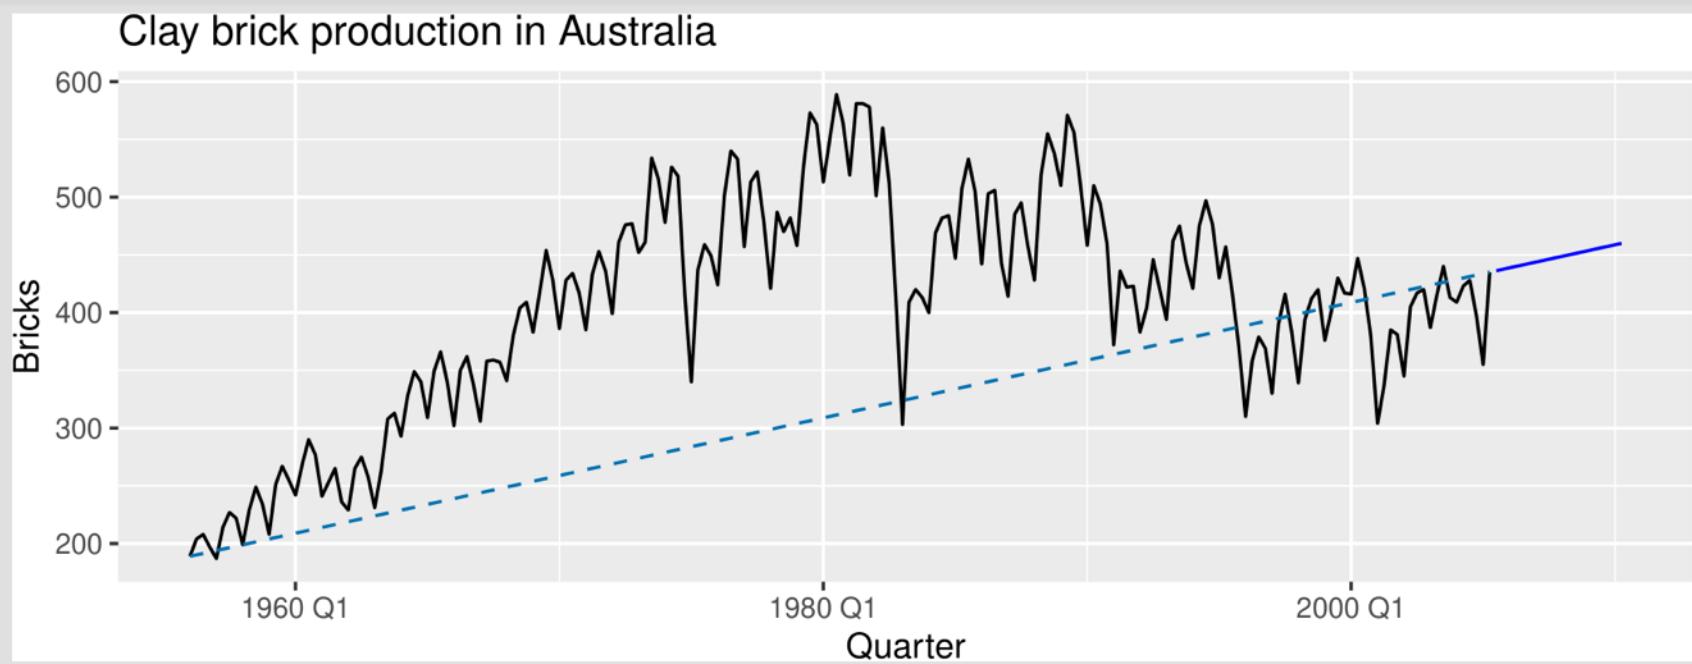
- Forecasts equal to last value plus average change.
- Forecasts:

$$\begin{aligned}\hat{y}_{T+h|T} &= y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) \\ &= y_T + \frac{h}{T-1} (y_T - y_1).\end{aligned}$$

- Equivalent to extrapolating a line drawn between first and last observations.

# Some simple forecasting methods

## Drift method



# Model fitting

The `model()` function trains models to data.

```
brick_fit <- aus_production %>%
  filter(!is.na(Bricks)) %>%
  model(
    Seasonal_naive = SNAIVE(Bricks),
    Naive = NAIVE(Bricks),
    Drift = RW(Bricks ~ drift()),
    Mean = MEAN(Bricks)
  )

## # A mable: 1 x 4
##   Seasonal_naive     Naive          Drift      Mean
##   <model> <model> <model> <model>
## 1 <SNAIVE> <NAIVE> <RW w/ drift> <MEAN>
```

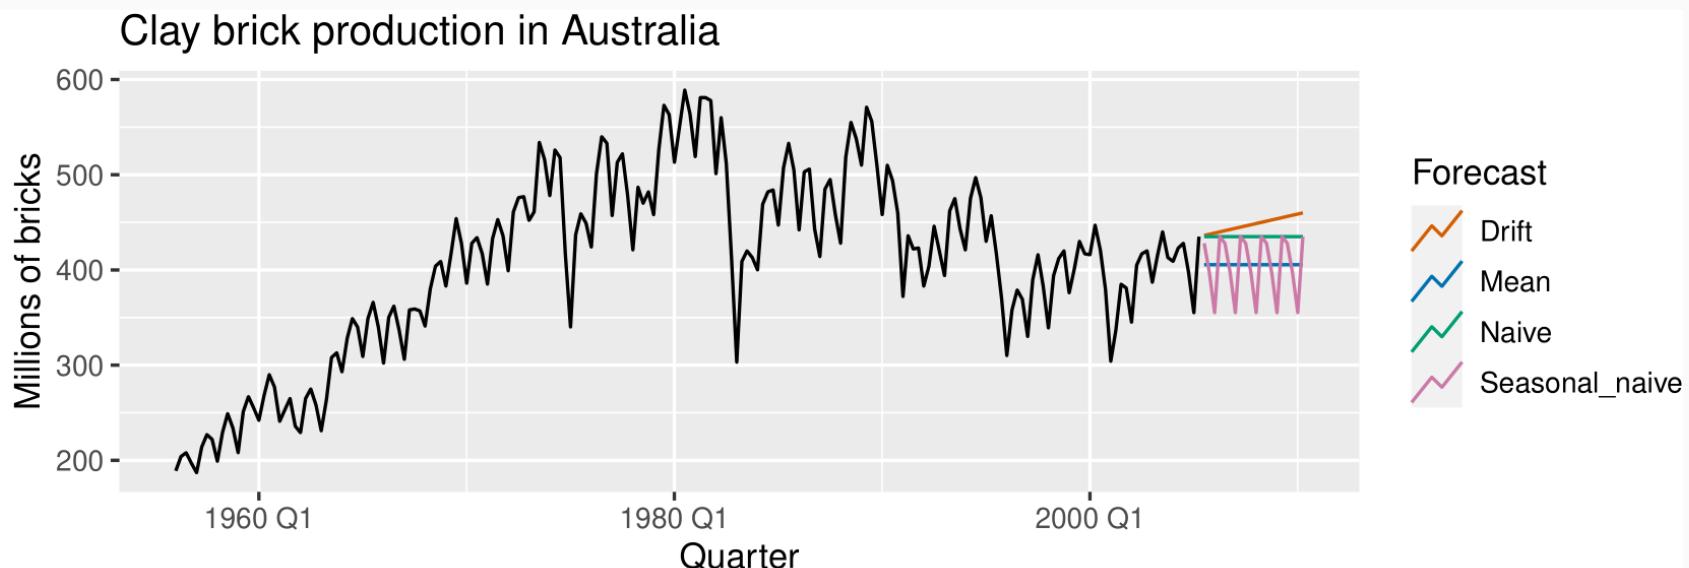
# Producing forecasts

```
brick_fc <- brick_fit %>%
  forecast(h = "5 years")

## # A fable: 80 x 4 [1Q]
## # Key:     .model [4]
##   .model      Quarter     Bricks .mean
##   <chr>       <qtr>     <dist> <dbl>
## 1 Seasonal_naive 2005 Q3 N(428, 2336) 428
## 2 Seasonal_naive 2005 Q4 N(397, 2336) 397
## 3 Seasonal_naive 2006 Q1 N(355, 2336) 355
## 4 Seasonal_naive 2006 Q2 N(435, 2336) 435
## # ... with 76 more rows
```

# Visualising forecasts

```
brick_fc %>%
  autoplot(aus_production, level = NULL) +
  labs(title = "Clay brick production in Australia",
       y = "Millions of bricks") +
  guides(colour = guide_legend(title = "Forecast"))
```



# Facebook closing stock price

```
# Extract training data
fb_stock <- gafa_stock %>%
  filter(Symbol == "FB") %>%
  mutate(trading_day = row_number()) %>%
  update_tsibble(index=trading_day, regular=TRUE)

# Specify, estimate and forecast
fb_stock %>%
  model(
    Mean = MEAN(Close),
    Naive = NAIVE(Close),
    Drift = RW(Close ~ drift())
  ) %>%
  forecast(h=42) %>%
  autoplot(fb_stock, level = NULL) +
  labs(title = "Facebook closing stock price", y="$US") +
  guides(colour=guide_legend(title="Forecast"))
```

# Facebook closing stock price



# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Fitted values

- $\hat{y}_{t|t-1}$  is the forecast of  $y_t$  based on observations  $y_1, \dots, y_{t-1}$ .
- We call these “fitted values”.
- Sometimes drop the subscript:  $\hat{y}_t \equiv \hat{y}_{t|t-1}$ .
- Often not true forecasts since parameters are estimated on all data.

## For example:

- $\hat{y}_t = \bar{y}$  for average method.
- $\hat{y}_t = y_{t-1} + (y_T - y_1)/(T - 1)$  for drift method.

# Forecasting residuals

**Residuals in forecasting:** difference between observed value and its fitted value:  $e_t = y_t - \hat{y}_{t|t-1}$ .

## Assumptions

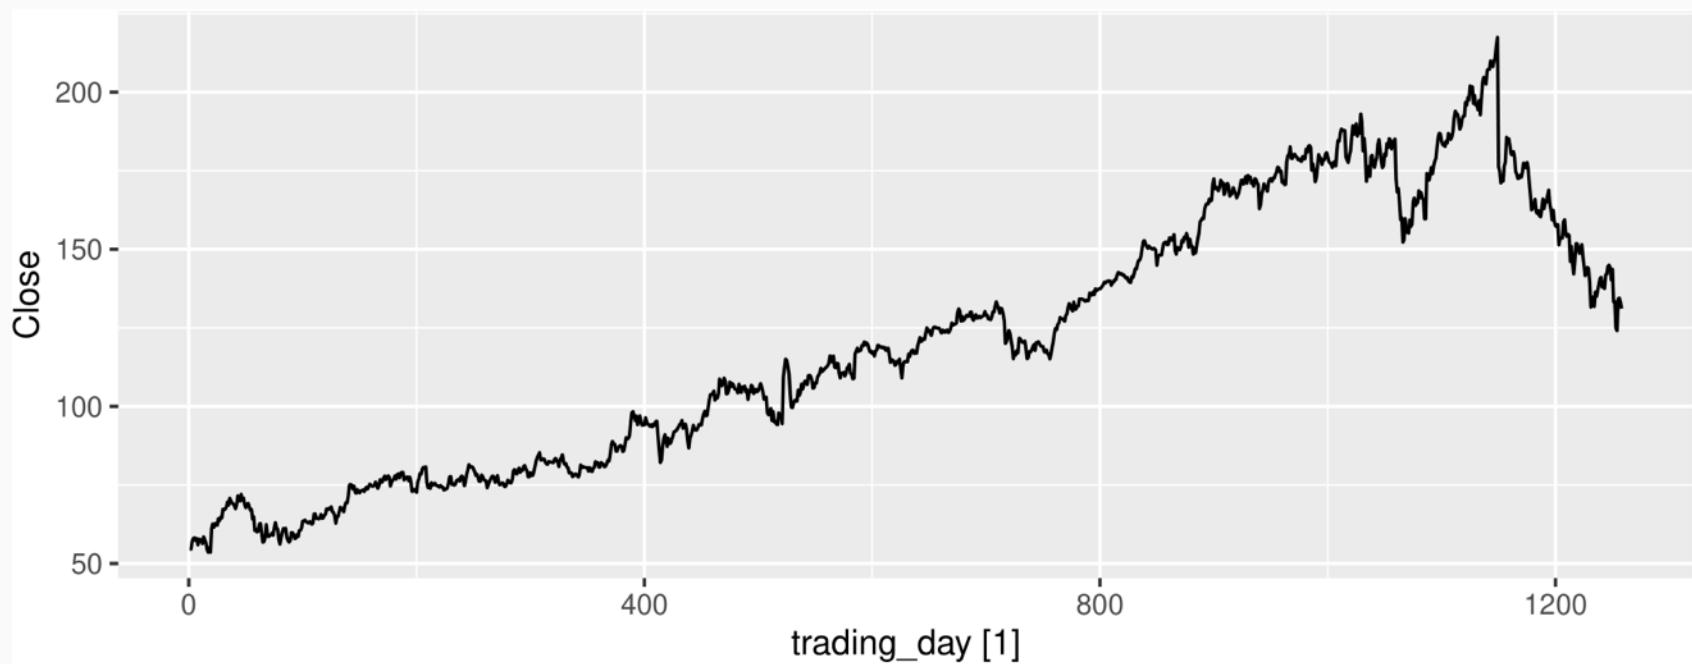
- 1  $\{e_t\}$  uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2  $\{e_t\}$  have mean zero. If they don't, then forecasts are biased.

## Useful properties (for distributions & prediction intervals)

- 3  $\{e_t\}$  have constant variance.
- 4  $\{e_t\}$  are normally distributed.

# Facebook closing stock price

```
fb_stock %>% autoplot(Close)
```



# Facebook closing stock price

```
fit <- fb_stock %>% model(NAIVE(Close))
augment(fit)

## # A tsibble: 1,258 x 7 [1]
## # Key:     Symbol, .model [1]
##   Symbol .model      trading_day Close .fitted .resid .innov
##   <chr>  <chr>        <int> <dbl>    <dbl>    <dbl>    <dbl>
## 1 FB    NAIVE(Close)     1  54.7     NA     NA     NA
## 2 FB    NAIVE(Close)     2  54.6    54.7 -0.150 -0.150
## 3 FB    NAIVE(Close)     3  57.2    54.6   2.64   2.64
## 4 FB    NAIVE(Close)     4  57.9    57.2   0.720  0.720
## 5 FB    NAIVE(Close)     5  58.2    57.9   0.310  0.310
## 6 FB    NAIVE(Close)     6  57.2    58.2  -1.01  -1.01
## 7 FB    NAIVE(Close)     7  57.9    57.2   0.720  0.720
## 8 FB    NAIVE(Close)     8  55.9    57.9  -2.03  -2.03
## 9 FB    NAIVE(Close)     9  57.7    55.9   1.83   1.83
## 10 FB   NAIVE(Close)    10  57.6    57.7  -0.140 -0.140
## # ... with 1,248 more rows
```

# Facebook closing stock price

```
fit <- fb_stock %>% model(NAIVE(Close))  
augment(fit)
```

```
## # A tsibble: 1,258 x 7 [1]  
## # Key:     Symbol, .model [1]  
##   Symbol .model      trading_day Close .fitted .resid .innov  
##   <chr>  <chr>        <int> <dbl>   <dbl>   <dbl>   <dbl>  
## 1 FB    NAIVE(Close)      1  54.7     NA     NA     NA  
## 2 FB    NAIVE(Close)      2  54.6   54.7 -0.150 -0.150  
## 3 FB    NAIVE(Close)      3  57.2   54.6   2.64   2.64  
## 4 FB    NAIVE(Close)      4  57.9   57.2   0.720  0.720  
## 5 FB    NAIVE(Close)      5  58.2   57.9   0.310  0.310  
## 6 FB    NAIVE(Close)      6  57.2   58.2 -1.01  -1.01  
## 7 FB    NAIVE(Close)      7  57.9   57.2   0.720  0.720  
## 8 FB    NAIVE(Close)      8  55.9   57.9 -2.03  -2.03  
## 9 FB    NAIVE(Close)      9  57.7   55.9   1.83   1.83  
## 10 FB   NAIVE(Close)     10  57.6   57.7 -0.140 -0.140  
## # ... with 1,248 more rows
```

$\hat{y}_{t|t-1}$

$e_t$

**Naïve forecasts:**

$$\hat{y}_{t|t-1} = y_{t-1}$$

$$e_t = y_t - \hat{y}_{t|t-1} = y_t - y_{t-1}$$

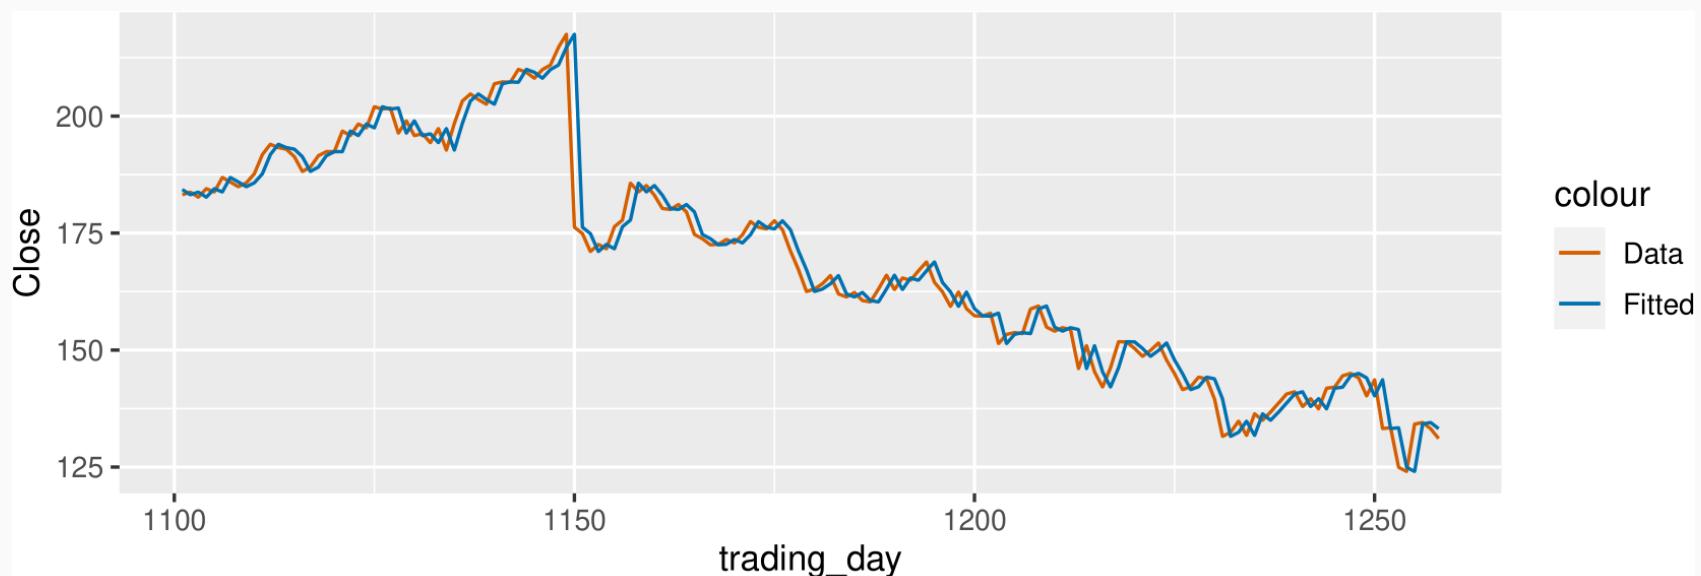
# Facebook closing stock price

```
augment(fit) %>%
  ggplot(aes(x = trading_day)) +
  geom_line(aes(y = Close, colour = "Data")) +
  geom_line(aes(y = .fitted, colour = "Fitted"))
```



# Facebook closing stock price

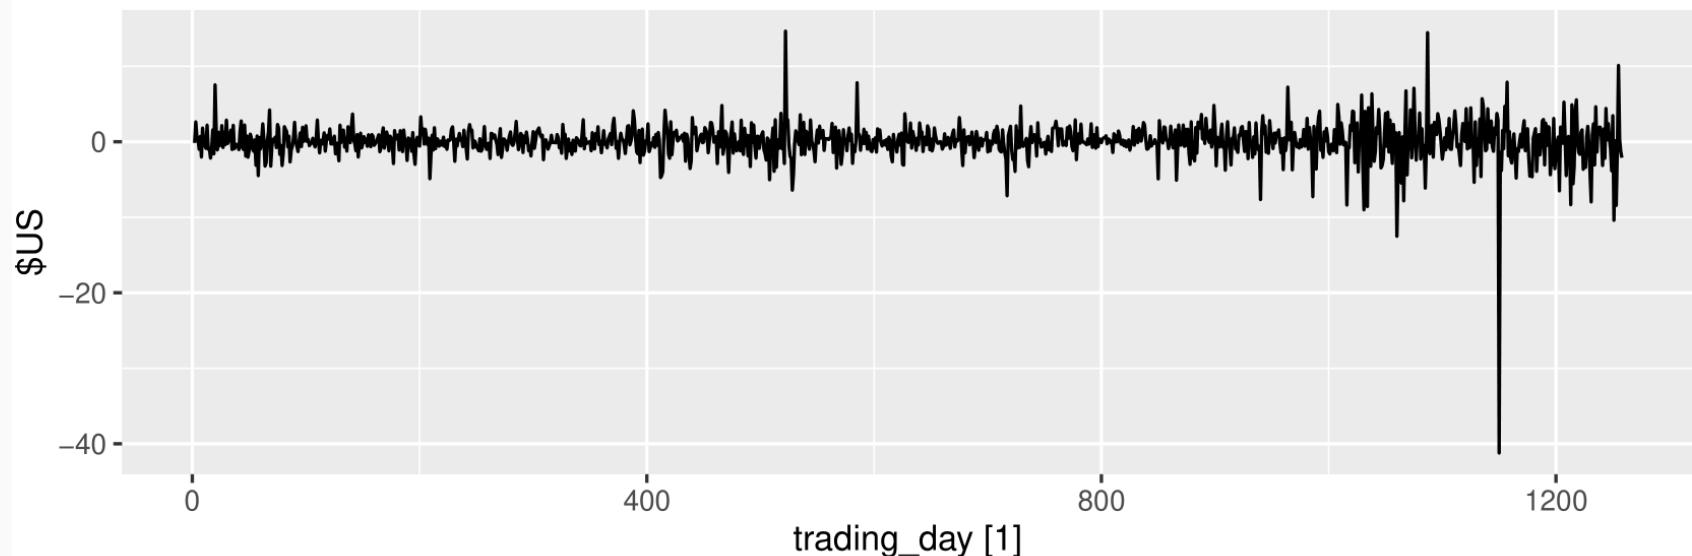
```
augment(fit) %>%
  filter(trading_day > 1100) %>%
  ggplot(aes(x = trading_day)) +
  geom_line(aes(y = Close, colour = "Data")) +
  geom_line(aes(y = .fitted, colour = "Fitted"))
```



# Facebook closing stock price

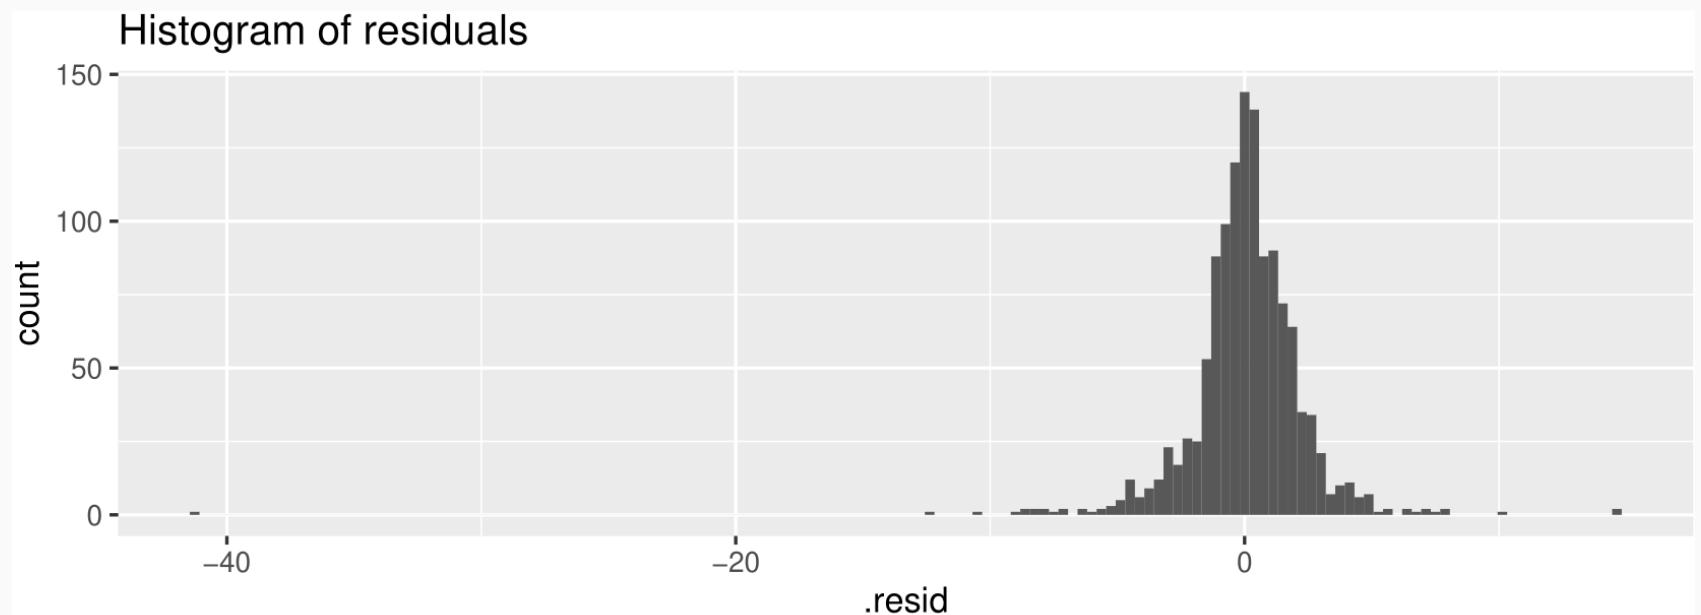
```
augment(fit) %>%
  autoplot(.resid) +
  labs(y = "$US",
       title = "Residuals from naïve method")
```

Residuals from naïve method



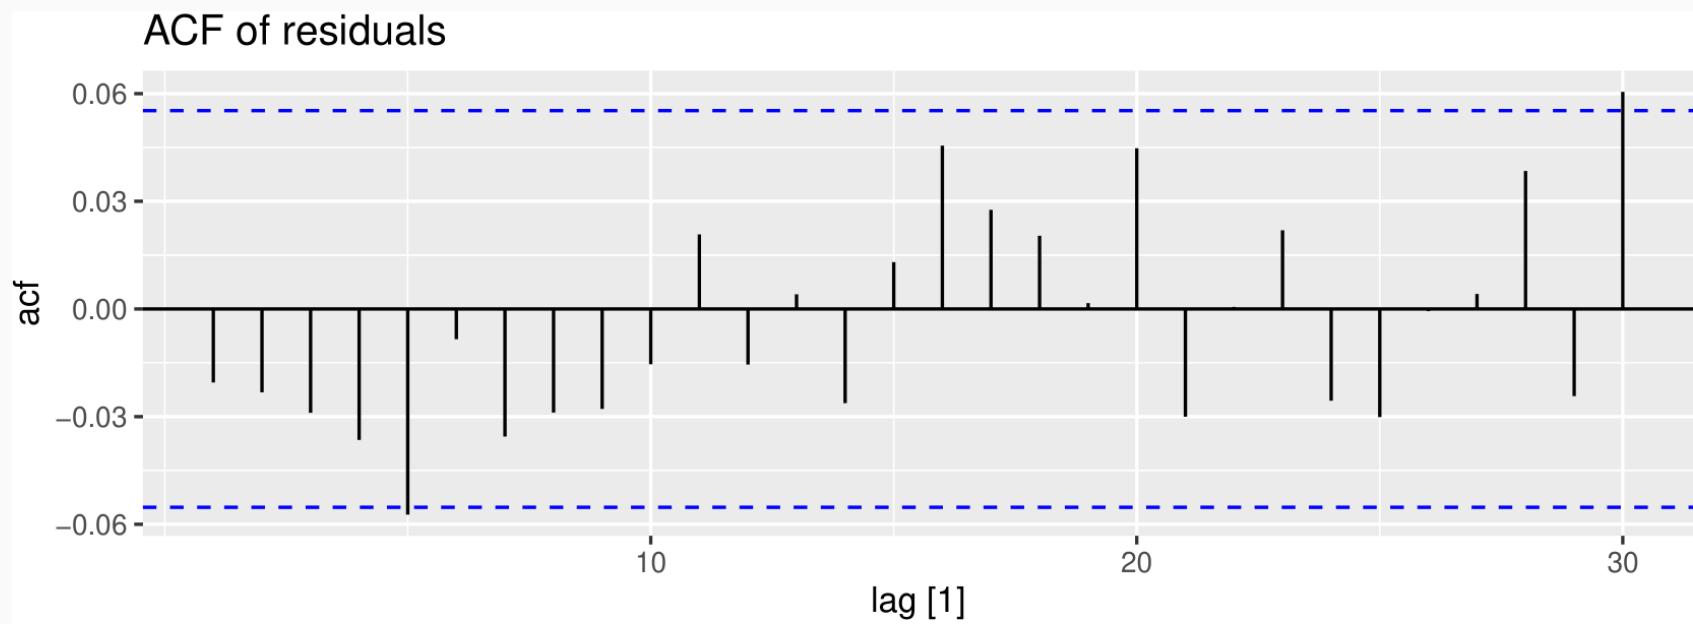
# Facebook closing stock price

```
augment(fit) %>%
  ggplot(aes(x = .resid)) +
  geom_histogram(bins = 150) +
  labs(title = "Histogram of residuals")
```



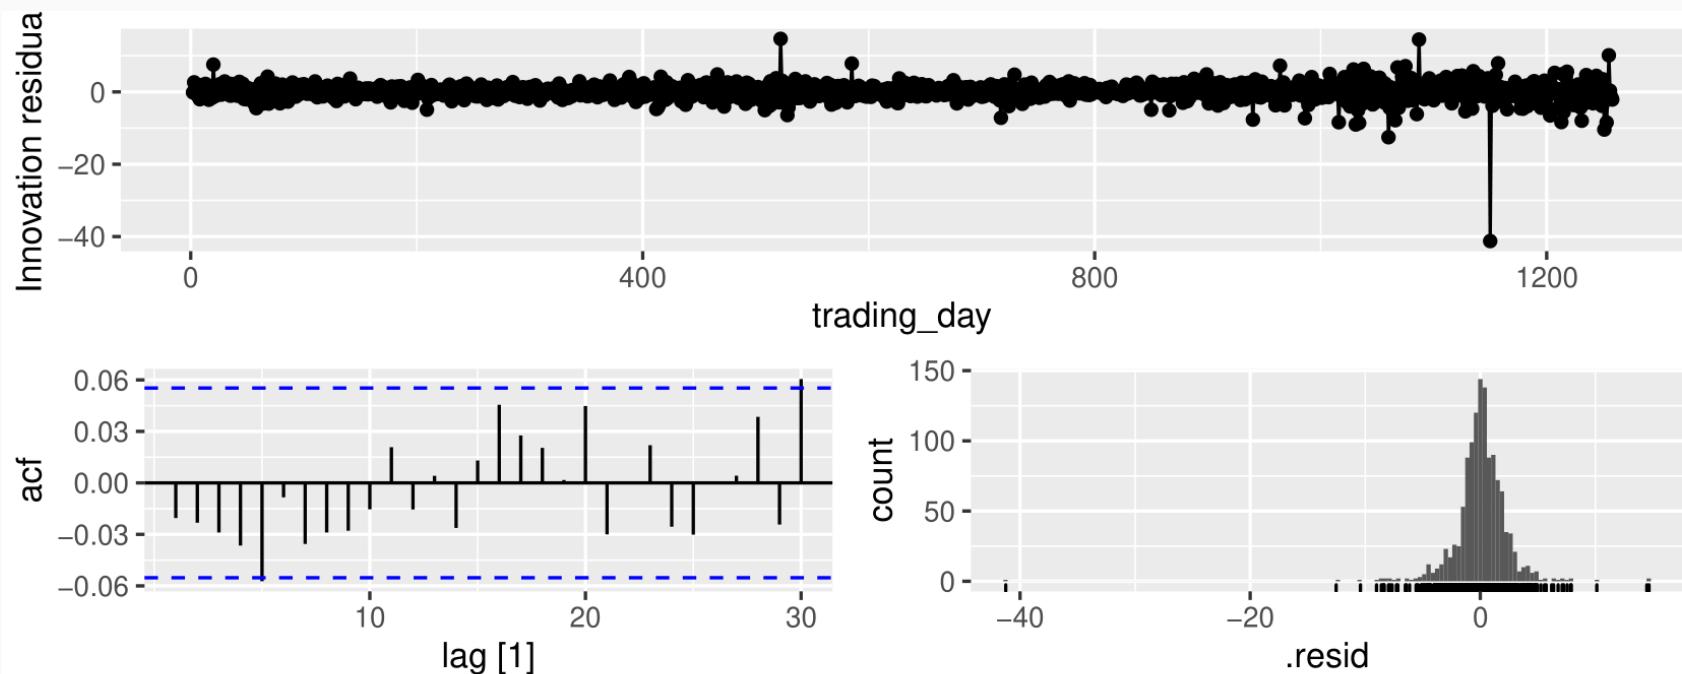
# Facebook closing stock price

```
augment(fit) %>%
  ACF(.resid) %>%
  autoplot() + labs(title = "ACF of residuals")
```



# gg\_tsresiduals() function

```
gg_tsresiduals(fit)
```



# ACF of residuals

- We assume that the residuals are white noise (uncorrelated, mean zero, constant variance). If they aren't, then there is information left in the residuals that should be used in computing forecasts.
- So a standard residual diagnostic is to check the ACF of the residuals of a forecasting method.
- We expect these to look like white noise.

# Portmanteau tests

Consider a *whole set* of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Box-Pierce test

$$Q = T \sum_{k=1}^{\ell} r_k^2$$

where  $\ell$  is max lag being considered and  $T$  is number of observations.

- If each  $r_k$  close to zero,  $Q$  will be **small**.
- If some  $r_k$  values large (positive or negative),  $Q$  will be **large**.

# Portmanteau tests

Consider a *whole set* of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Ljung-Box test

$$Q^* = T(T + 2) \sum_{k=1}^{\ell} (T - k)^{-1} r_k^2$$

where  $\ell$  is max lag being considered and  $T$  is number of observations.

- My preferences:  $\ell = 10$  for non-seasonal data,  $h = 2m$  for seasonal data.
- Better performance, especially in small samples.

# Portmanteau tests

- If data are WN,  $Q^*$  has  $\chi^2$  distribution with  $(\ell - K)$  degrees of freedom where  $K$  = no. parameters in model.
- When applied to raw data, set  $K = 0$ .
- $\text{lag} = \ell$ ,  $\text{dof} = K$

```
augment(fit) %>%
  features(.resid, ljung_box, lag=10, dof=0)
```

```
## # A tibble: 1 x 4
##   Symbol .model      lb_stat lb_pvalue
##   <chr>  <chr>       <dbl>     <dbl>
## 1 FB    NAIVE(Close)  12.1     0.276
```

# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Forecast distributions

- A forecast  $\hat{y}_{T+h|T}$  is (usually) the mean of the conditional distribution  $y_{T+h} \mid y_1, \dots, y_T$ .
- Most time series models produce normally distributed forecasts.
- The forecast distribution describes the probability of observing any future value.

# Forecast distributions

Assuming residuals are normal, uncorrelated,  $\text{sd} = \hat{\sigma}$ :

**Mean:**  $\hat{y}_{T+h|T} \sim N(\bar{y}, (1 + 1/T)\hat{\sigma}^2)$

**Naïve:**  $\hat{y}_{T+h|T} \sim N(y_T, h\hat{\sigma}^2)$

**Seasonal naïve:**  $\hat{y}_{T+h|T} \sim N(y_{T+h-m(k+1)}, (k + 1)\hat{\sigma}^2)$

**Drift:**  $\hat{y}_{T+h|T} \sim N(y_T + \frac{h}{T-1}(y_T - y_1), h\frac{T+h}{T}\hat{\sigma}^2)$

where  $k$  is the integer part of  $(h - 1)/m$ .

Note that when  $h = 1$  and  $T$  is large, these all give the same approximate forecast variance:  $\hat{\sigma}^2$ .

# Prediction intervals

- A prediction interval gives a region within which we expect  $y_{T+h}$  to lie with a specified probability.
- Assuming forecast errors are normally distributed, then a 95% PI is

$$\hat{y}_{T+h|T} \pm 1.96 \hat{\sigma}_h$$

where  $\hat{\sigma}_h$  is the st dev of the  $h$ -step distribution.

- When  $h = 1$ ,  $\hat{\sigma}_h$  can be estimated from the residuals.

# Prediction intervals

```
brick_fc %>% hilo(level = 95)
```

```
## # A tsibble: 80 x 5 [1Q]
## # Key:   .model [4]
##       .model     Quarter    Bricks .mean      `95%` 
##       <chr>      <qtr>     <dist> <dbl>      <hilo>
## 1 Seasonal_naive 2005 Q3 N(428, 2336)  428 [333, 523]95
## 2 Seasonal_naive 2005 Q4 N(397, 2336)  397 [302, 492]95
## 3 Seasonal_naive 2006 Q1 N(355, 2336)  355 [260, 450]95
## 4 Seasonal_naive 2006 Q2 N(435, 2336)  435 [340, 530]95
## 5 Seasonal_naive 2006 Q3 N(428, 4672)  428 [294, 562]95
## 6 Seasonal_naive 2006 Q4 N(397, 4672)  397 [263, 531]95
## 7 Seasonal_naive 2007 Q1 N(355, 4672)  355 [221, 489]95
## 8 Seasonal_naive 2007 Q2 N(435, 4672)  435 [301, 569]95
## 9 Seasonal_naive 2007 Q3 N(428, 7008)  428 [264, 592]95
```

# Prediction intervals

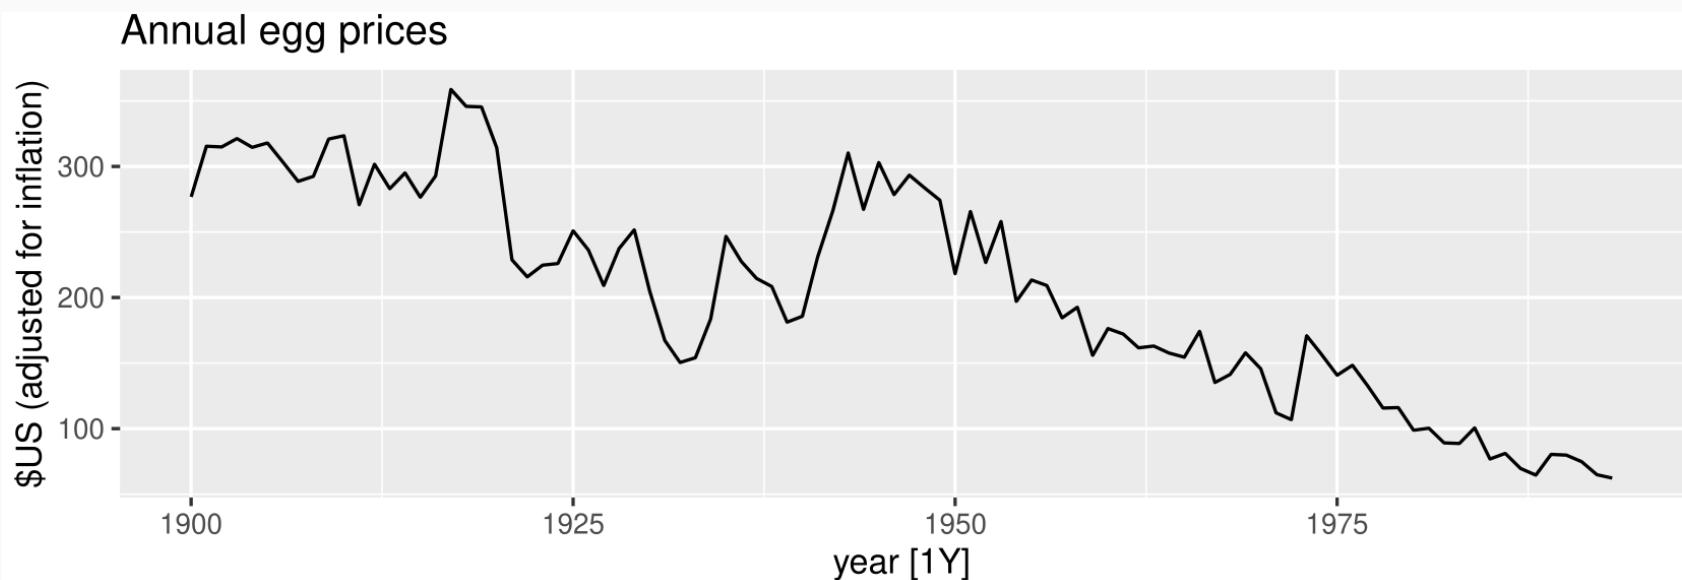
- Point forecasts often useless without a measure of uncertainty (such as prediction intervals).
- Prediction intervals require a stochastic model (with random errors, etc).
- For most models, prediction intervals get wider as the forecast horizon increases.
- Use level argument to control coverage.
- Check residual assumptions before believing them.
- Usually too narrow due to unaccounted uncertainty.

# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Modelling with transformations

```
eggs <- prices %>%
  filter(!is.na(eggs)) %>% select(eggs)
eggs %>% autoplot() +
  labs(title="Annual egg prices",
       y="$US (adjusted for inflation)")
```



# Modelling with transformations

Transformations used in the left of the formula will be automatically back-transformed. To model log-transformed egg prices, you could use:

```
fit <- eggs %>%
  model(RW(log(eggs) ~ drift()))
fit

## # A mable: 1 x 1
##   `RW(log(eggs) ~ drift())` 
##                 <model>
## 1             <RW w/ drift>
```

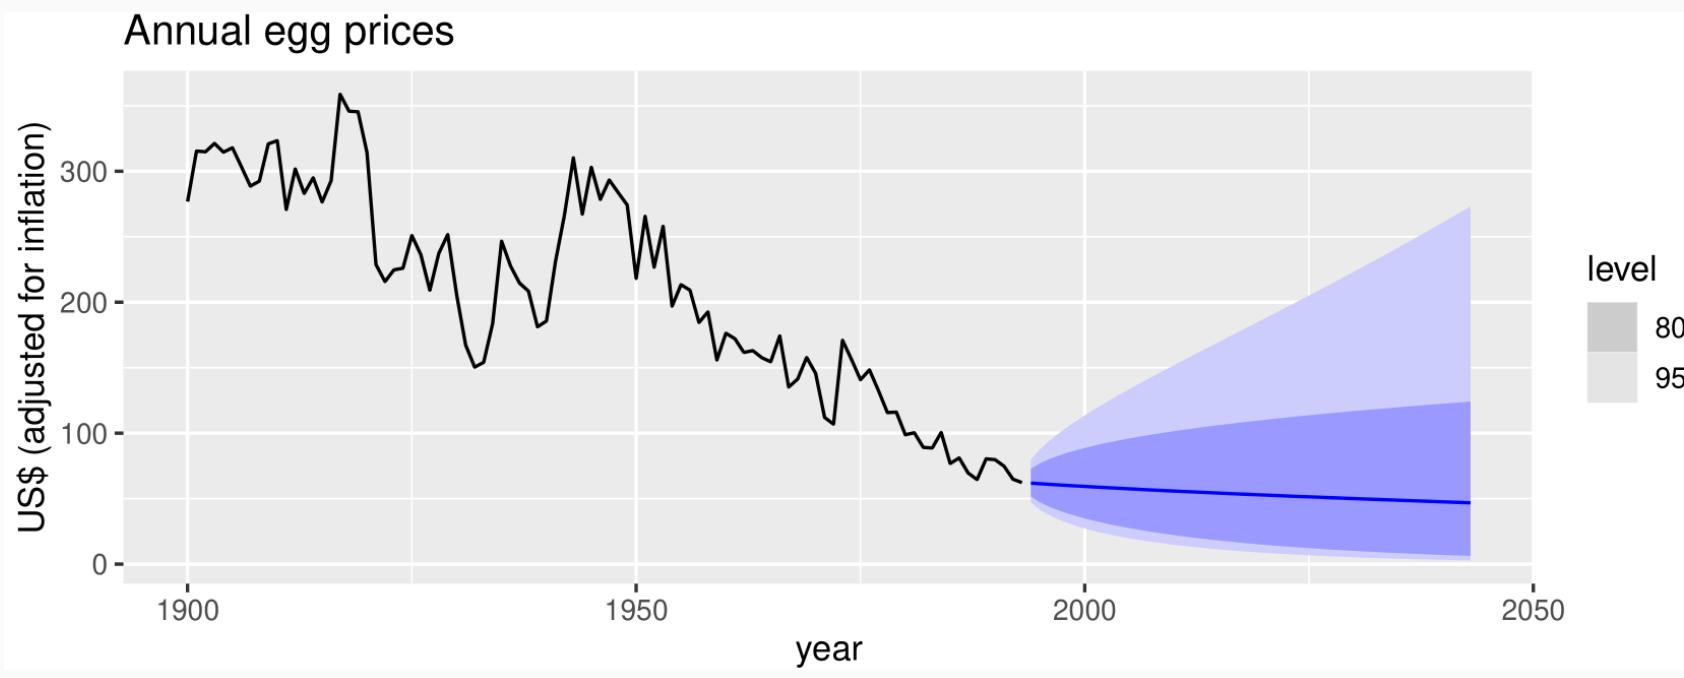
# Forecasting with transformations

```
fc <- fit %>%
  forecast(h = 50)
fc

## # A fable: 50 x 4 [1Y]
## # Key:     .model [1]
##   .model                 year          eggs .mean
##   <chr>                 <dbl>        <dist> <dbl>
## 1 RW(log(eggs) ~ drift()) 1994 t(N(4.1, 0.018)) 61.8
## 2 RW(log(eggs) ~ drift()) 1995 t(N(4.1, 0.036)) 61.4
## 3 RW(log(eggs) ~ drift()) 1996 t(N(4.1, 0.054)) 61.0
## 4 RW(log(eggs) ~ drift()) 1997 t(N(4.1, 0.073)) 60.5
## 5 RW(log(eggs) ~ drift()) 1998 t(N(4.1, 0.093)) 60.1
## 6 RW(log(eggs) ~ drift()) 1999 t(N(4, 0.11)) 59.7
## 7 RW(log(eggs) ~ drift()) 2000 t(N(4, 0.13)) 59.3
## 8 RW(log(eggs) ~ drift()) 2001 t(N(4, 0.15)) 59.0
```

# Forecasting with transformations

```
fc %>% autoplot(eggs) +  
  labs(title="Annual egg prices",  
       y="US$ (adjusted for inflation)")
```



# Bias adjustment

- Back-transformed point forecasts are medians.
- Back-transformed PI have the correct coverage.

## Back-transformed means

Let  $X$  be have mean  $\mu$  and variance  $\sigma^2$ .

Let  $f(x)$  be back-transformation function, and  $Y = f(X)$ .

Taylor series expansion about  $\mu$ :

$$f(X) = f(\mu) + (X - \mu)f'(\mu) + \frac{1}{2}(X - \mu)^2f''(\mu).$$

$$E[Y] = E[f(X)] = f(\mu) + \frac{1}{2}\sigma^2f''(\mu)$$

# Bias adjustment

**Box-Cox back-transformation:**

$$y_t = \begin{cases} \exp(w_t) & \lambda = 0; \\ (\lambda W_t + 1)^{1/\lambda} & \lambda \neq 0. \end{cases}$$

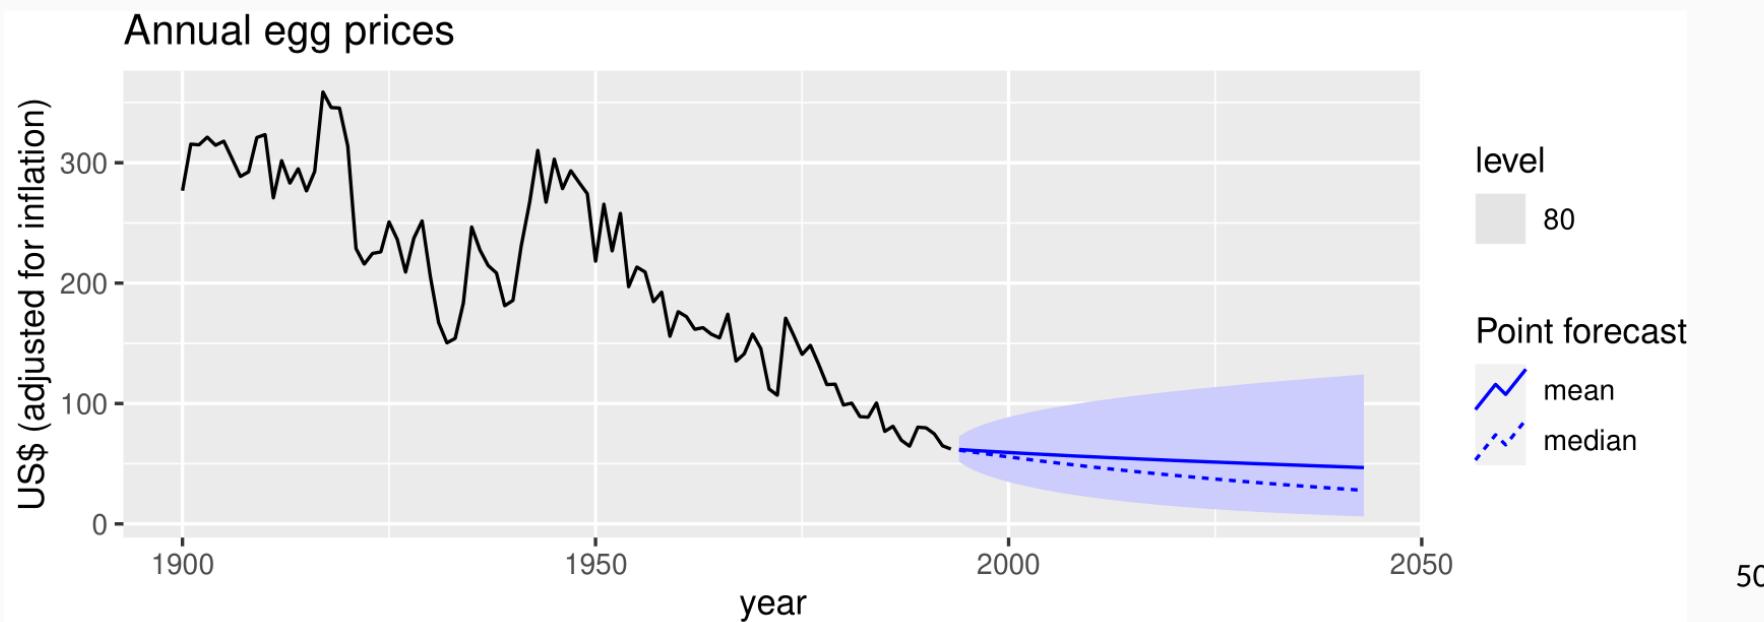
$$f(x) = \begin{cases} e^x & \lambda = 0; \\ (\lambda x + 1)^{1/\lambda} & \lambda \neq 0. \end{cases}$$

$$f''(x) = \begin{cases} e^x & \lambda = 0; \\ (1 - \lambda)(\lambda x + 1)^{1/\lambda - 2} & \lambda \neq 0. \end{cases}$$

$$E[Y] = \begin{cases} e^\mu \left[ 1 + \frac{\sigma^2}{2} \right] & \lambda = 0; \\ (\lambda \mu + 1)^{1/\lambda} \left[ 1 + \frac{\sigma^2(1-\lambda)}{2(\lambda \mu + 1)^2} \right] & \lambda \neq 0. \end{cases}$$

# Bias adjustment

```
fc %>%
  autoplot(eggs, level = 80, point_forecast = lst(mean, median)) +
  labs(title="Annual egg prices",
       y="US$ (adjusted for inflation)")
```



# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Forecasting and decomposition

$$y_t = \hat{S}_t + \hat{A}_t$$

- $\hat{A}_t$  is seasonally adjusted component
  - $\hat{S}_t$  is seasonal component.
- 
- Forecast  $\hat{S}_t$  using SNAIVE.
  - Forecast  $\hat{A}_t$  using non-seasonal time series method.
  - Combine forecasts of  $\hat{S}_t$  and  $\hat{A}_t$  to get forecasts of original data.

# US Retail Employment

```
us_retail_employment <- us_employment %>%
  filter(year(Month) >= 1990, Title == "Retail Trade") %>%
  select(-Series_ID)
us_retail_employment
```

```
## # A tsibble: 357 x 3 [1M]
##       Month Title     Employed
##       <mth> <chr>     <dbl>
## 1 1990 Jan Retail Trade 13256.
## 2 1990 Feb Retail Trade 12966.
## 3 1990 Mar Retail Trade 12938.
## 4 1990 Apr Retail Trade 13012.
## 5 1990 May Retail Trade 13108.
## 6 1990 Jun Retail Trade 13183.
## 7 1990 Jul Retail Trade 13170.
## 8 1990 Aug Retail Trade 13160.
## 9 1990 Sep Retail Trade 12112
```

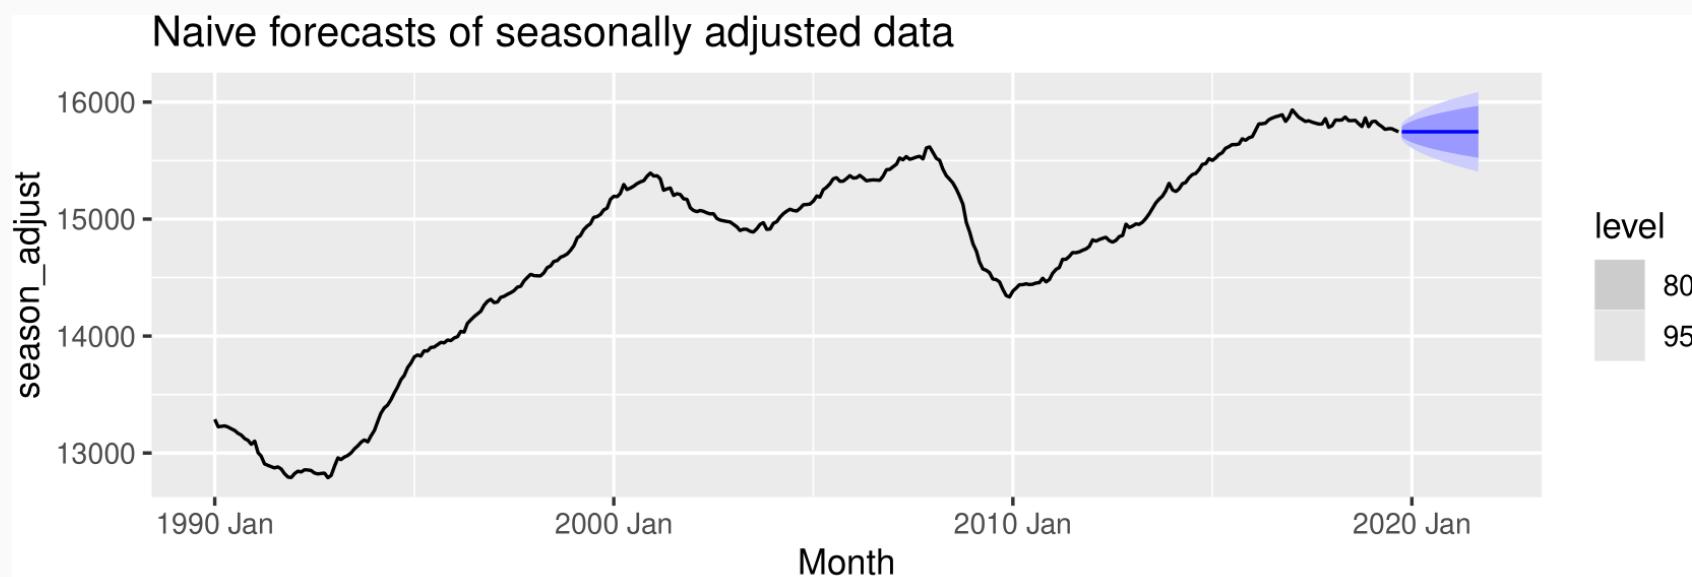
# US Retail Employment

```
dcmp <- us_retail_employment %>%
  model(STL(Employed)) %>%
  components() %>% select(-.model)
dcmp
```

```
## # A tsibble: 357 x 6 [1M]
##       Month Employed trend season_year remainder season_adjust
##       <mth>    <dbl>  <dbl>      <dbl>    <dbl>        <dbl>
## 1 1990 Jan     13256. 13288.    -33.0     0.836     13289.
## 2 1990 Feb     12966. 13269.   -258.     -44.6     13224.
## 3 1990 Mar     12938. 13250.   -290.     -22.1     13228.
## 4 1990 Apr     13012. 13231.   -220.      1.05     13232.
## 5 1990 May     13108. 13211.   -114.      11.3     13223.
## 6 1990 Jun     13183. 13192.   -24.3      15.5     13207.
## 7 1990 Jul     13170. 13172.   -23.2      21.6     13193.
## 8 1990 Aug     13160. 13151.   -9.52     17.8     13169.
## 9 1990 Sep     12112. 12121.   -20.5      22.0     12152.
```

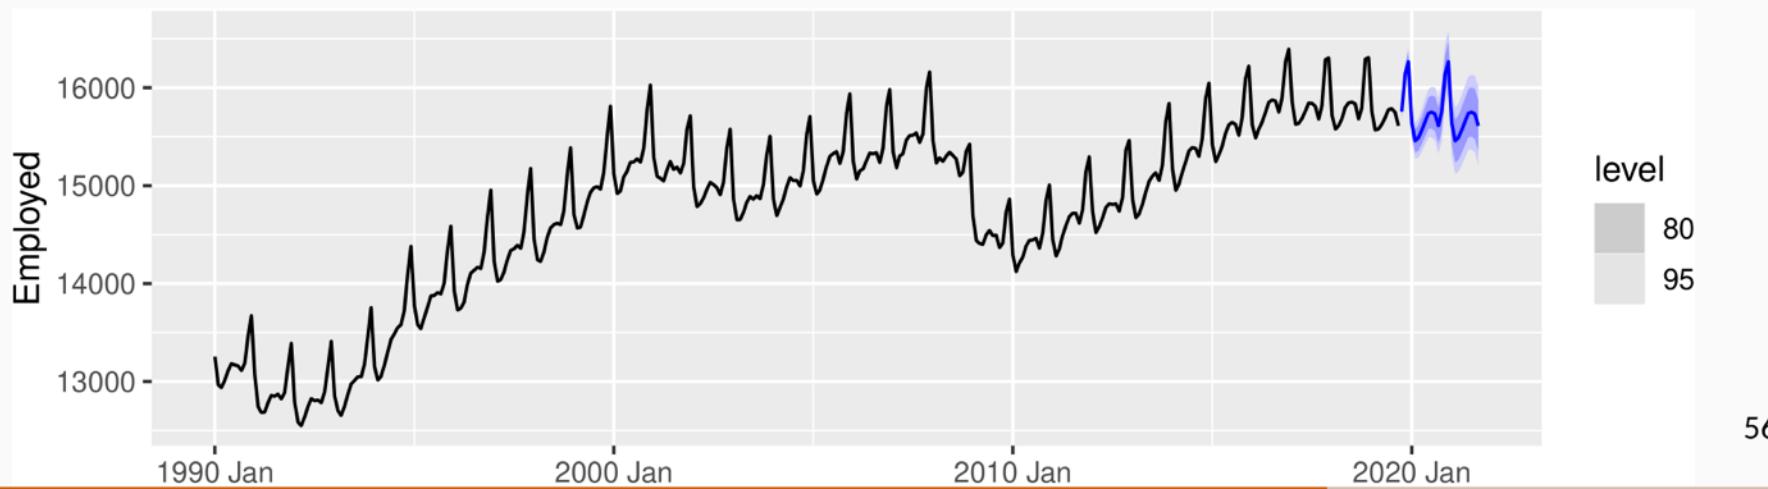
# US Retail Employment

```
dcmp %>%
  model(NAIVE(season_adjust)) %>%
  forecast() %>%
  autoplot(dcmp) +
  labs(title = "Naive forecasts of seasonally adjusted data")
```



# US Retail Employment

```
us_retail_employment %>%
  model(stlf = decomposition_model(
    STL(Employed ~ trend(window = 7), robust = TRUE),
    NAIVE(season_adjust)
  )) %>%
  forecast() %>%
  autoplot(us_retail_employment)
```



# Decomposition models

`decomposition_model()` creates a decomposition model

- You must provide a method for forecasting the `season_adjust` series.
- A seasonal naive method is used by default for the seasonal components.
- The variances from both the seasonally adjusted and seasonal forecasts are combined.

# Outline

- 1 A tidy forecasting workflow
- 2 Some simple forecasting methods
- 3 Residual diagnostics
- 4 Distributional forecasts and prediction intervals
- 5 Forecasting with transformations
- 6 Forecasting and decomposition
- 7 Evaluating forecast accuracy

# Training and test sets



- A model which fits the training data well will not necessarily forecast well.
- A perfect fit can always be obtained by using a model with enough parameters.
- Over-fitting a model to data is just as bad as failing to identify a systematic pattern in the data.
- The test set must not be used for *any* aspect of model development or calculation of forecasts.
- Forecast accuracy is based only on the test set.

# Forecast errors

Forecast “error”: the difference between an observed value and its forecast.

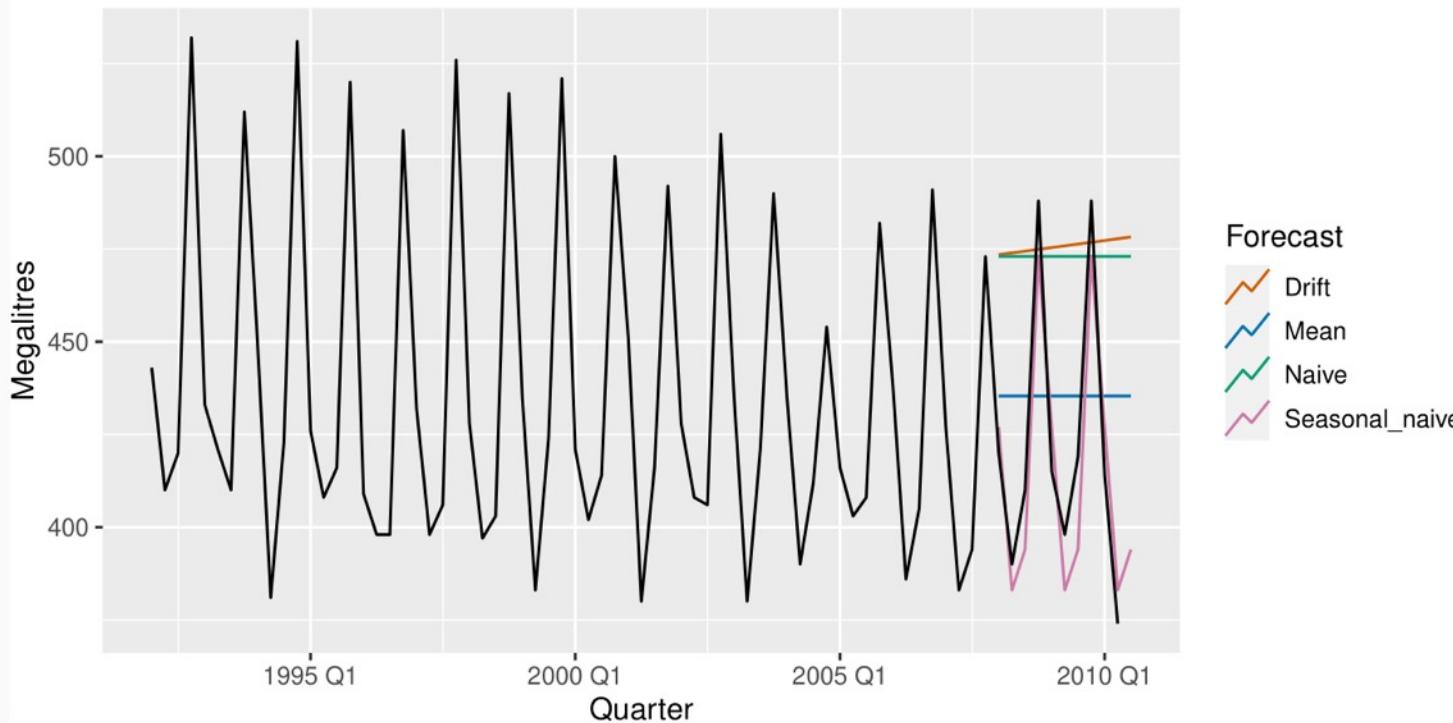
$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by  $\{y_1, \dots, y_T\}$

- Unlike residuals, forecast errors on the test set involve multi-step forecasts.
- These are *true* forecast errors as the test data is not used in computing  $\hat{y}_{T+h|T}$ .

# Measures of forecast accuracy

Forecasts for quarterly beer production



# Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE =  $\text{mean}(|e_{T+h}|)$

MSE =  $\text{mean}(e_{T+h}^2)$

RMSE =  $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE =  $100\text{mean}(|e_{T+h}| / |y_{T+h}|)$

# Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}| / |y_{T+h}|)$$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $t$ , and  $y$  has a natural zero.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}| / Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

For non-seasonal time series,

$$Q = (T - 1)^{-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naïve method.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}| / Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

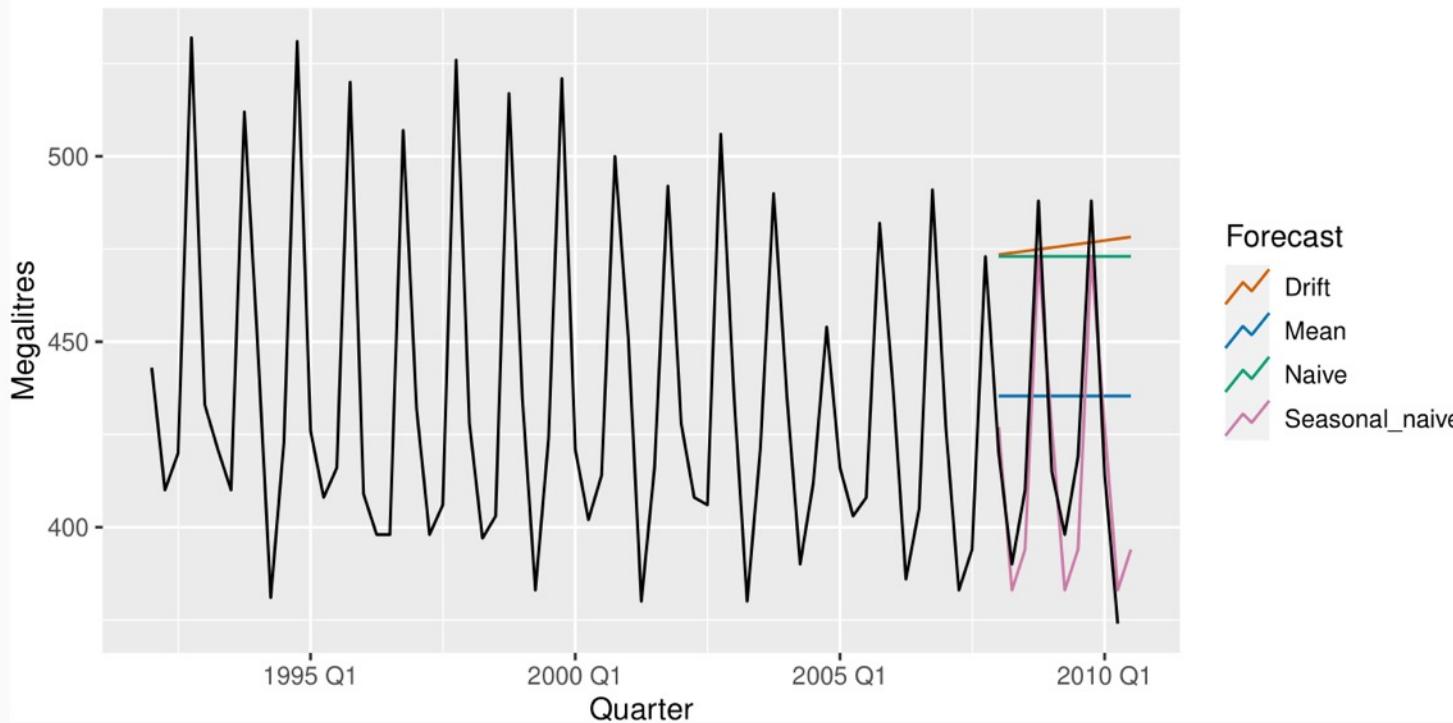
For seasonal time series,

$$Q = (T - m)^{-1} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naïve method.

# Measures of forecast accuracy

Forecasts for quarterly beer production



# Measures of forecast accuracy

```
recent_production <- aus_production %>%
  filter(year(Quarter) >= 1992)
train <- recent_production %>%
  filter(year(Quarter) <= 2007)
beer_fit <- train %>%
  model(
    Mean = MEAN(Beer),
    Naive = NAIVE(Beer),
    Seasonal_naive = SNAIVE(Beer),
    Drift = RW(Beer ~ drift()))
)
beer_fc <- beer_fit %>%
  forecast(h = 10)
```

# Measures of forecast accuracy

```
accuracy(beer_fit)
```

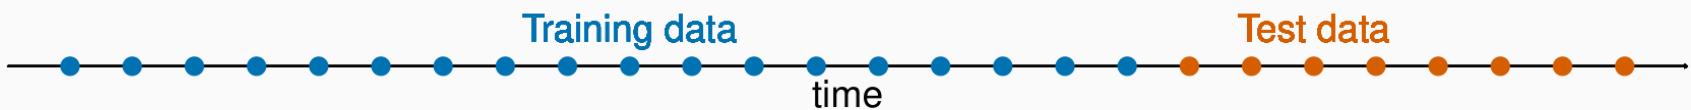
```
## # A tibble: 4 x 6
##   .model      .type    RMSE    MAE    MAPE    MASE
##   <chr>      <chr>    <dbl>   <dbl>   <dbl>   <dbl>
## 1 Drift     Training  65.3   54.8  12.2   3.83
## 2 Mean      Training  43.6   35.2  7.89   2.46
## 3 Naive     Training  65.3   54.7  12.2   3.83
## 4 Seasonal_naive Training 16.8   14.3  3.31   1
```

```
accuracy(beer_fc, recent_production)
```

```
## # A tibble: 4 x 6
##   .model      .type    RMSE    MAE    MAPE    MASE
##   <chr>      <chr>    <dbl>   <dbl>   <dbl>   <dbl>
## 1 Drift     Test     64.9   58.9  14.6   4.12
## 2 Mean      Test     38.4   34.8  8.28   2.44
## 3 Naive     Test     62.7   57.4  14.2   4.01
## 4 Seasonal_naive Test    14.3  13.4  3.17  0.937
```

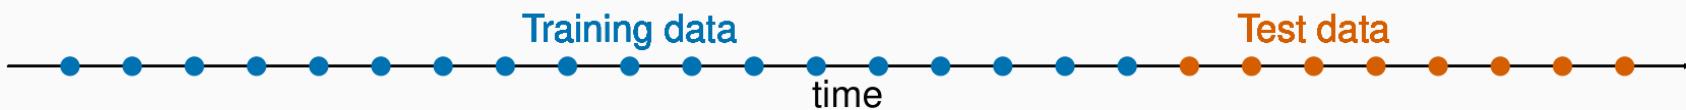
# Time series cross-validation

## Traditional evaluation

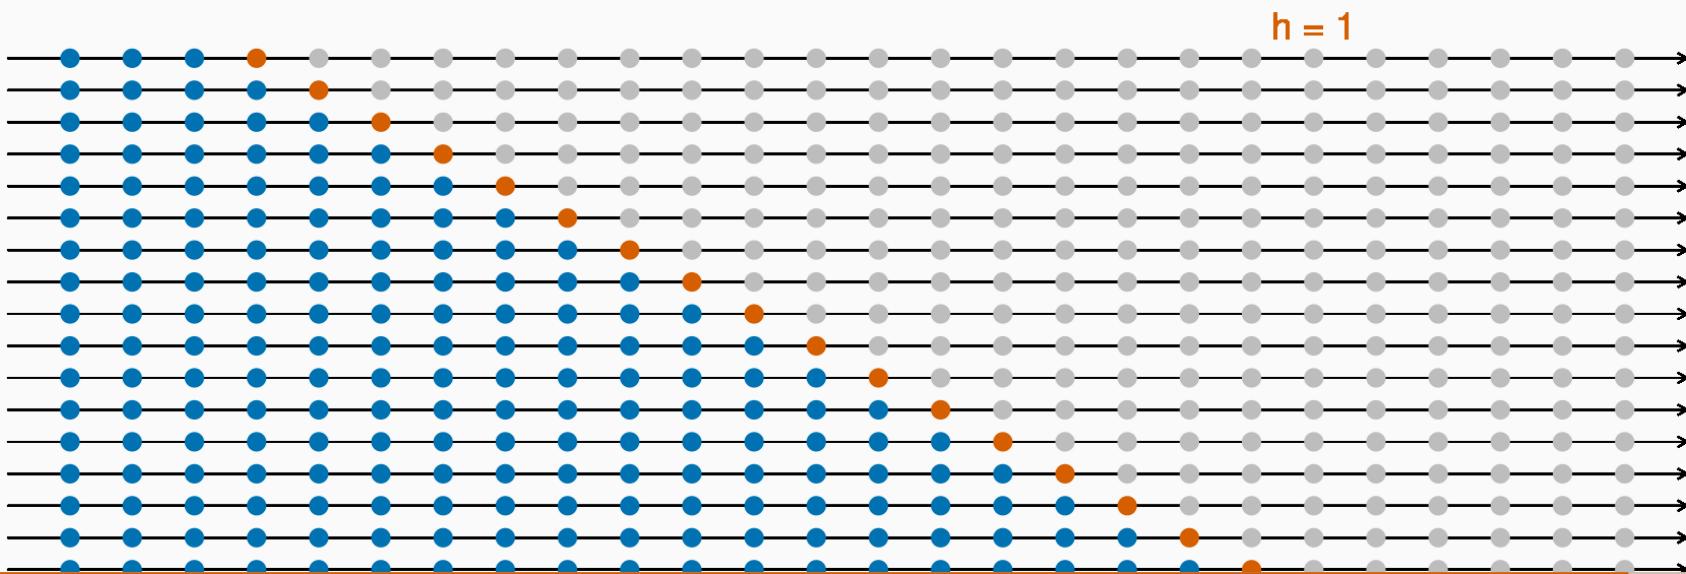


# Time series cross-validation

## Traditional evaluation

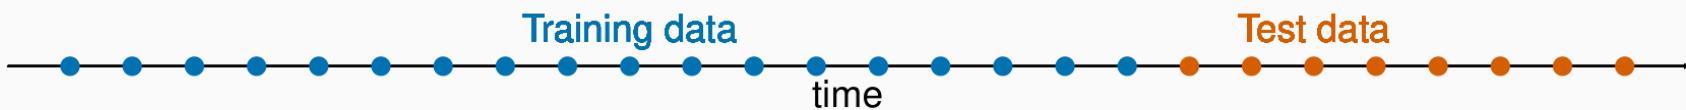


## Time series cross-validation

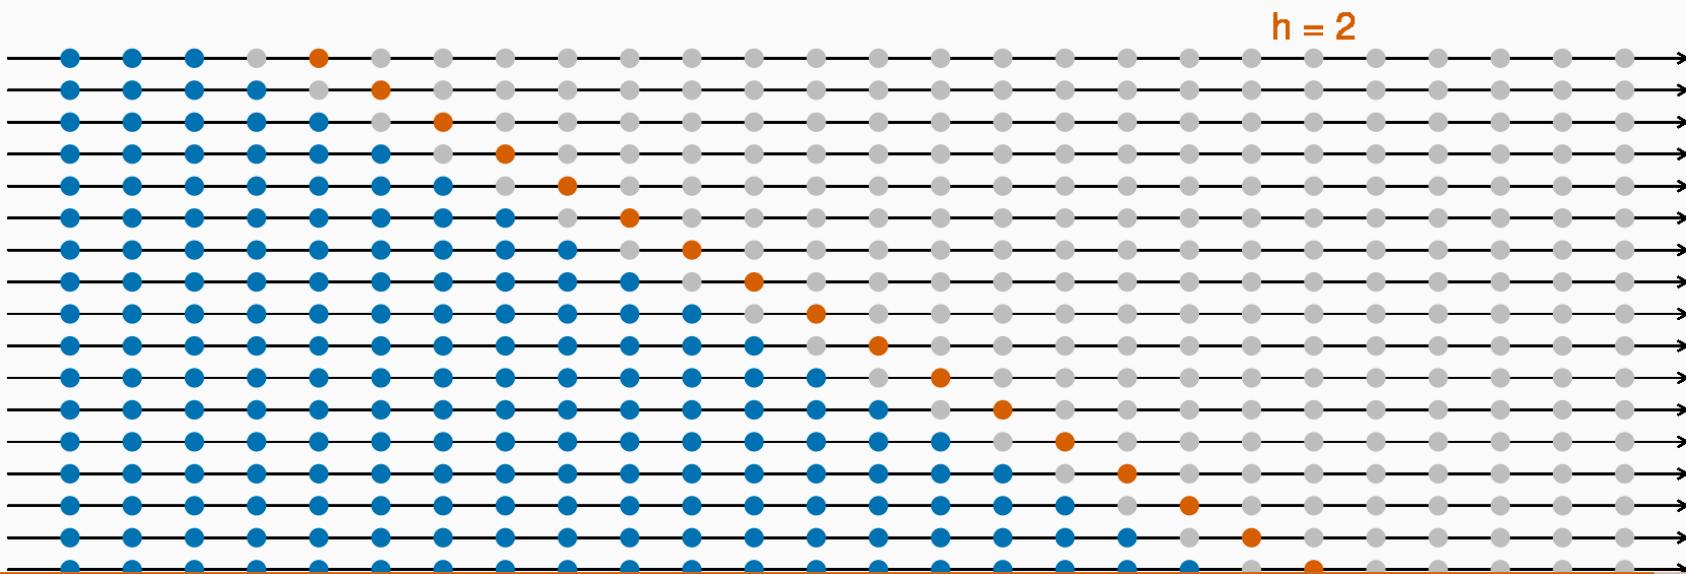


# Time series cross-validation

## Traditional evaluation

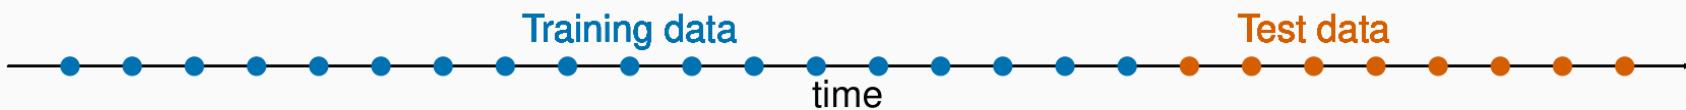


## Time series cross-validation

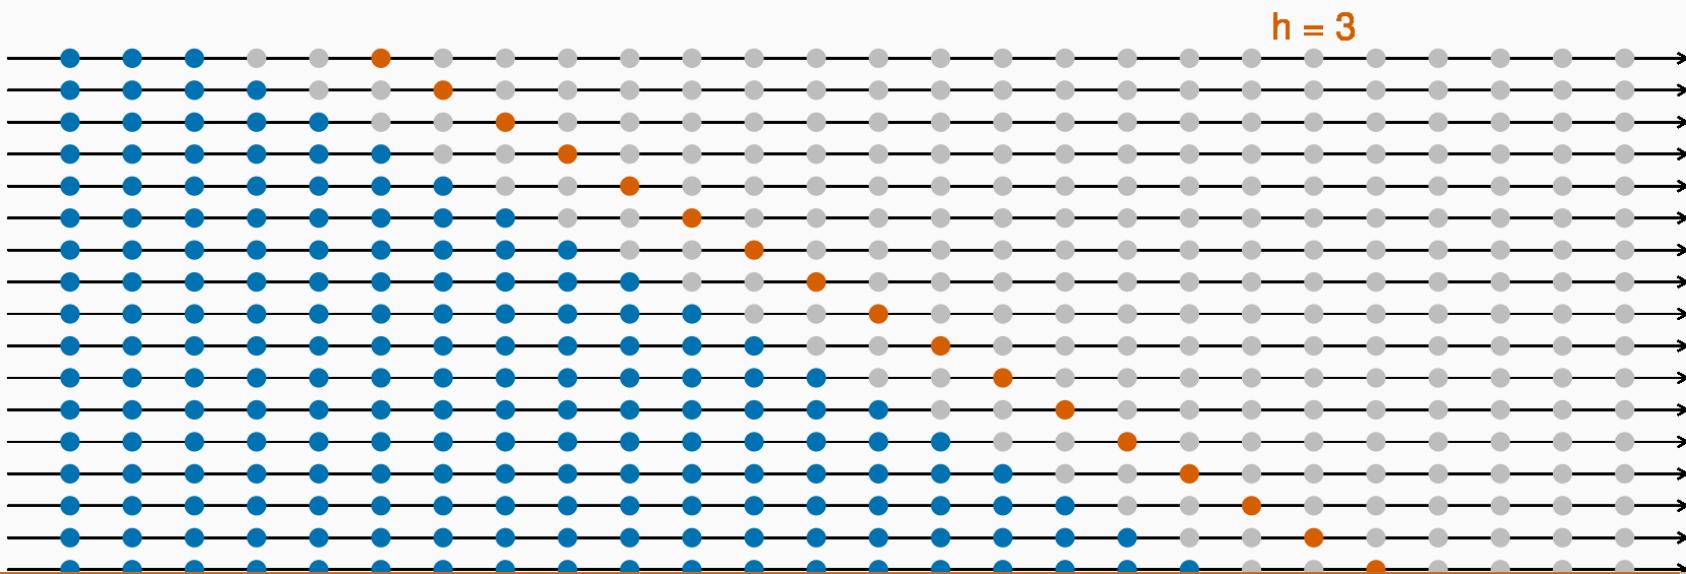


# Time series cross-validation

## Traditional evaluation



## Time series cross-validation

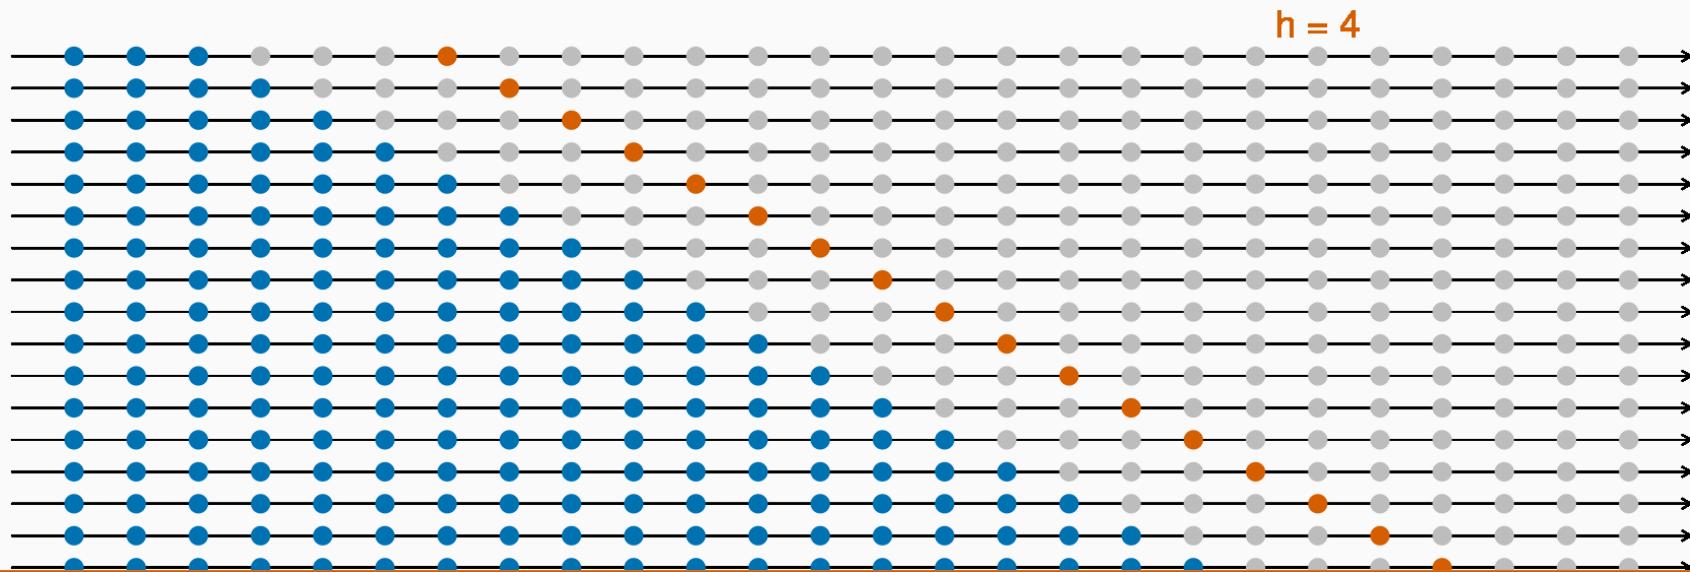


# Time series cross-validation

## Traditional evaluation

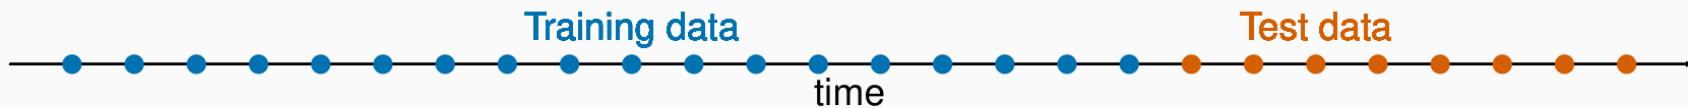


## Time series cross-validation

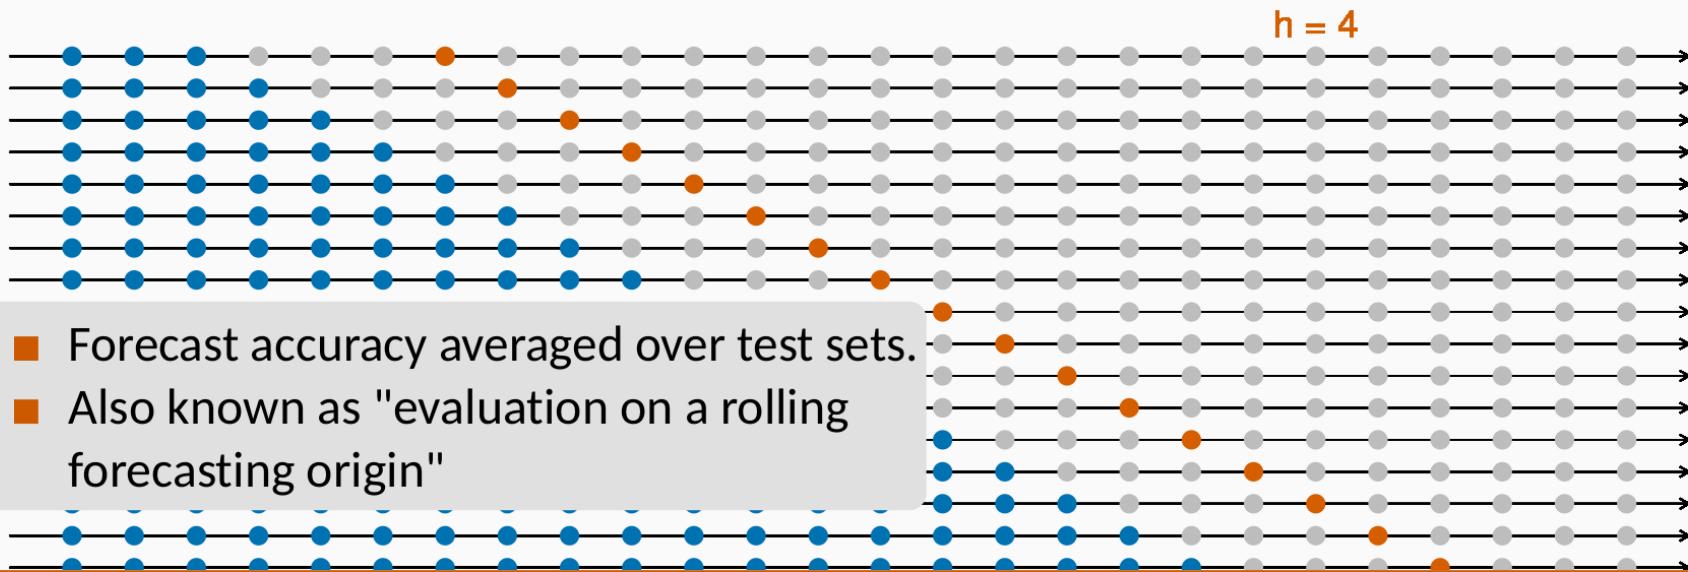


# Time series cross-validation

## Traditional evaluation



## Time series cross-validation



# Time series cross-validation

Stretch with a minimum length of 3, growing by 1 each step.

```
fb_stretch <- fb_stock %>%
  stretch_tsibble(.init = 3, .step = 1) %>%
  filter(.id != max(.id))
```

```
## # A tsibble: 790,650 x 4 [1]
## # Key:      .id [1,255]
##   Date       Close trading_day   .id
##   <date>     <dbl>      <int> <int>
## 1 2014-01-02  54.7        1     1
## 2 2014-01-03  54.6        2     1
## 3 2014-01-06  57.2        3     1
## 4 2014-01-02  54.7        1     2
## 5 2014-01-03  54.6        2     2
## 6 2014-01-06  57.2        3     2
## 7 2014-01-07  57.9        4     2
```

# Time series cross-validation

Estimate RW w/ drift models for each window.

```
fit_cv <- fb_stretch %>%
  model(RW(Close ~ drift()))

## # A mable: 1,255 x 3
## # Key:     .id, Symbol [1,255]
##     .id Symbol `RW(Close ~ drift())` 
##   <int> <chr>           <model>
## 1     1 FB           <RW w/ drift>
## 2     2 FB           <RW w/ drift>
## 3     3 FB           <RW w/ drift>
## 4     4 FB           <RW w/ drift>
## # ... with 1,251 more rows
```

# Time series cross-validation

Produce one step ahead forecasts from all models.

```
fc_cv <- fit_cv %>%
  forecast(h=1)

## # A fable: 1,255 x 5 [1]
## # Key:     .id, Symbol [1,255]
##     .id Symbol trading_day     Close .mean
##   <int> <chr>      <dbl>     <dist> <dbl>
## 1     1 FB          4 N(58, 3.9) 58.4
## 2     2 FB          5 N(59, 2) 59.0
## 3     3 FB          6 N(59, 1.5) 59.1
## 4     4 FB          7 N(58, 1.8) 57.7
## # ... with 1,251 more rows
```

# Time series cross-validation

```
# Cross-validated  
fc_cv %>% accuracy(fb_stock)  
# Training set  
fb_stock %>% model(RW(Close ~ drift())) %>% accuracy()
```

	RMSE	MAE	MAPE
Cross-validation	2.418	1.469	1.266
Training	2.414	1.465	1.261

A good way to choose the best forecasting model is to find the model with the smallest RMSE computed using time series cross-validation.